

# Apache Spark – Apache HBase Connector

**Feature Rich and Efficient Access to HBase  
through Spark SQL**

Weiqing Yang

Mingjie Tang

October, 2017



## About Authors

### ◆ Weiqing Yang

- Contribute to Apache Spark, Apache Hadoop, Apache HBase, Apache Ambari
- Software Engineer at Hortonworks

### ◆ Mingjie Tang

- SparkSQL, Spark Mllib, Spark Streaming, Data Mining, Machine Learning
- Software Engineer at Hortonworks

### ◆ ... All Other SHC Contributors

# Agenda

Motivation

Overview

Architecture & Implementation

Usage & Demo

# Motivation

- ◆ Limited Spark Support in HBase Upstream
  - RDD level
  - But Spark Is Moving to DataFrame/Dataset
- ◆ Existing Connectors in DataFrame Level
  - Complicated Design
    - Embedding Optimization Plan inside Catalyst Engine
    - Stability Impact with Coprocessor
    - Serialized RDD Lineage to HBase
  - Heavy Maintenance Overhead

# Overview

# Apache Spark– Apache HBase Connector (SHC)

## ◆ Combine Spark and HBase

- Spark Catalyst Engine for Query Plan and Optimization
- HBase as Fast Access KV Store
- Implement Standard External Data Source with Build-in Filter, Maintain Easily

## ◆ Full Fledged DataFrame Support

- Spark SQL
- Integrated Language Query

## ◆ High Performance

- Partition Pruning, Data Locality, Column Pruning, Predicate Pushdown
- Use Spark UnhandledFilters API
- Cache Spark HBase Connections

# Data Coder & Data Schema

## Support Different Data Coders

- PrimitiveType: Native Support Java Primitive Types
- Avro: Native Support Avro Encoding/Decoding
- Phoenix: Phoenix Encoding/Decoding
- Plug-In Data Coder
- Can Run on the Top of Existing HBase Tables


## Support Composite Key

- ```
def cat = s"""{
  "table":{"namespace":"default", "name":"shcExampleTable", "tableCoder":"Phoenix"},
  "rowkey":"key1:key2",
  "columns":{
    "col00":{"cf":"rowkey", "col":"key1", "type":"string"},
    "col01":{"cf":"rowkey", "col":"key2", "type":"int"},
    ...
  }
}
```

# Architecture & Implementation



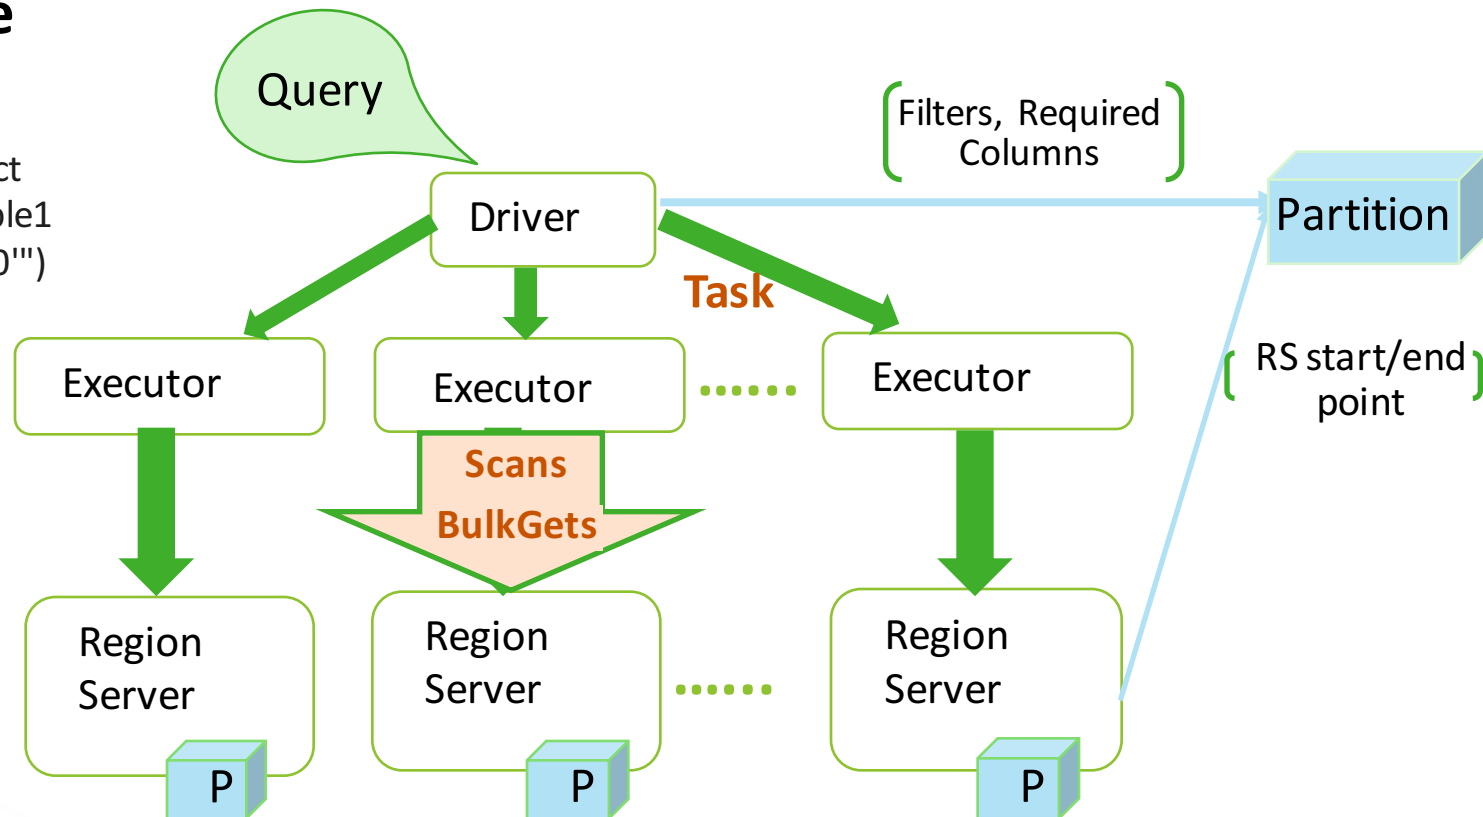
# Architecture



Picture 1. SHC architecture

# Architecture

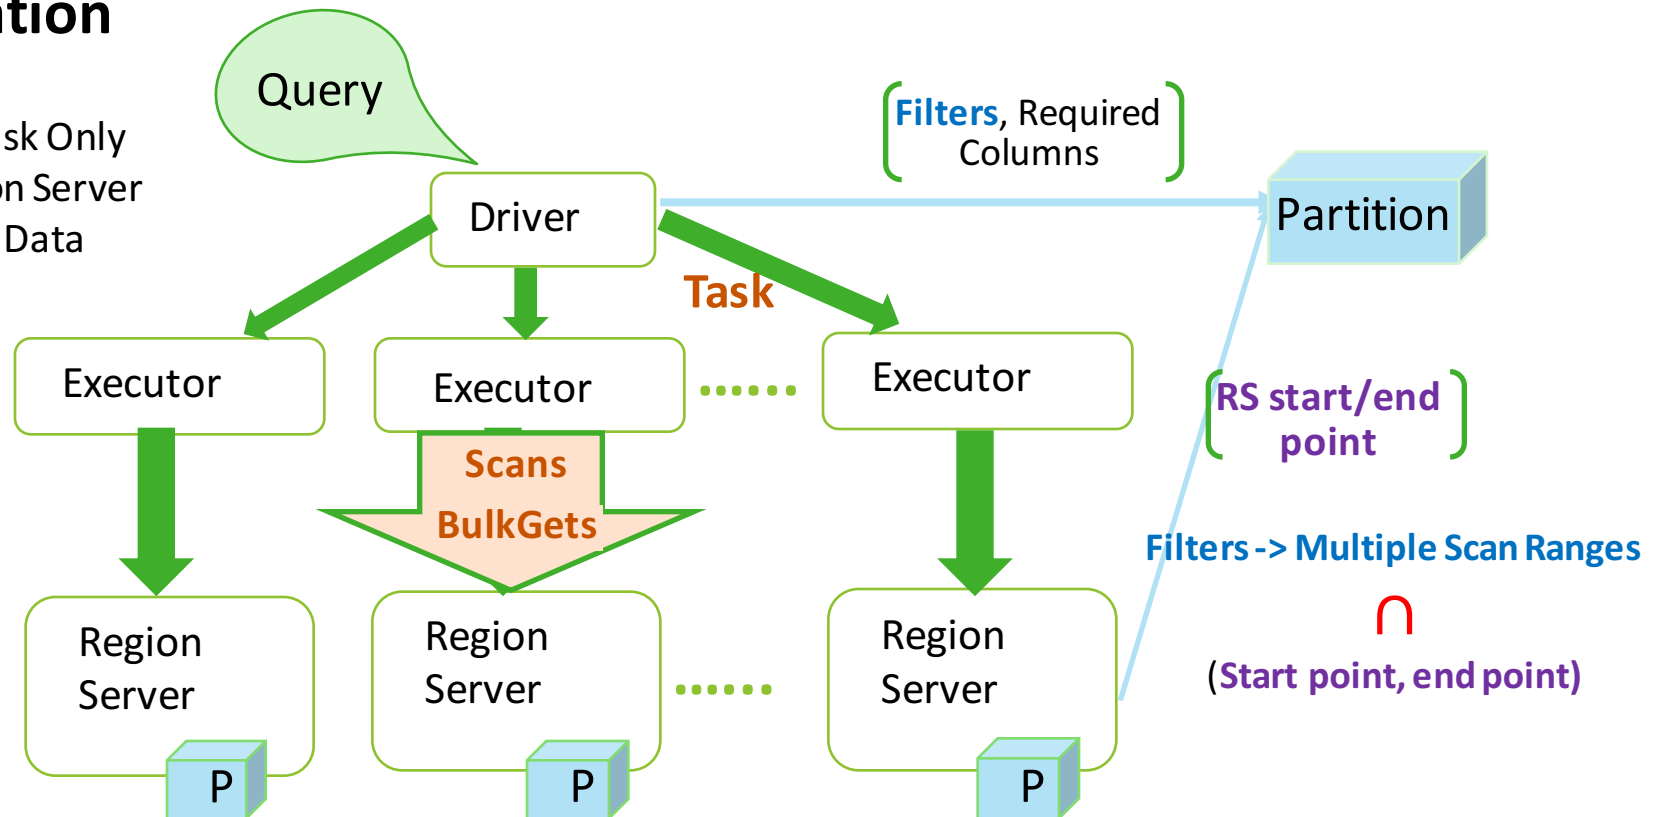
```
sqlContext.sql("select  
count(col1) from table1  
where key < 'row050'")
```



Picture 1. SHC architecture

# Implementation

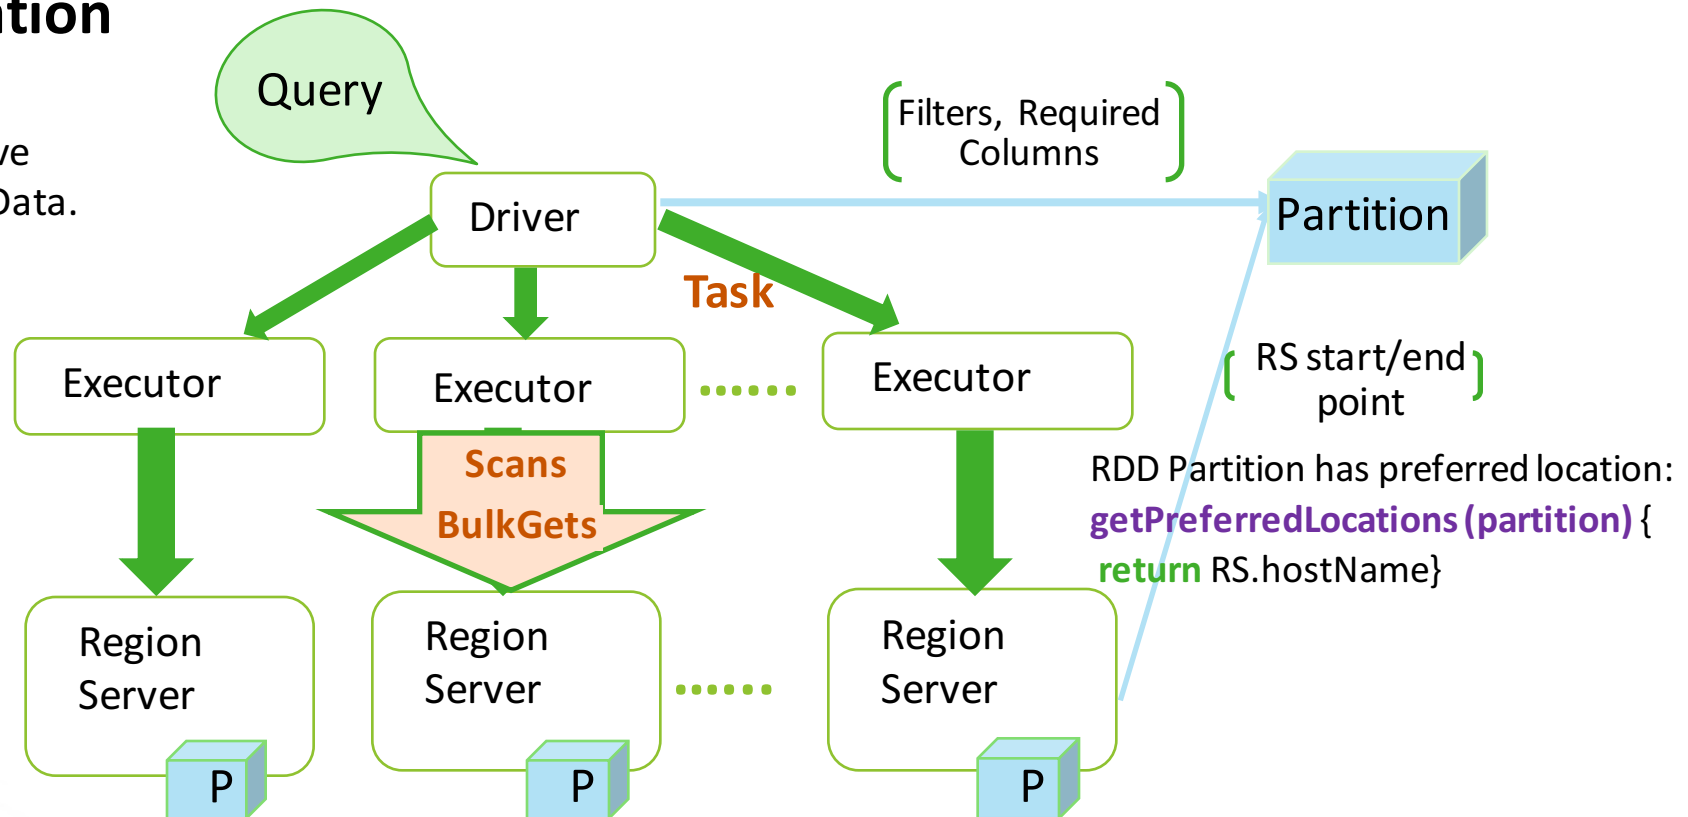
**Partition Pruning:** Task Only Performed in Region Server Holding Requested Data



Picture 1. SHC architecture

# Implementation

**Data Locality:** Move  
Computation to Data.

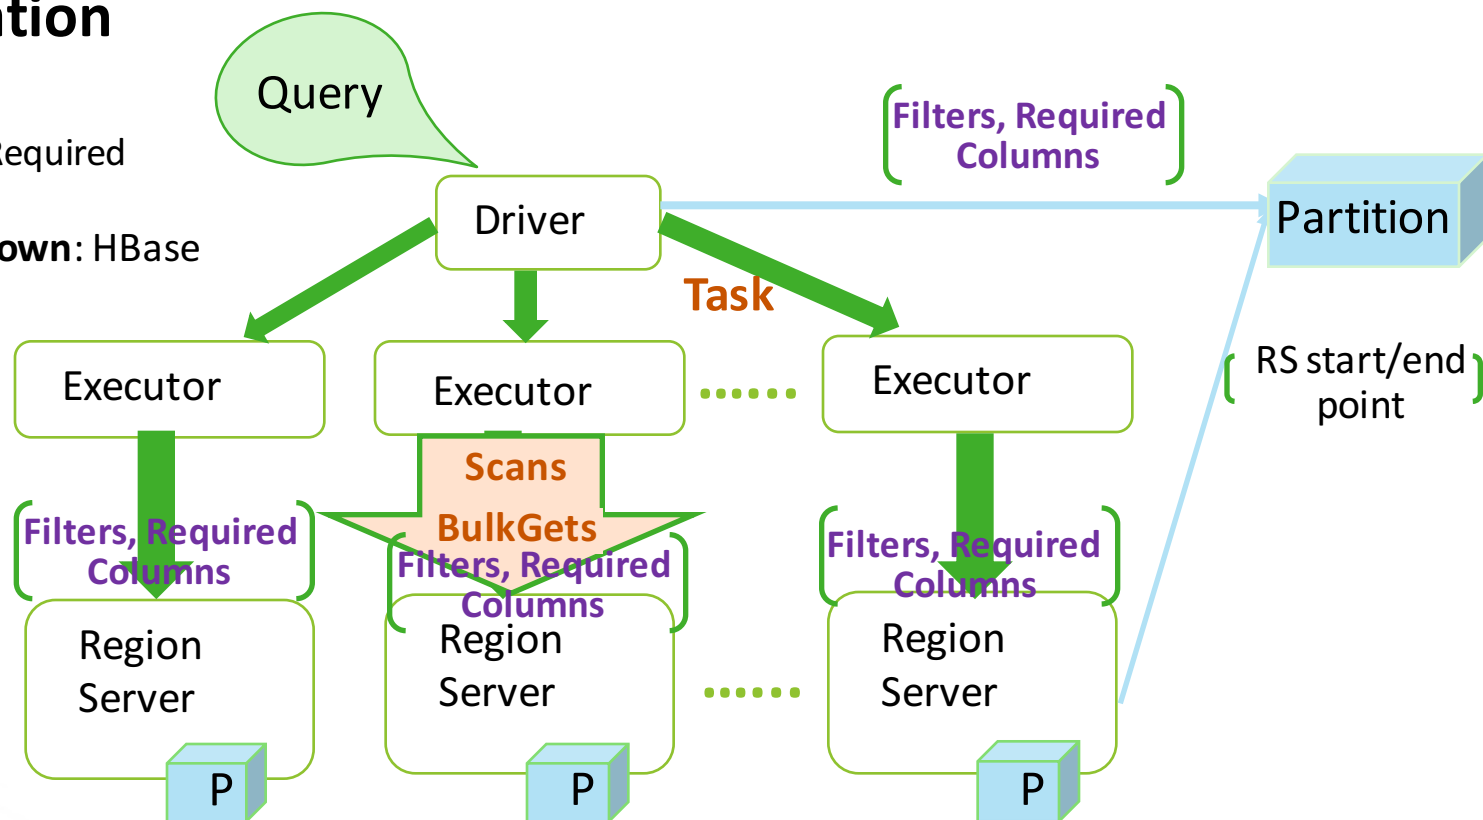


Picture 1. SHC architecture

# Implementation

**Column Pruning:** Required Column

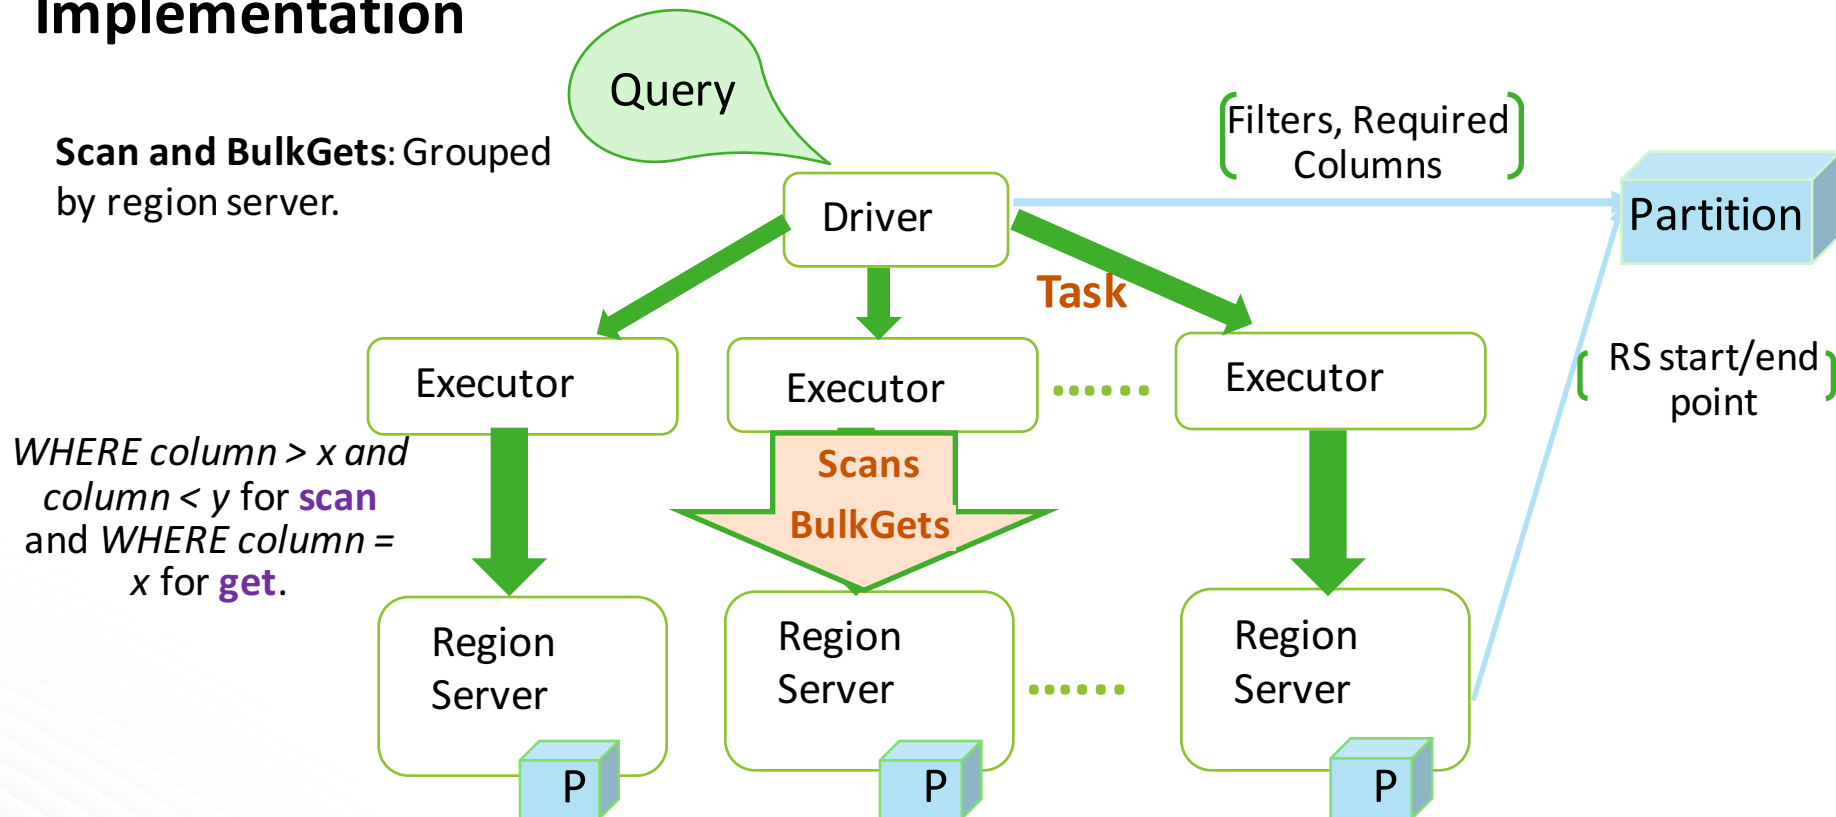
**Predicate Pushdown:** HBase built-in Filters



Picture 1. SHC architecture

# Implementation

**Scan and BulkGets:** Grouped by region server.



Picture 1. SHC architecture

# Usage & Demo

# How to Use SHC?

## ◆ Github

- <https://github.com/hortonworks-spark/shc>

## ◆ SHC Examples

- <https://github.com/hortonworks-spark/shc/tree/master/examples>

## ◆ Apache HBase Jira

- <https://issues.apache.org/jira/browse/HBASE-14789>



# Demo

- ◆ Interactive Jobs through Spark Shell
- ◆ Batch Jobs

# Acknowledgement

- ◆ HBase Community & Spark Community
- ◆ All SHC Contributors, Zhan Zhang

# Reference

## ◆ Hortonworks Public Repo

- <http://repo.hortonworks.com/content/repositories/releases/com/hortonworks/>

## ◆ Apache Spark

- <http://spark.apache.org/>

## ◆ Apache HBase

- <https://hbase.apache.org/>

# Thanks

## Q & A

Emails:  
[wyang@hortonworks.com](mailto:wyang@hortonworks.com)

# BACKUP

# Kerberos Cluster

- ◆ Kerberos Ticket
  - `kinit -kt foo.keytab foouser` or Principle/Keytab
- ◆ Long Running Service
  - `--principal`, `--keytab`
- ◆ Multiple Secure HBase Clusters
  - Spark only Supports Single Secure HBase Cluster
  - Use SHC Credential Manager
  - Refer [LRJobAccessing2Clusters Example in github](#)

# Usage

Define the catalog for the schema mapping:

```
def catalog = s"""{
  |"table":{"namespace":"default", "name":"phoenixTable",
  |  "tableCoder":"Phoenix", "version":"2.0"},
  |"rowkey":"key",
  |"columns":{
  |"col0":{"cf":"rowkey", "col":"key", "type":"string"},
  |"col1":{"cf":"cf1", "col":"col1", "type":"boolean"},
  |"col2":{"cf":"cf2", "col":"col2", "type":"double"},
  |"col3":{"cf":"cf3", "col":"col3", "type":"float"},
  |"col4":{"cf":"cf4", "col":"col4", "type":"int"},
  |"col5":{"cf":"cf5", "col":"col5", "type":"bigint"},
  |"col6":{"cf":"cf6", "col":"col6", "type":"smallint"},
  |"col7":{"cf":"cf7", "col":"col7", "type":"string"},
  |"col8":{"cf":"cf8", "col":"col8", "type":"tinyint"}
  |}
  |}""".stripMargin
```

## Usage

### ◆ Prepare the data and populate the HBase table

```
val data = (0 to 255).map { i => HBaseRecord(i, "extra") }
```

```
sc.parallelize(data).toDF.write.options(  
  Map(HBaseTableCatalog.tableCatalog -> catalog, HBaseTableCatalog.newTable -> "5"))  
  .format("org.apache.spark.sql.execution.datasources.hbase")  
  .save()
```



# Usage

## ◆ Load the DataFrame

```
def withCatalog(cat: String): DataFrame = {  
  sqlContext  
    .read  
    .options(Map(HBaseTableCatalog.tableCatalog->cat))  
    .format("org.apache.spark.sql.execution.datasources.hbase")  
    .load()  
}
```

```
val df = withCatalog(catalog)
```

# Usage

## ◆ Query

### Language integrated query:

```
val s = df.filter((( $"col0" <= "row050" && $"col0" > "row040" ) ||  
    $"col0" === "row005" && ( $"col4" === 1 || $"col4" === 42 ) )  
    .select( "col0", "col1", "col4" )
```

### SQL:

```
val s = df.filter((( $"col0" <= "row050" && $"col0" > "row040" )  
df.registerTempTable( "table" )  
sqlContext.sql( "select count(col1) from table" ).show
```

# Usage

## ◆ Work with different data sources

*// Part 1: write data into Hive table and read data from it*

```
val df1 = sql("SELECT * FROM shcHiveTable")
```

*// Part 2: read data from Hbase table*

```
val df2 = withCatalog(cat)
```

*// Part 3: join the two dataframes*

```
val s1 = df1.filter($"key" <= "40").select("key", "col1")
```

```
val s2 = df2.filter($"key" <= "20" && $"key" >= "1").select("key", "col2")
```

```
val result = s1.join(s2, Seq("key"))
```

```
result.show()
```