

Build, Scale, and Deploy Deep Learning Pipelines Using Apache Spark

Sue Ann Hong, Databricks

Tim Hunter, Databricks

#EUdd3



About Us

Sue Ann Hong

- Software engineer @ Databricks
- Ph.D. from CMU in Machine Learning



Tim Hunter

- Software engineer @ Databricks
- Ph.D. from UC Berkeley in Machine Learning
- Very early Spark user



This talk

- Deep Learning at scale: current state
- Deep Learning Pipelines: the vision
- End-to-end workflow with DL Pipelines
- Future

Deep Learning at Scale

: current state

What is Deep Learning?

- A set of machine learning techniques that use layers that transform numerical inputs
 - Classification
 - Regression
 - Arbitrary mapping
- Popular in the 80's as Neural Networks
- Recently came back thanks to advances in data collection, computation techniques, and hardware.

Success of Deep Learning

Tremendous success for applications with complex data

- AlphaGo
- Image interpretation
- Automatic translation
- Speech recognition

But requires a lot of effort

- No exact science around deep learning
- Success requires many engineer-hours
- Low level APIs with steep learning curve
- Not well integrated with other enterprise tools
- Tedious to distribute computations

What does Spark offer?

Very little in Apache Spark MLlib itself (multilayer perceptron)

Many Spark packages

Integrations with existing DL libraries

- Deep Learning Pipelines (from Databricks)
- Caffe (CaffeOnSpark)
- Keras (Elephas)
- mxnet
- Paddle
- TensorFlow (TensorFlow on Spark, TensorFrames)
- CNTK (mmlspark)

Implementations of DL on Spark


- BigDL
- DeepDist
- DeepLearning4J
- MLlib
- SparkCL
- SparkNet

Deep Learning in industry


- Currently limited adoption
- Huge potential beyond the industrial giants
- How do we accelerate the road to massive availability?

Deep Learning Pipelines

Deep Learning Pipelines: Deep Learning with Simplicity


- Open-source Databricks library
- Focuses on *ease of use* and *integration*
 - without sacrificing *performance*
- Primary language: Python 
- Uses Apache Spark for *scaling out* common tasks
- Integrates with MLlib Pipelines to capture the ML workflow concisely

A typical Deep Learning workflow



- **Load data** (images, text, time series, ...)
- **Interactive work**
- **Train**
 - Select an architecture for a neural network
 - Optimize the weights of the NN
- **Evaluate** results, potentially re-train
- **Apply:**
 - Pass the data through the NN to produce new features or output

A typical Deep Learning workflow



Load data

- Image loading in Spark

Interactive work

- Pre-trained models

Train

- Transfer learning Part 1

- Distributed tuning Part 2

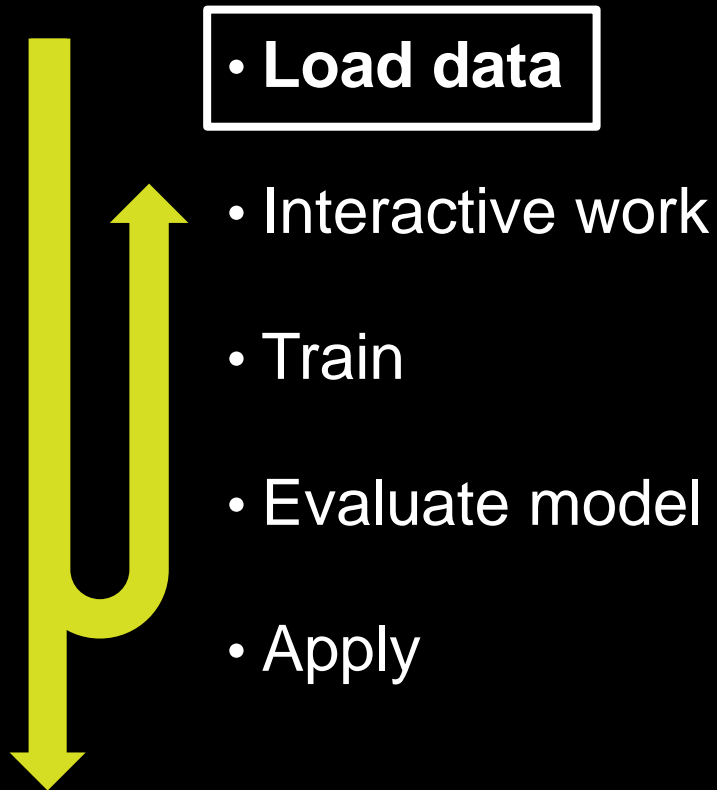
Evaluate

Apply

- Distributed batch prediction
- Deploying models in SQL

End-to-End Workflow with Deep Learning Pipelines

Deep Learning Pipelines



Adds support for images in Spark

```
from sparkdl import readImages
image_df = readImages(sample_img_dir)
```


filePath	image
dbfs:/tmp/flower_photos/sample/100080576_f52e8ee070_n.jpg	▶ {"mode": "RGB", "height": 263, "width": 320, "nChannels": 3, "data": "h4eFioqljo6OkZC"} ▶ {"mode": "RGB", "height": 263, "width": 320, "nChannels": 3, "data": "h4eFioqljo6OkZC"}
dbfs:/tmp/flower_photos/sample/100930342_92e8746431_n.jpg	▶ {"mode": "RGB", "height": 209, "width": 320, "nChannels": 3, "data": "Ey4PEC8QDi8QI"} ▶ {"mode": "RGB", "height": 209, "width": 320, "nChannels": 3, "data": "Ey4PEC8QDi8QI"}

- ImageSchema, reader, conversion functions to/from numpy arrays
- Most of the tools we'll describe work on ImageSchema columns

Upcoming: built-in support in Spark

- Spark-21866
- Contributing image format & reading to Spark
- Targeted for Spark 2.3
- Joint work with Microsoft

Deep Learning Pipelines



- Load data

- **Interactive work**

- Train

- Evaluate model

- Apply

Applying popular models

- Popular pre-trained models accessible through MLlib Transformers

```
predictor = DeepImagePredictor(inputCol="image",  
                                outputCol="predicted_labels",  
                                modelName="InceptionV3")  
  
predictions_df = predictor.transform(image_df)
```

Applying popular models

```
predictor = DeepImagePredictor(inputCol="image",  
                                outputCol="predicted_labels",  
                                modelName="InceptionV3")  
  
predictions_df = predictor.transform(image_df)
```

filePath

dbfs:/tmp/flower_photos/sample/100080576_f52e8ee070_n.jpg



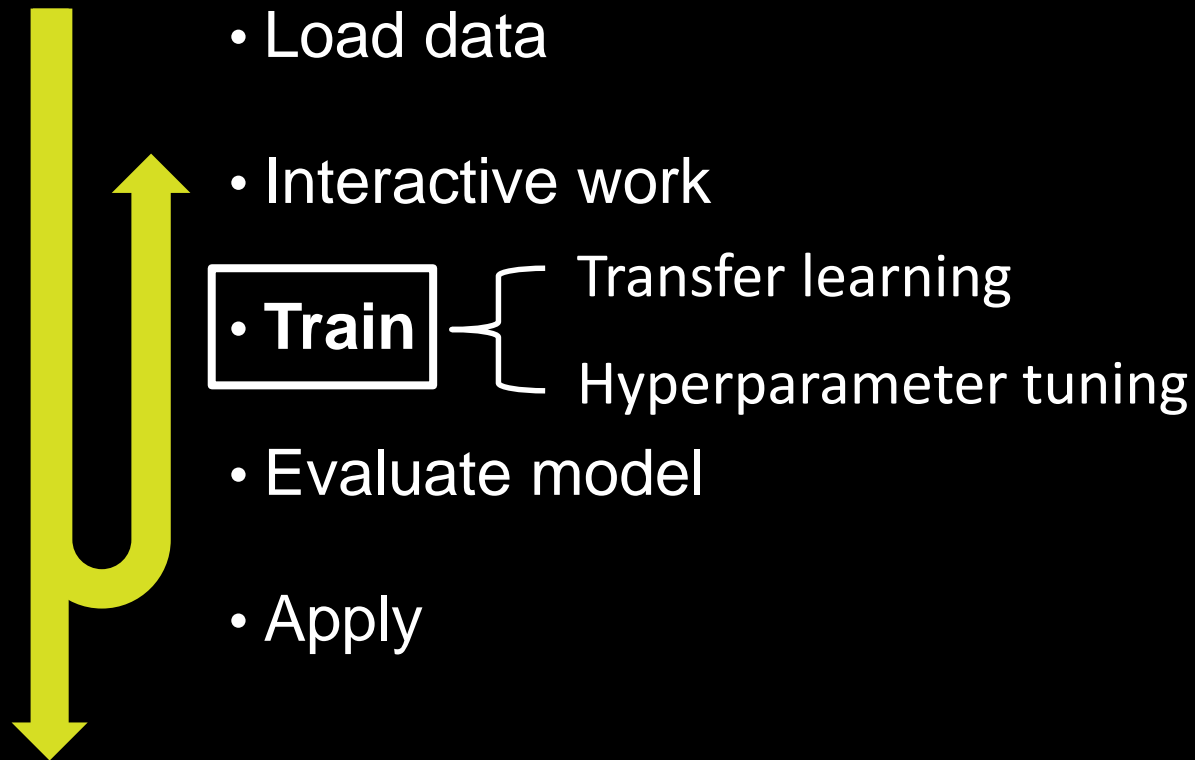
predicted_labels

```
▶ [{"class": "n11939491", "description": "daisy", "probability": 0.8805494},  
  {"class": "n02219486", "description": "ant", "probability": 0.0020712544},  
  {"class": "n02206856", "description": "bee", "probability": 0.00084249553},  
  {"class": "n02165456", "description": "ladybug", "probability": 0.000675952},  
  {"class": "n03691459", "description": "loudspeaker", "probability": 0.00063085736},  
  {"class": "n02190166", "description": "fly", "probability": 0.0006220306},  
  {"class": "n02281406", "description": "sulphur_butterfly", "probability": 0.0006154293},  
  {"class": "n07930864", "description": "cup", "probability": 0.00055486656},  
  {"class": "n02112018", "description": "Pomeranian", "probability": 0.0004993835},  
  {"class": "n07745940", "description": "strawberry", "probability": 0.00048750438}]
```


dbfs:/tmp/flower_photos/sample/100930342_92e8746431_n.jpg

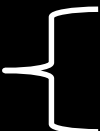
```
▶ [{"class": "n03930313", "description": "picket_fence", "probability": 0.18473865},
```

Deep Learning Pipelines



Deep Learning Pipelines

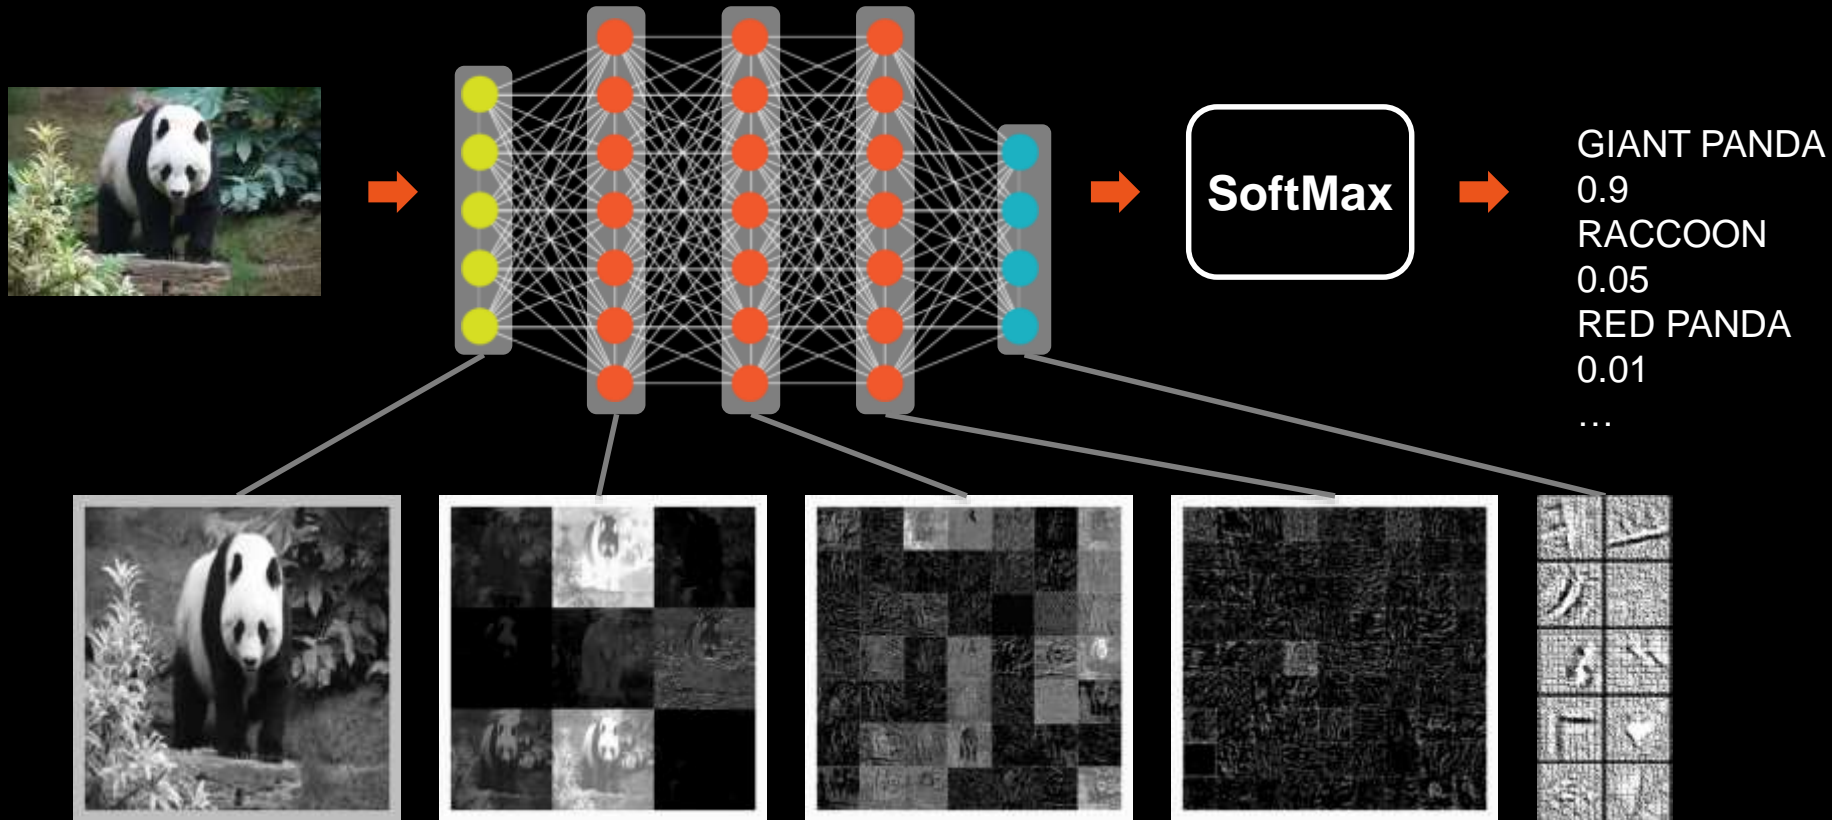


- Load data
- Interactive work
- **Train** 
 - Transfer learning**
 - Hyperparameter tuning
- Evaluate model
- Apply

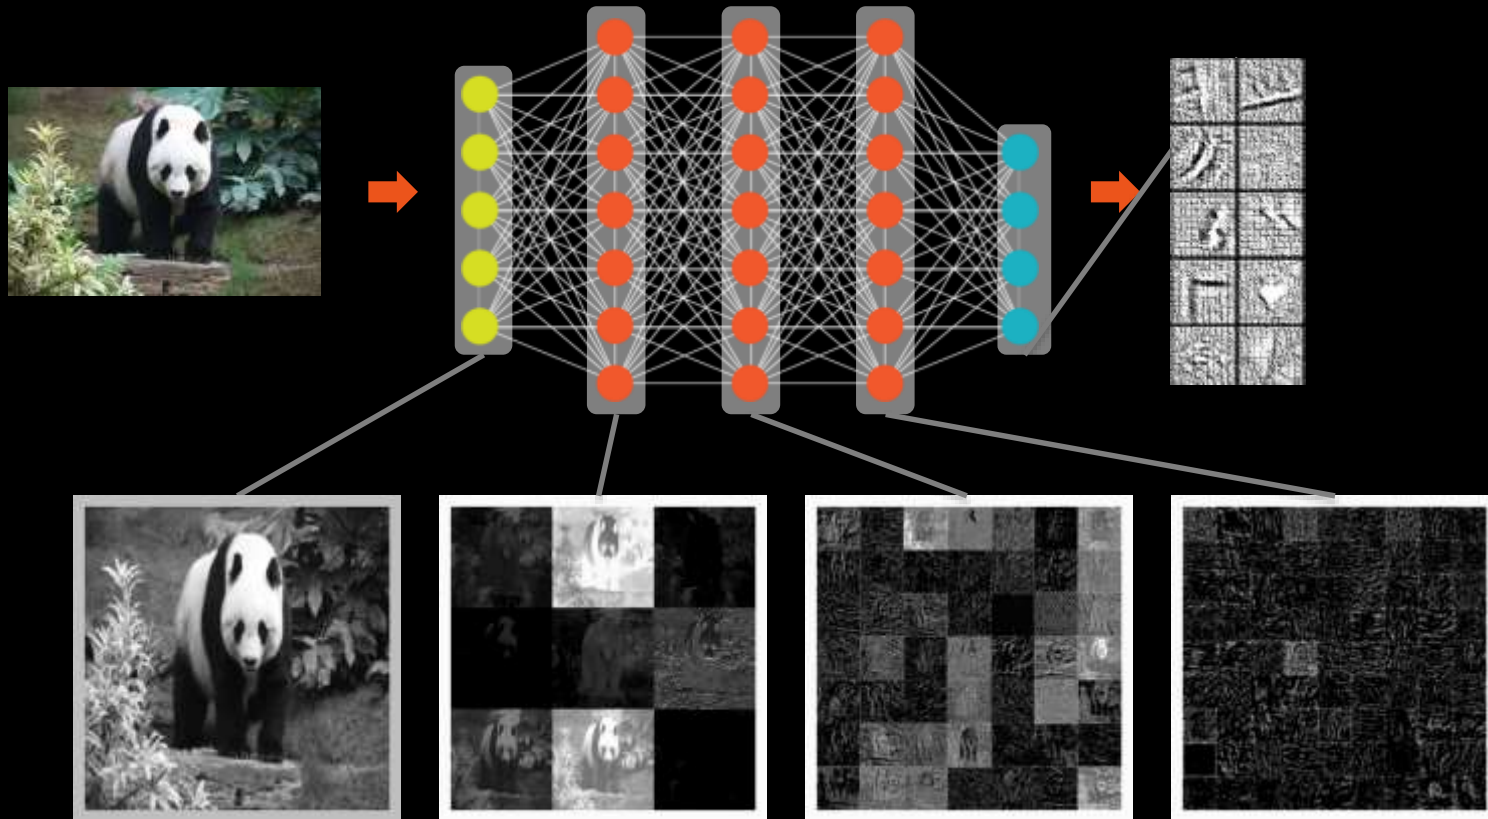
Transfer learning

- Pre-trained models may not be directly applicable
 - New domain, e.g. shoes
- Training from scratch requires
 - Enormous amounts of data
 - A lot of compute resources & time
- Idea: intermediate representations learned for one task may be useful for other related tasks


Transfer Learning



Transfer Learning




Transfer Learning



**Classifie
r**

Transfer Learning



Classifier



Rose: 0.7


Daisy: 0.3

0.01
...

MLlib Pipelines primer


- MLlib: the machine learning library included with Spark
- Transformer
 - Takes in a Spark dataframe
 - Returns a Spark dataframe with new column(s) containing “transformed” data
 - e.g. a Model is a Transformer
- Estimator
 - A learning algorithm, e.g. `lr = LogisticRegression()`
 - Produces a Model via `lr.fit()`
- Pipeline: a sequence of Transformers and Estimators

Transfer Learning as a Pipeline



Rose / Daisy

Transfer Learning as a Pipeline



**MLlib
Pipeline**

Transfer Learning as a Pipeline


```
featurizer = DeepImageFeaturizer(inputCol="image",  
                                   outputCol="features",  
                                   modelName="InceptionV3")  
  
lr = LogisticRegression(labelCol="label")  
  
p = Pipeline(stages=[featurizer, lr])  
  
p_model = p.fit(train_df)
```

Transfer Learning

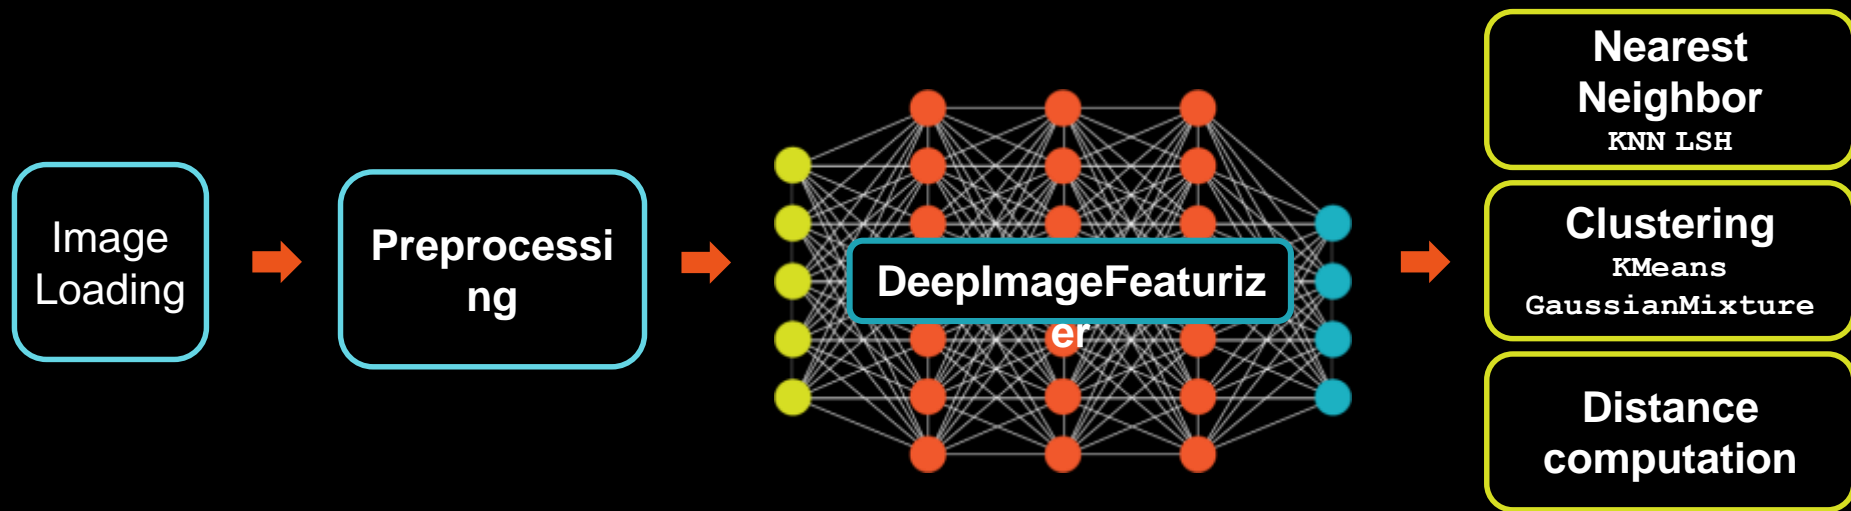
- Usually for classification tasks
 - Similar task, new domain
- But other forms of learning leveraging learned representations can be loosely considered transfer learning



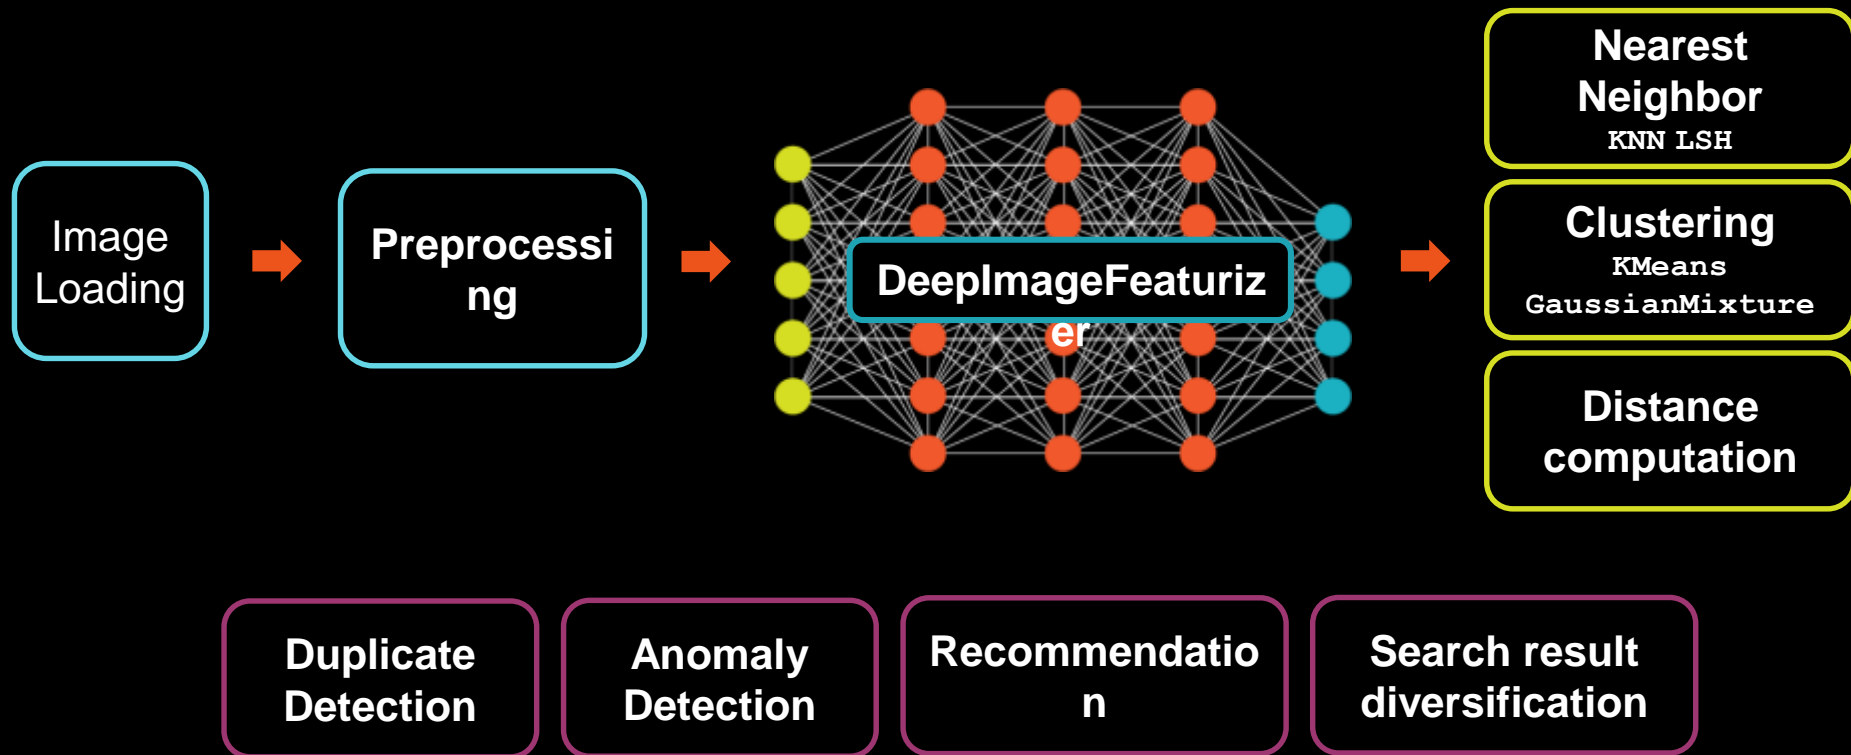
Featurization for similarity-based ML



Featurization for similarity-based ML

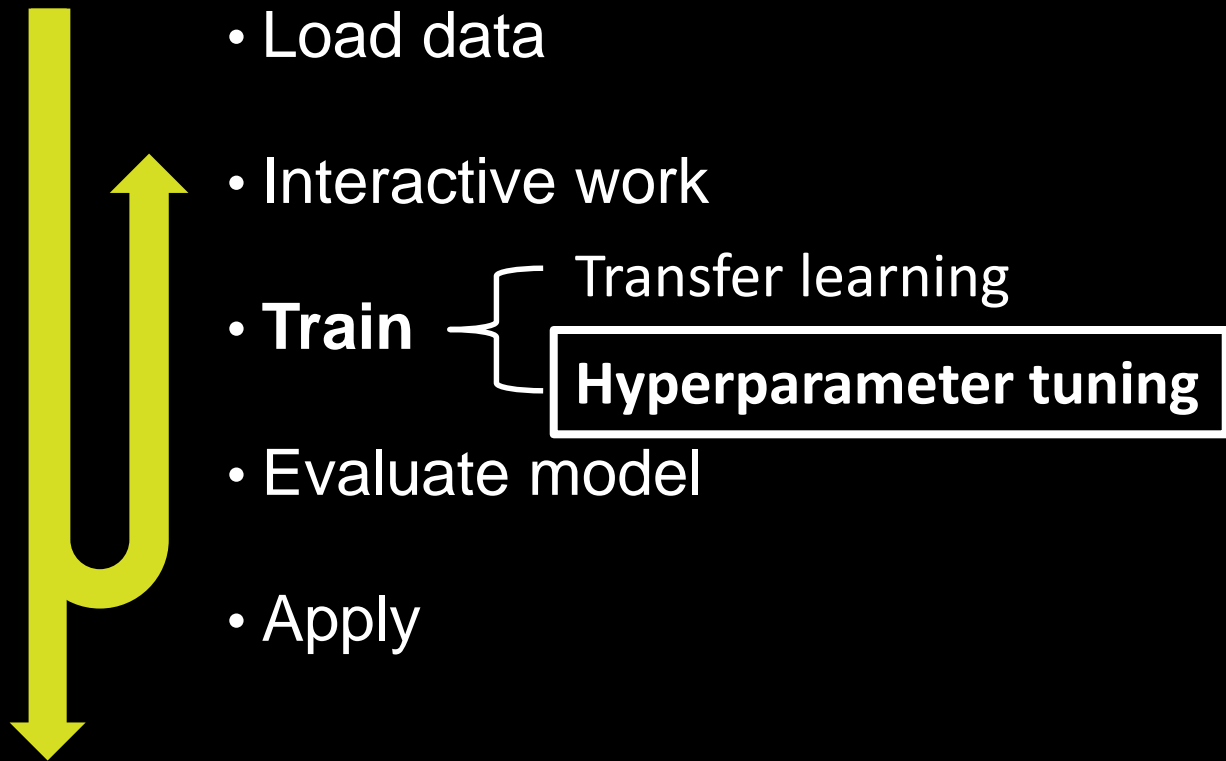


Featurization for similarity-based ML




Break?

Deep Learning Pipelines



Distributed Hyperparameter Tuning



MLlib Components

Estimators

- Learning algorithms
- Implement

```
model = fit(train_df, params)
```

 e.g. learning rate

Evaluators

- Metric used to measure model's goodness on validation data
- e.g. `BinaryClassificationEvaluator` computes accuracy

Model Selection (Parameter Search)

Estimators

Map of params

Evaluators

```
paramGrid = ( ParamGridBuilder()  
    .addGrid(hashingTF.numFeatures, [10, 100])  
    .addGrid(lr.regParam, [0.1, 0.01])  
    .build() )
```

```
cv = CrossValidator(  
    estimator=pipeline,  
    estimatorParamMaps=paramGrid,  
    evaluator=BinaryClassificationEvaluator())
```

```
best_model = cv.fit(train_df)
```


Deep Learning Estimators

- Transfer learning Pipelines
- KerasImageFileEstimator

Deep Learning Estimators

- **Transfer learning Pipelines**
- KerasImageFileEstimator

Distributed Tuning Transfer Learning



Distributed Tuning Transfer Learning

```
pipeline = Pipeline([DeepImageFeaturizer(), LogisticRegression()])

paramGrid = ( ParamGridBuilder()
               .addGrid(modelName=["ResNet50", "InceptionV3"]) )

cv = CrossValidator(estimator= pipeline,
                    estimatorParamMaps=paramGrid,
                    evaluator=BinaryClassificationEvaluator(),
                    numFolds=3)

best_model = cv.fit(train_df)
```

Deep Learning Estimators

- Transfer learning Pipelines
- **KerasImageFileEstimator**

Keras

- A popular, declarative interface to build DL models
- High level, expressive API in python
- Executes on TensorFlow, Theano, CNTK

```
model = Sequential()  
model.add(Dense(32, input_dim=784))  
model.add(Activation('relu'))
```

Keras Estimator

```
model = Sequential()  
model.add(...)  
model.save(model_filename)  
  
estimator = KerasImageFileEstimator(  
    kerasOptimizer="adam",  
    kerasLoss="categorical_crossentropy",  
    kerasFitParams={"batch_size":100},  
    modelFile=model_filename)  
model = model.fit(dataframe)
```


Keras Estimator in Model Selection

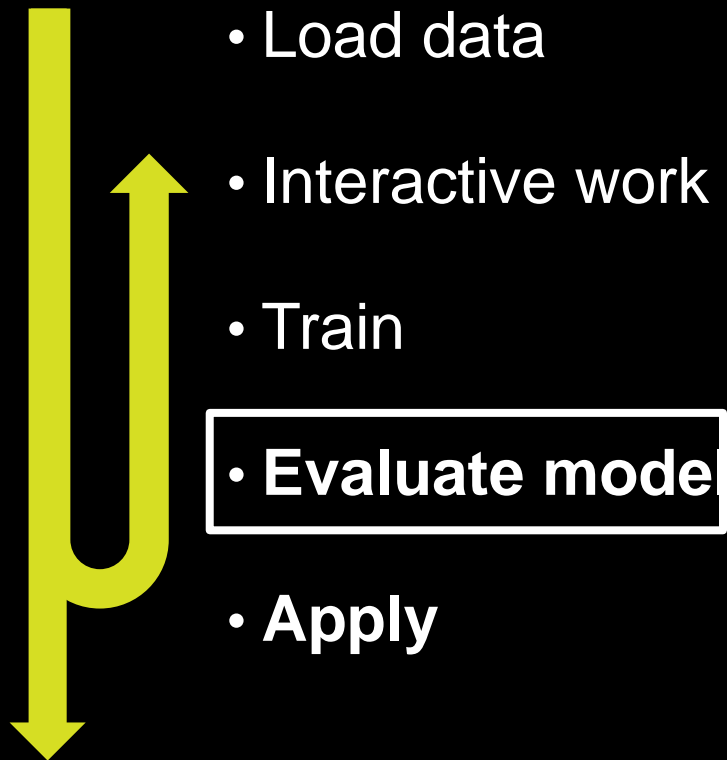
```
estimator = KerasImageFileEstimator(  
    kerasOptimizer="adam",  
    kerasLoss="categorical_crossentropy",  
    kerasFitParams={"batch_size":100},  
    modelFile=model_filename)  
model = model.fit(dataframe)
```

```
estimator = KerasImageFileEstimator(  
    kerasOptimizer="adam",  
    kerasLoss="categorical_crossentropy")  
  
paramGrid = ( ParamGridBuilder()  
    .addGrid(kerasFitParams=[{"batch_size":100}, {"batch_size":200}])  
    .addGrid(modelFile=[model1, model2]) )
```

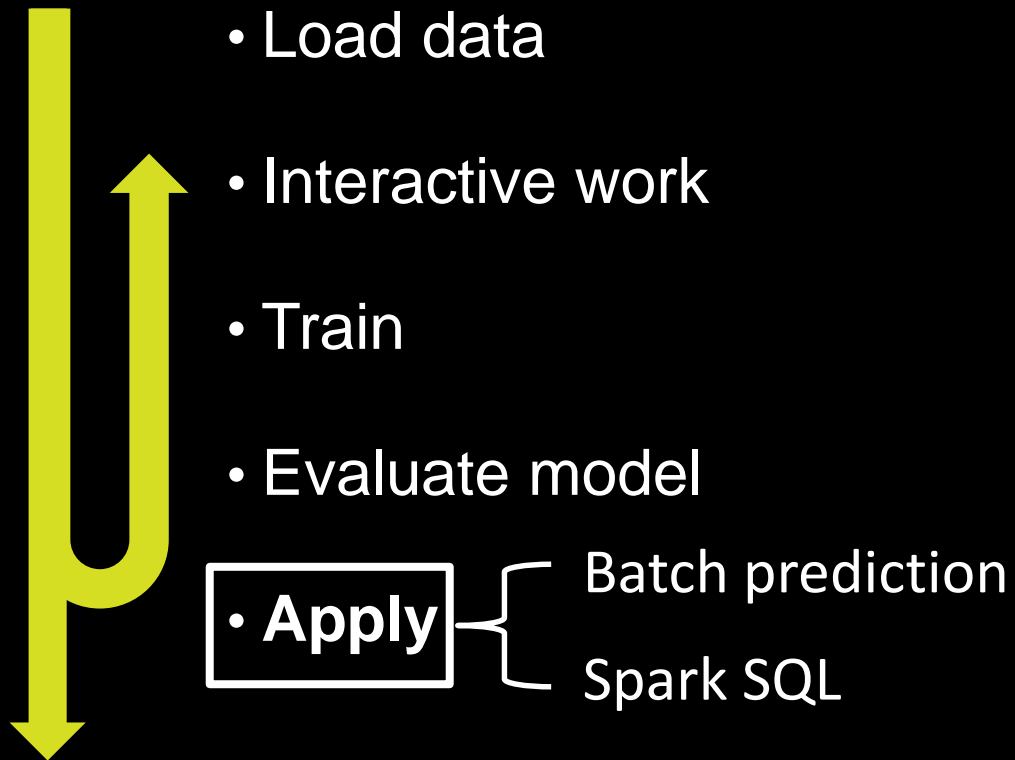
Keras Estimator in Model Selection

```
estimator = KerasImageFileEstimator(  
    kerasOptimizer="adam",  
    kerasLoss="categorical_crossentropy")  
  
paramGrid = ( ParamGridBuilder()  
    .addGrid(kerasFitParams=[{"batch_size":100}, {"batch_size":200}])  
    .addGrid(modelFile=[model1, model2]) )  
  
cv = CrossValidator(estimator=estimator,  
    estimatorParamMaps=paramGrid,  
    evaluator=BinaryClassificationEvaluator(),  
    numFolds=3)  
  
best_model = cv.fit(train_df)
```


Deep Learning Pipelines



Deep Learning Pipelines



Deep Learning Pipelines



- Load data
- Interactive work
- Train
- Evaluate model
- **Apply** {

Batch prediction

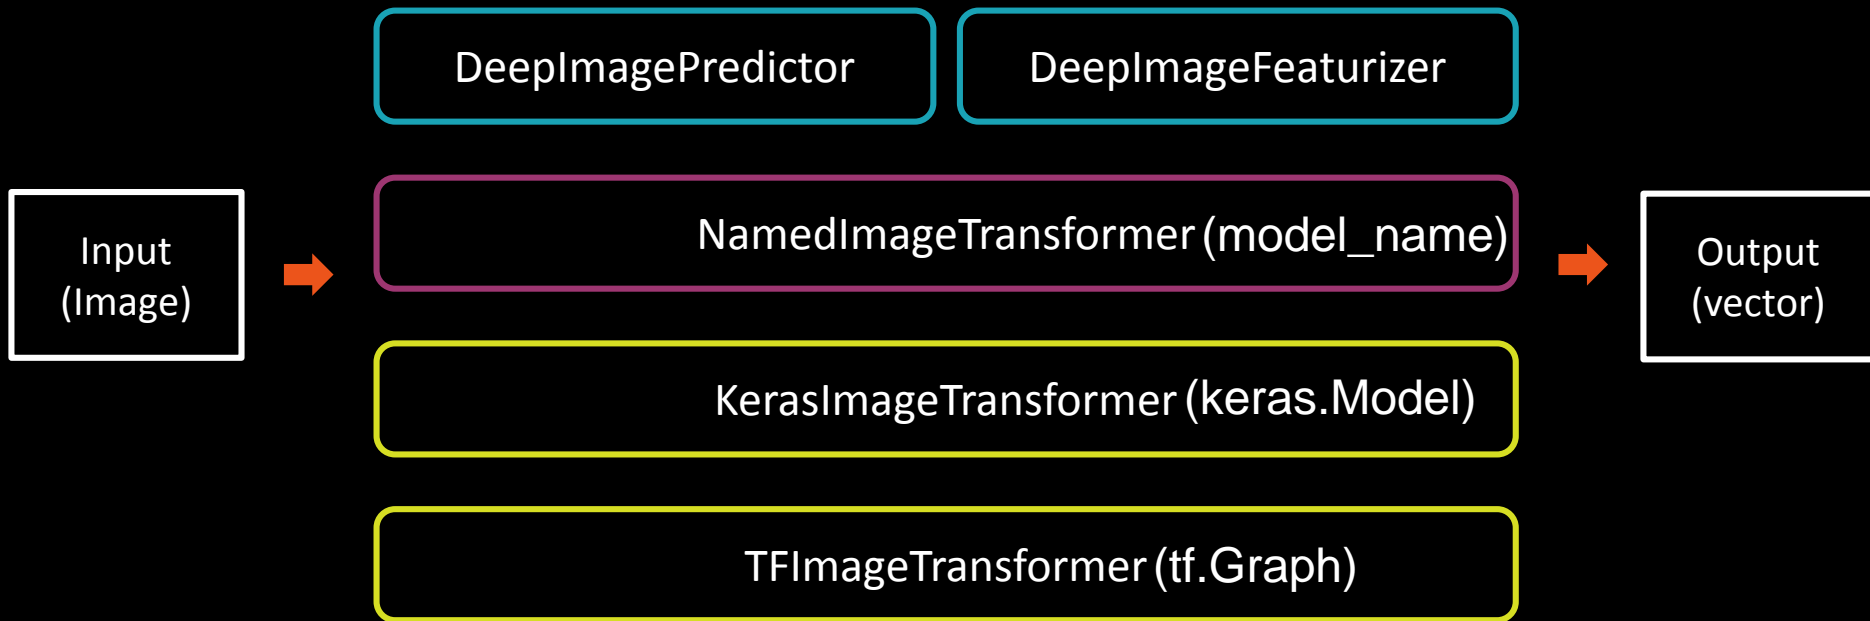
Spark SQL

Batch prediction as an MLlib Transformer

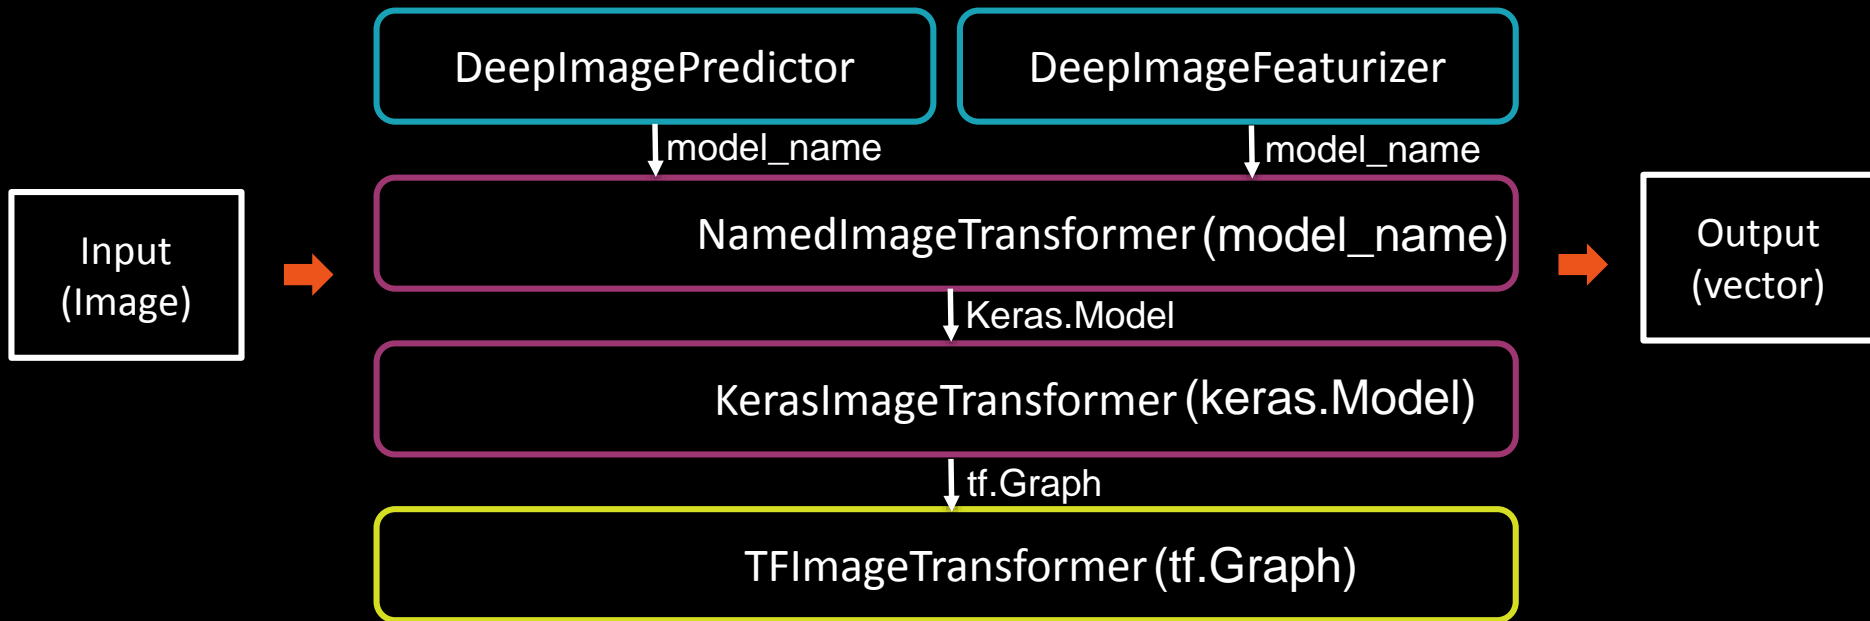
- Recall a model is a Transformer in MLlib

```
predictor = XXTransformer(inputCol="image",  
                           outputCol="predictions",  
                           modelSpecification={...})  
predictions = predictor.transform(test_df)
```

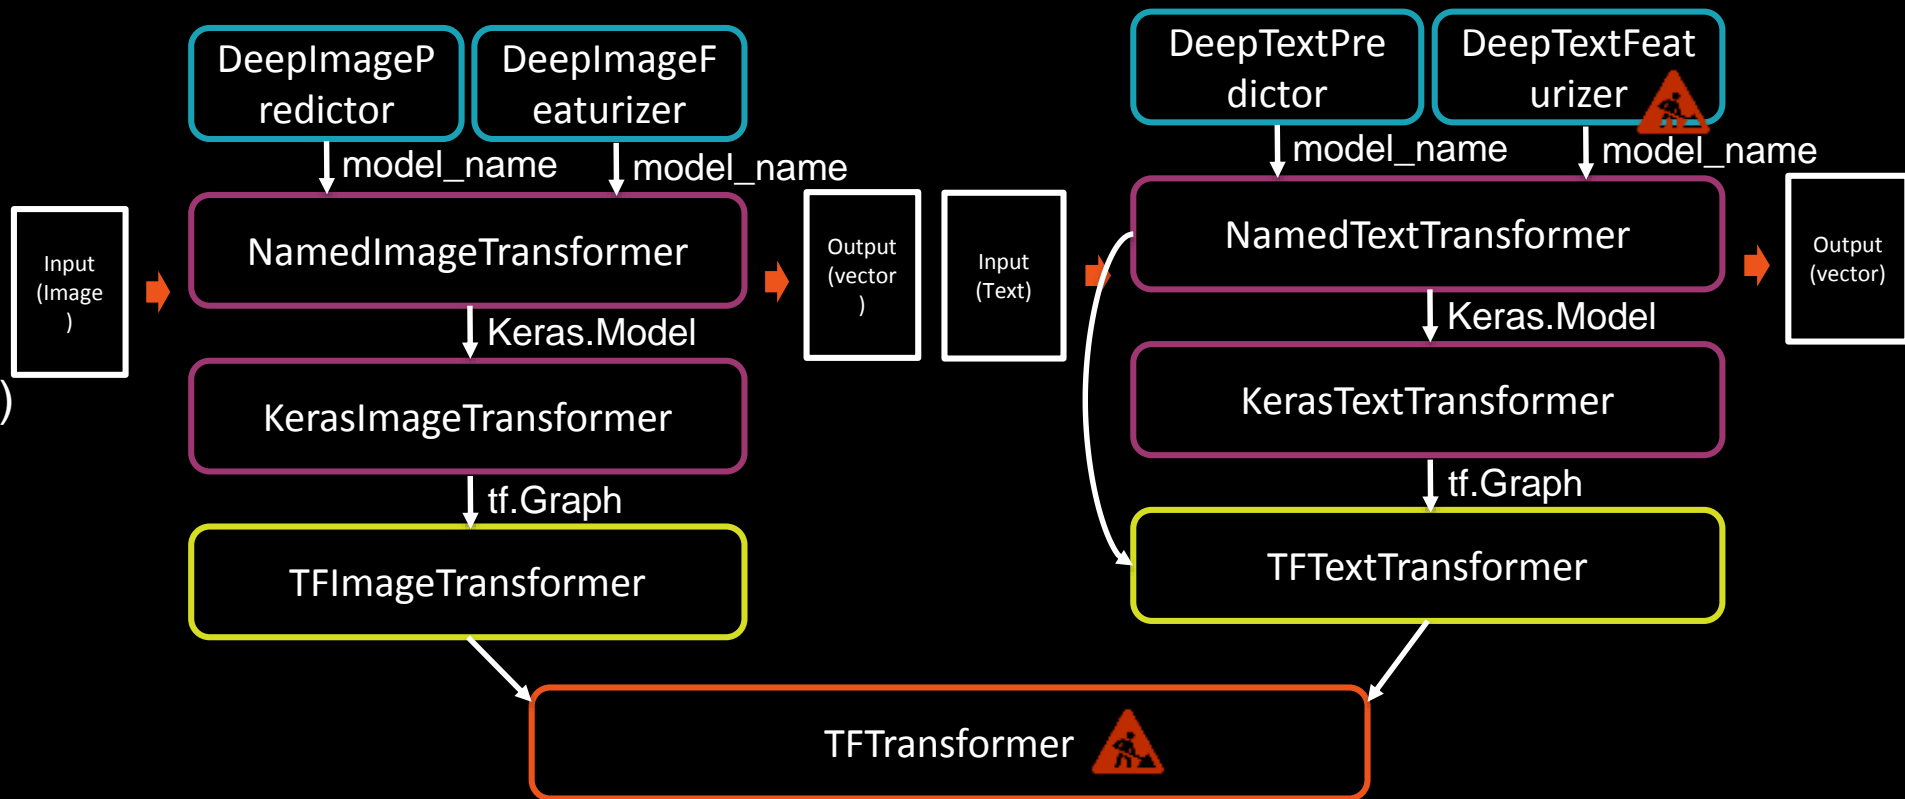
Hierarchy of DL transformers for images



Hierarchy of DL transformers for images



Hierarchy of DL transformers



Defined by `tf.Graph`


1. Convert ImageSchema data into a vector
2. Use tensorframes to
 - Efficiently distribute `tf.Graph` to workers
 - Apply the graph to the partitioned data
3. Return the result as a `spark.ml.Vector`

Aside: high performance with Spark

- TensorFlow (and other frameworks) have 2 mechanisms to ingest data:
 - memory-based API (tensors)
 - file-based API (Queue, ...)
- DLP makes all data transfers in memory
- Spark responsible for reading and assembling data
- Low-level transformations handled by TensorFrames


High performance with Spark

- Every DL transform is a TensorFlow graph




High performance with Spark

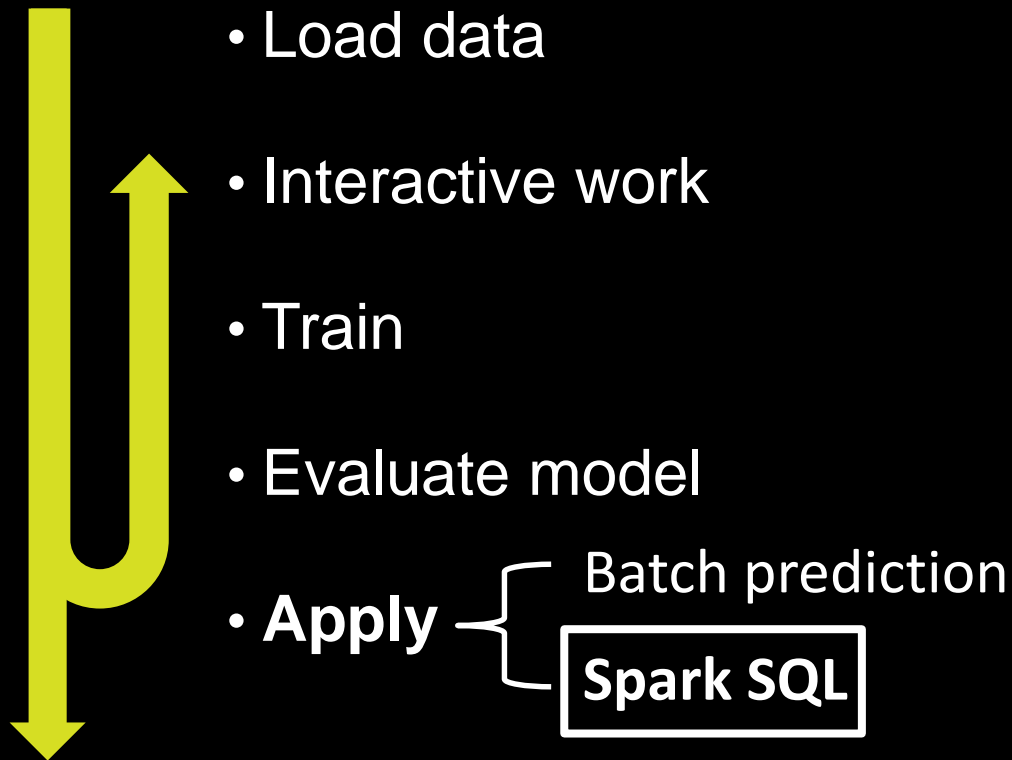
- Every DL transform is a TensorFlow graph
- Applied on each of the rows of the dataframe



Example: running batch inference



Deep Learning Pipelines



Shipping predictors in SQL

Take a trained model / Pipeline, register a SQL UDF usable by *anyone* in the organization

```
registerKerasUDF("my_object_recognition_function",  
                 keras_model_file="/mymodels/007model.h5")
```

In Spark SQL:

```
select image, my_object_recognition_function(image) as objects  
from traffic_imgs
```

This means you can apply deep learning models in streaming!

Almost done

Deep Learning Pipelines : Future

In progress 

- Scala API for DeepImageFeaturizer
- Text featurization (embeddings)
- TFTransformer for arbitrary vectors

Future

- Distributed training
- Support for more backends, e.g. MXNet, PyTorch, BigDL

Deep Learning without Deep Pockets

- Simple API for Deep Learning, integrated with MLlib
- Scales common tasks with transformers and estimators
- Embeds Deep Learning models in MLlib and SparkSQL
- Check out <https://github.com/databricks/spark-deep-learning> !

Thank you!

Questions?

Resources

Blog posts & webinars (<http://databricks.com/blog>)

- [Deep Learning Pipelines](#)
- [GPU acceleration in Databricks](#)
- [BigDL on Databricks](#)
- [Deep Learning and Apache Spark](#)

Docs for Deep Learning on Databricks (<http://docs.databricks.com>)

- [Getting started](#)
- [Deep Learning Pipelines Example](#)
- [Spark integration](#)