



# Deep Learning and Streaming in Apache Spark 2.x

Matei Zaharia  
@matei\_zaharia



# Welcome to Spark Summit Europe

Our largest European summit yet

102

talks

1200

attendees

11

tracks

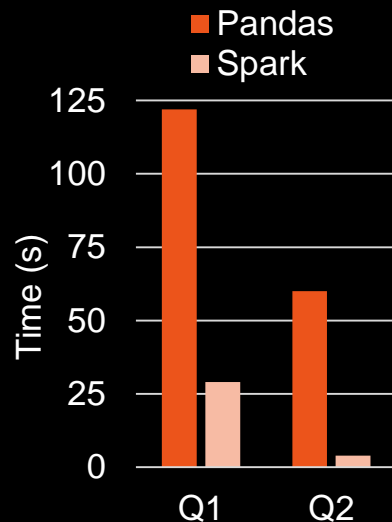
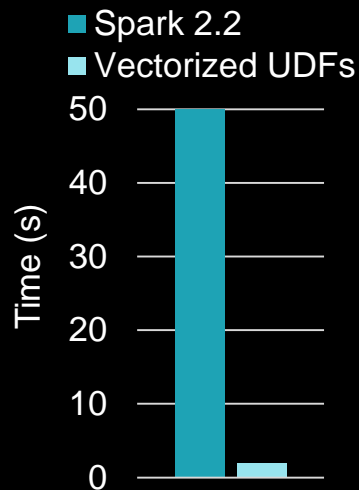
# What's New in Spark?

Cost-based optimizer (Spark 2.2)

Python and R improvements

- PyPI & CRAN packages (Spark 2.2)
- Python ML plugins (Spark 2.3)
- Vectorized Pandas UDFs (Spark 2.3)

Kubernetes support (targeting 2.3)

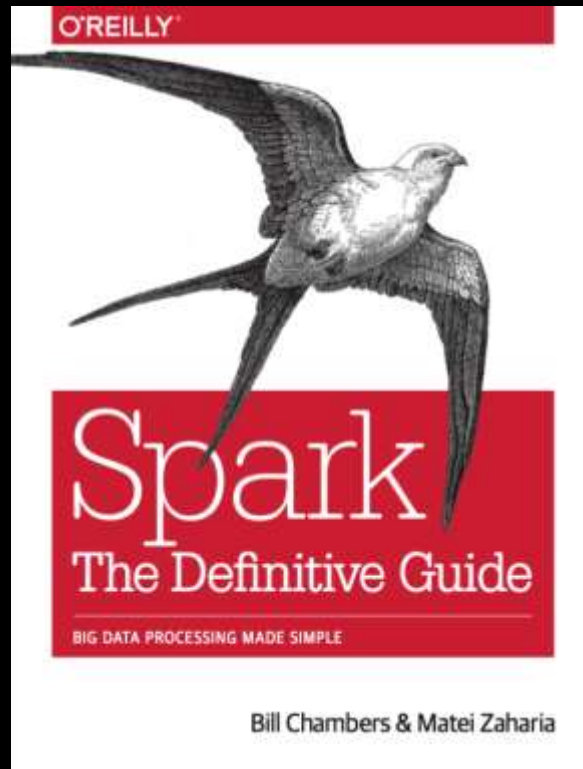


# Spark: The Definitive Guide

To be released this winter

Free preview chapters and  
code on Databricks website:

[dbricks.co/spark-guide](https://dbricks.co/spark-guide)



# Two Fast-Growing Workloads

## Streaming & Deep Learning

Both are important but **complex** with current tools

We think we can **simplify** both with Apache Spark!

# Why are Streaming and DL Hard?

Similar to early big data tools!

Tremendous potential, but very hard to use at first:

- **Low-level APIs** (MapReduce)
- **Separate systems** for each task (SQL, ETL, ML, etc)



# Spark's Approach

## 1) Composable, high-level APIs

- Build apps from components

## 2) Unified engine

- Run complete, end-to-end apps



# Expanding Spark to New Areas

**1 Structured Streaming**

**2 Deep Learning**



# Structured Streaming

Streaming today requires separate APIs & systems

Structured Streaming is a high-level, end-to-end API

- **Simple interface:** run any DataFrame or SQL code incrementally
- **Complete apps:** combine with batch & interactive queries
- **End-to-end reliability:** exactly-once processing

Became GA in Apache Spark 2.2

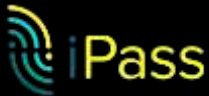
# Structured Streaming Use Cases

100s of customer apps in production on Databricks

Largest apps process tens of trillions of records per month



Monitor quality of live video streaming



Anomaly detection on millions of WiFi hotspots



Real-time game analytics at scale

# Example: YAHOO!

## Benchmark

```
KStream<String, ProjectedEvent> filtered = kEvents.filter((key, value) -> {  
    return value.event_type.equals("view");  
}).mapValues((value) -> {  
    return new ProjectedEvent(value.ad_id, value.event_time);  
});
```

```
KTable<String, String> kCampaigns = builder.table("campaigns", "cmp-state");  
KTable<String, CampaignAd> deserCampaigns = kCampaigns.mapValues((value) -> {  
    Map<String, String> campMap = Json.parser.readValue(value);  
    return new CampaignAd(campMap.get("ad_id"), campMap.get("campaign_id"));  
});  
KStream<String, String> joined =  
    filtered.join(deserCampaigns, (value1, value2) -> {  
        return value2.campaign_id;  
    },  
    Serdes.String(), Serdes.serdeFrom(new ProjectedEventSerializer(),  
        new ProjectedEventDeserializer()));
```

```
KStream<String, String> keyedData = joined.selectKey((key, value) -> value);  
KTable<Windowed<String>, Long> counts = keyedData.groupByKey()  
    .count(TimeWindows.of(10000), "time-windows");
```

APACHE  **Spark**™ DataFrames

```
events  
    .where("event_type = 'view'")  
    .join(table("campaigns"), "ad_id")  
    .groupBy(  
        window('event_time', "10 seconds"),  
        'campaign_id')  
    .count()
```

Batch Plan

Scan Files

Aggregate

Write to Sink

automatic  
transformation

Incremental Plan

Scan New  
Files

Stateful Agg.

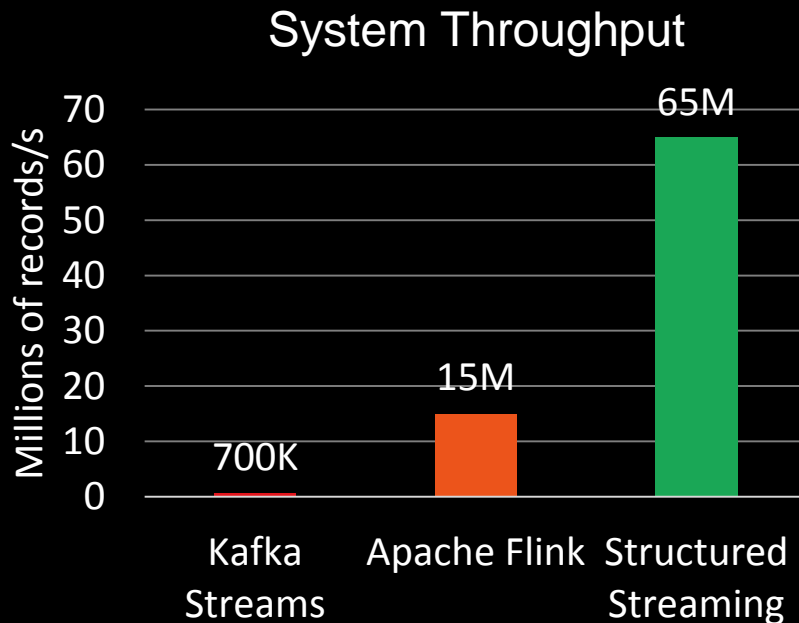
Update Sink

# Performance: YAHOO!

## Benchmark

Structured Streaming  
reuses the **Spark SQL  
Optimizer** and **Tungsten  
Engine**.

**4x**  
fewer nodes



# What About Latency?

**Continuous processing** mode to run without microbatches

- **<1 ms latency** (same as per-record streaming systems)
- No changes to user code
- Proposal in [SPARK-20928](#)

**Key idea:** same API can target both streaming & batch

**Find out more in today's deep dive**

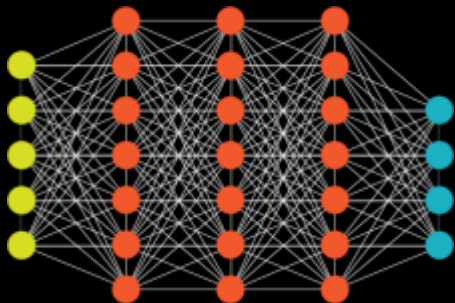
# Expanding Spark to New Areas

① **Structured Streaming**

② **Deep Learning**

# Deep Learning has Huge Potential

Unprecedented ability to work with unstructured data such as images and text



# But Deep Learning is Hard to Use

Current APIs (TensorFlow, Keras, etc) are low-level

- Build a computation graph from scratch

Scale-out requires manual parallelization

Hard to use models in larger applications

Very similar to early big data APIs



# Deep Learning on Spark

Image support in MLlib: SPARK-21866 (Spark 2.3)

DL framework integrations: TensorFlowOnSpark, MMLSpark, Intel BigDL

Higher-level APIs: Deep Learning Pipelines

# New in TensorFlowOnSpark


Library to run distributed TF on Spark clusters & data

- Built at Yahoo!, where it powers photos, videos & more

Yahoo! and Databricks collaborated to add:

- ML pipeline APIs
- Support for non-YARN and AWS clusters

[github.com/yahoo/TensorFlowOnSpark](https://github.com/yahoo/TensorFlowOnSpark)



talk  
tomorrow  
at 17:00

# Deep Learning Pipelines

Low-level DL frameworks are powerful, but common use cases should be much simpler to build

**Goal:** Enable an **order of magnitude** more users to build **production** apps using deep learning

# Deep Learning Pipelines

**Key idea:** High-level API built on ML Pipelines model

- Common use cases are just a few lines of code
- All operators automatically scale over Spark
- Expose models in batch, streaming & SQL apps

Uses existing DL engines (TensorFlow, Keras, etc)

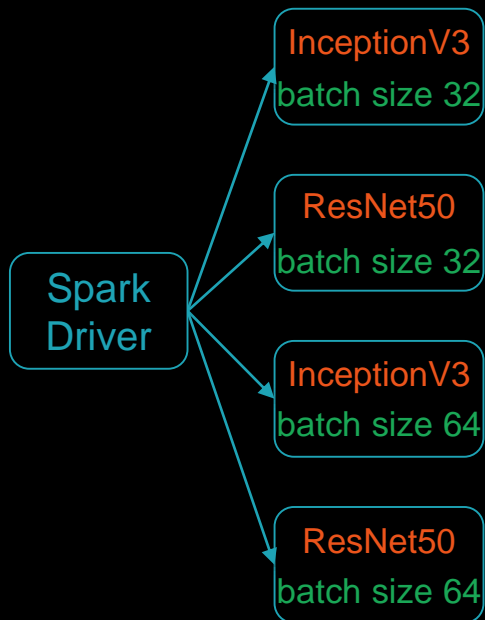
# Example: Using Existing Model

```
predictor = DeepImagePredictor(inputCol="image",  
                                outputCol="labels",  
                                modelName="InceptionV3")
```

```
predictions_df = predictor.transform(image_df)
```

```
SELECT image, my_predictor(image) AS labels  
FROM uploaded_images
```

# Example: Model Search



```
est = KerasImageFileEstimator()
```

```
grid = ParamGridBuilder() \  
  .addGrid(est.modelFile, ["InceptionV3", "ResNet50"]) \  
  .addGrid(est.kerasParams, [{'batch': 32}, {'batch': 64}]) \  
  .build()
```

```
CrossValidator(est, eval, grid).fit(image_df)
```



# Deep Learning Pipelines

## Demo

Sue Ann Hong