



Fast Data with Apache Ignite & Apache Spark

Christos Erotocritou, GridGain Systems

#EUstr10



...is a **distributed, memory-centric** data platform
with powerful & flexible **processing** APIs



Apache Ignite Memory-Centric Data Platform

Applications



Financial
Services




Telco




Travel &
Logistics



E-Commerce



Pharma &
Healthcare



IoT

SQL

Key/Value

Transactions

Compute

IgniteRDD

Streaming

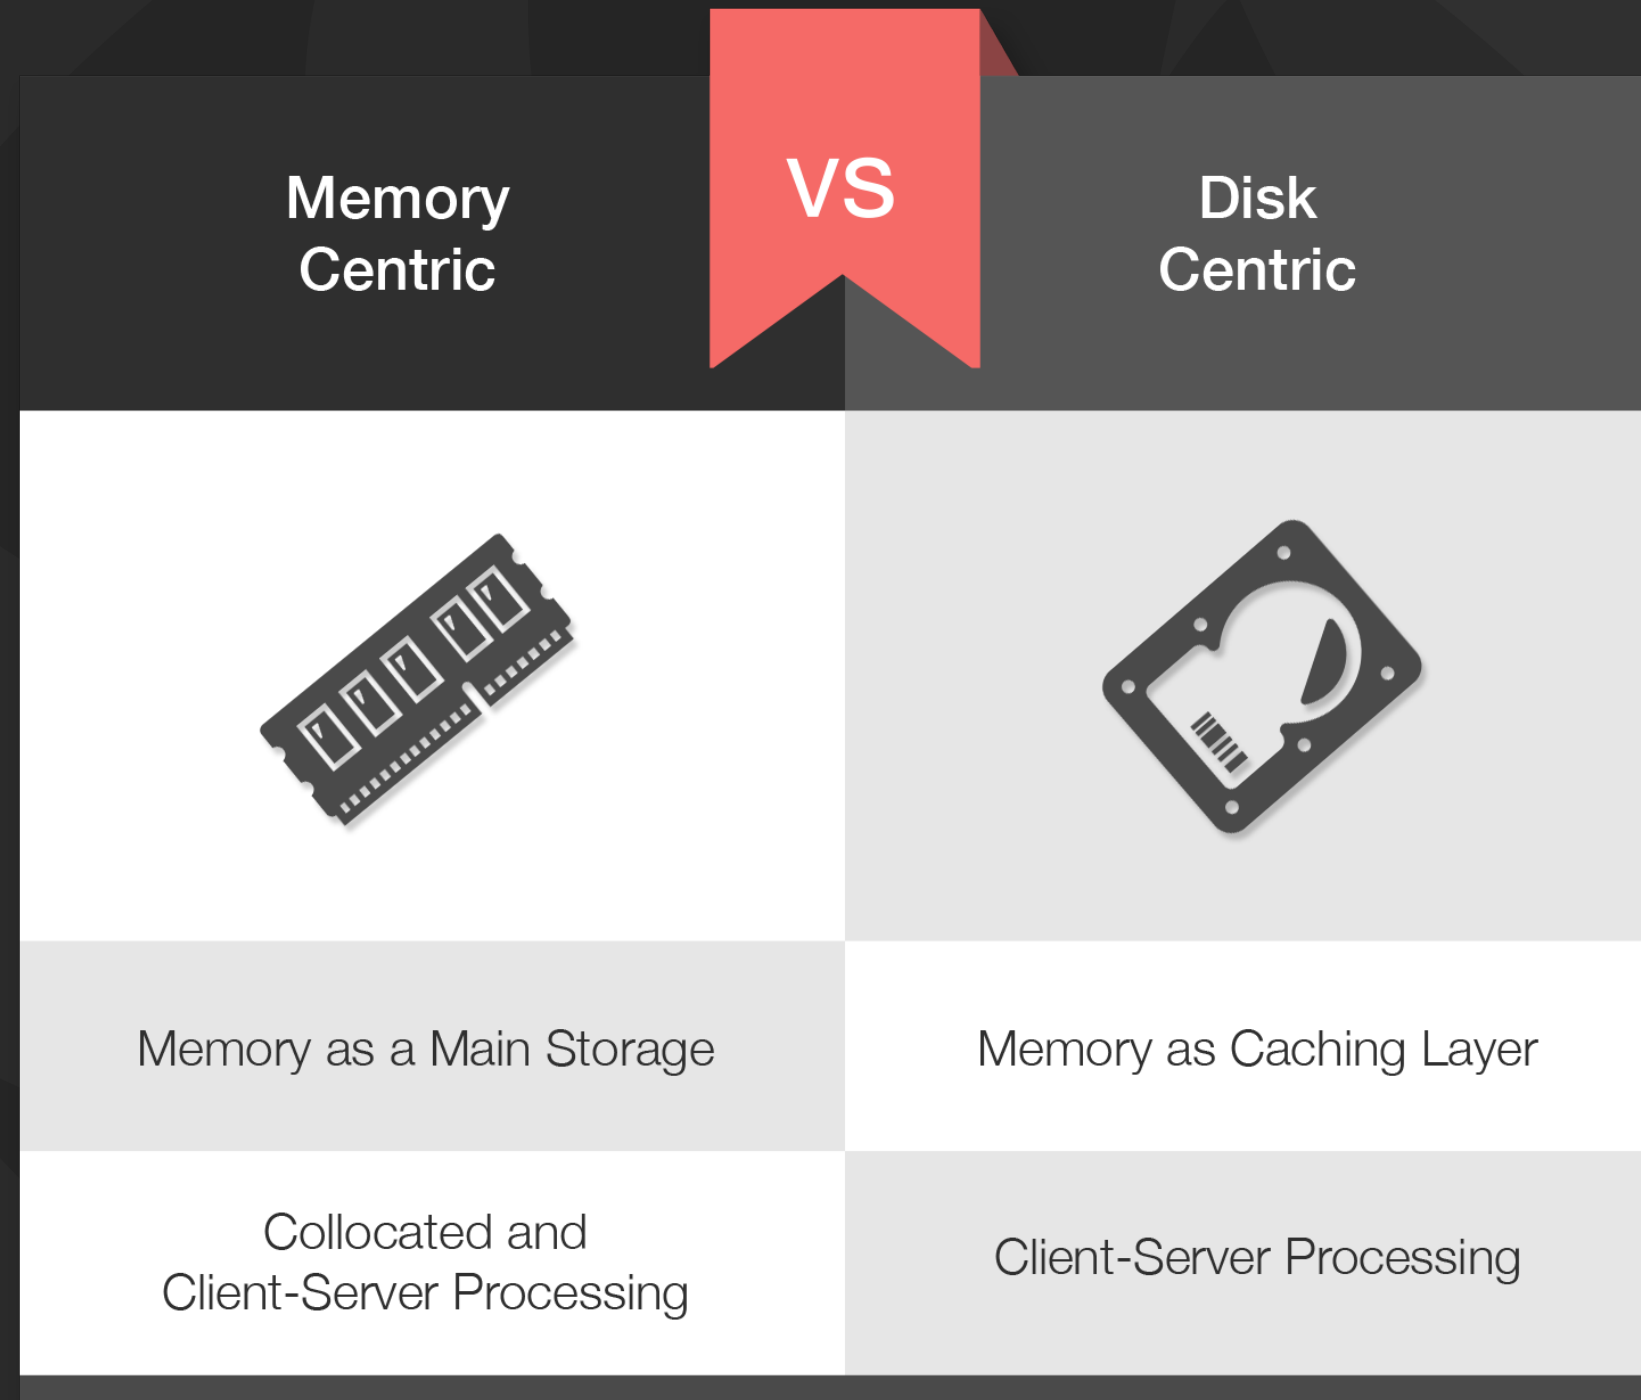
ML

 **Ignite Memory-Centric Storage**

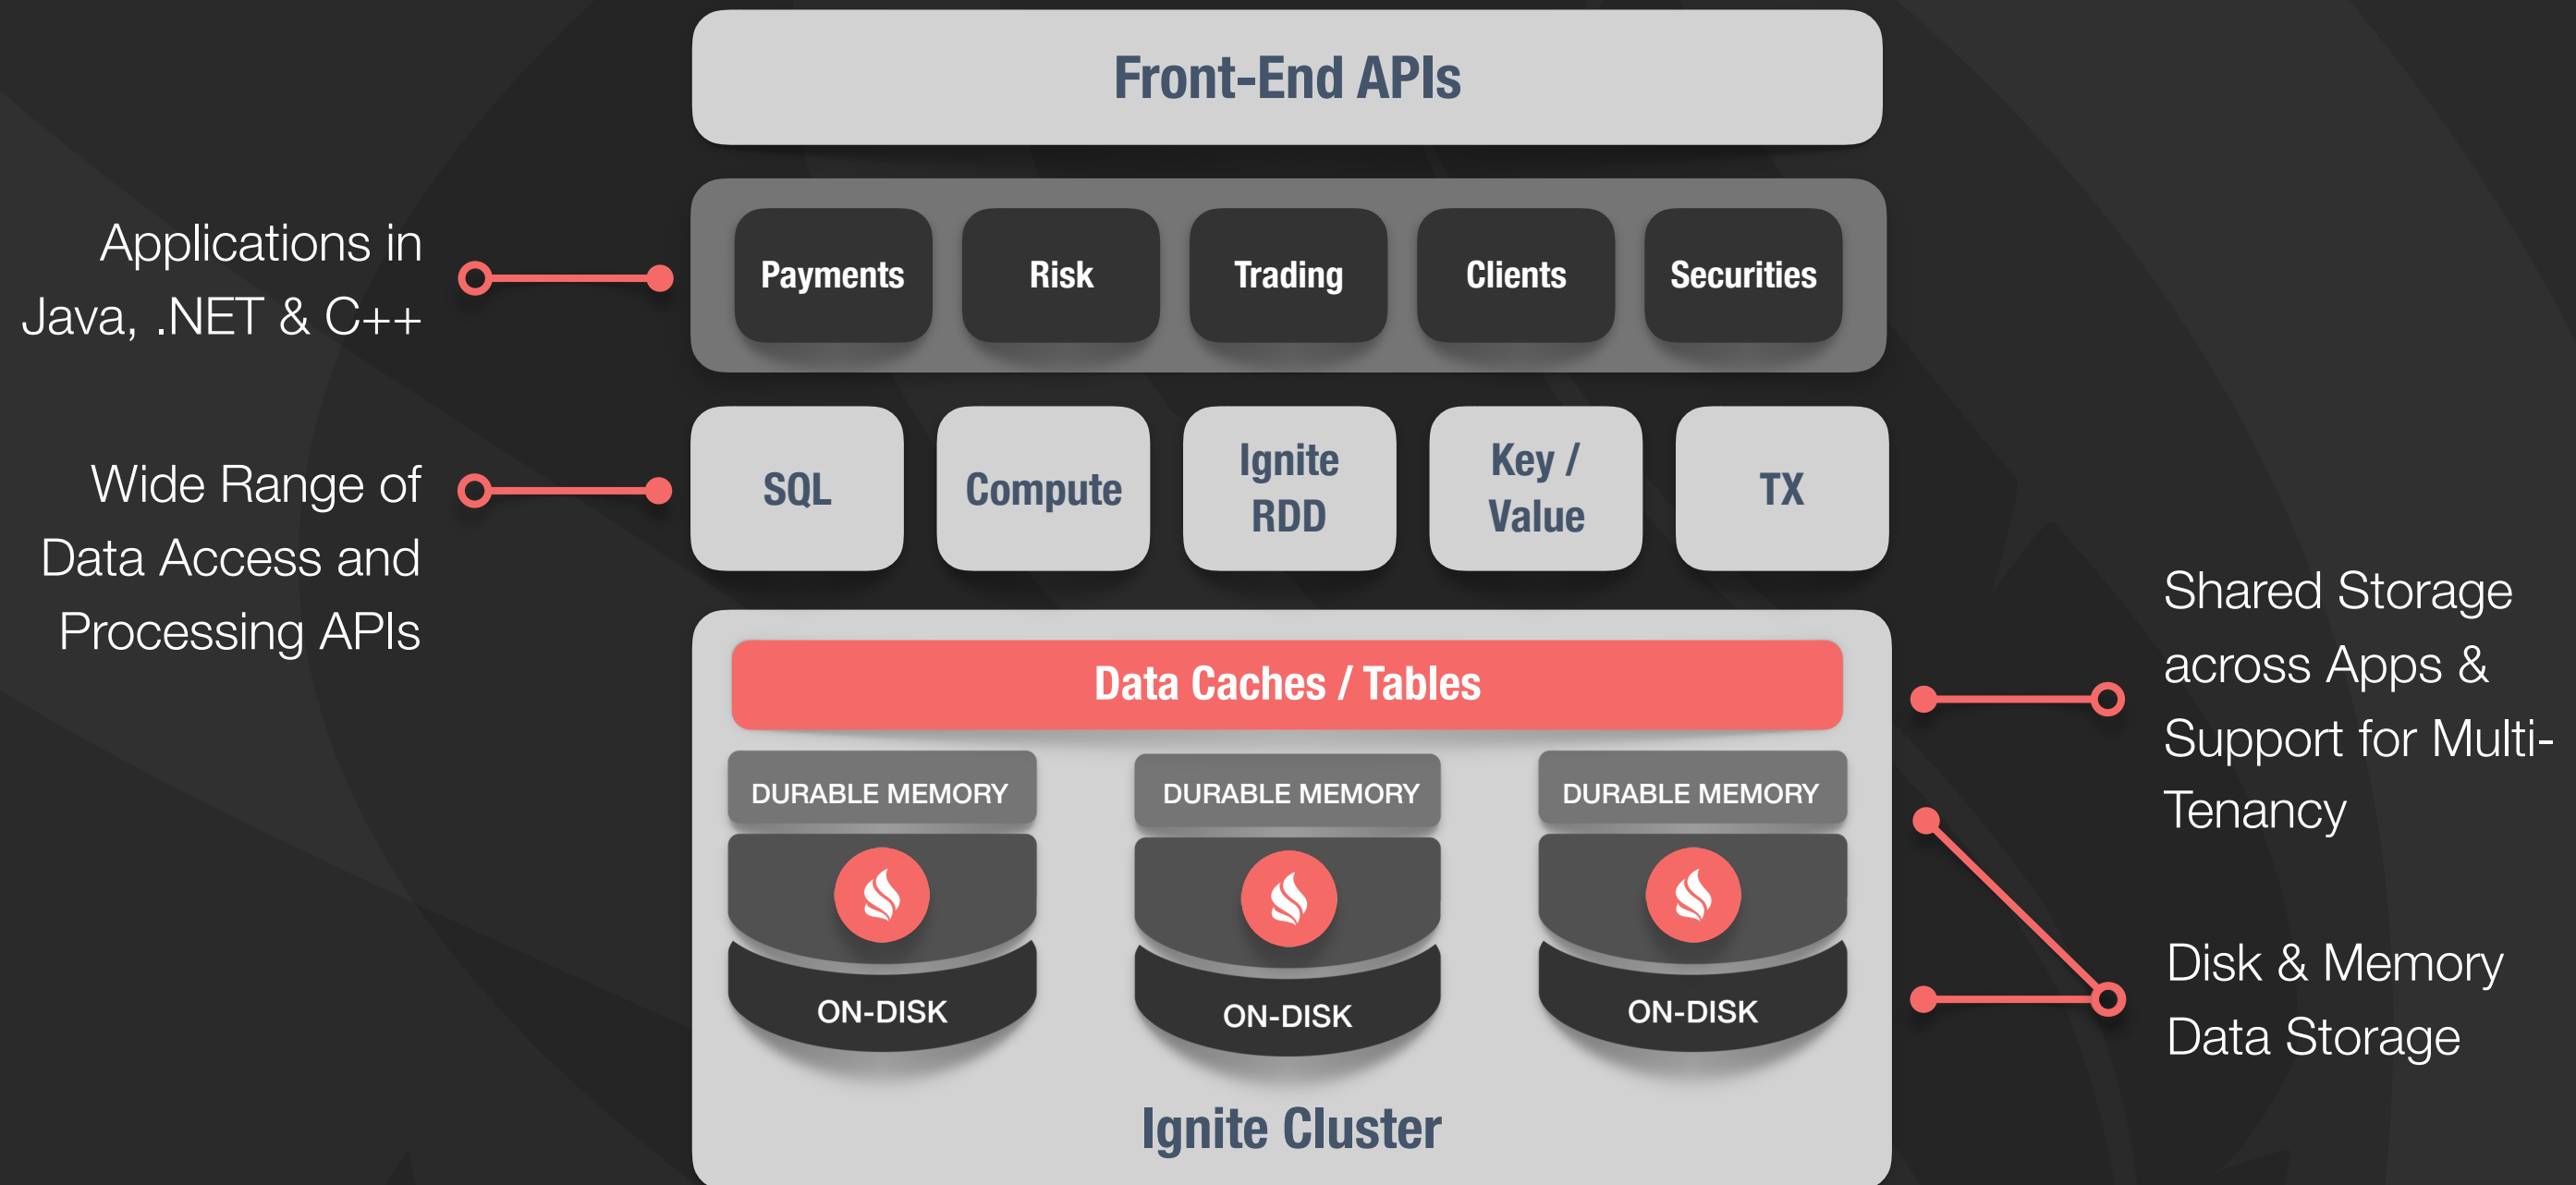
 **Ignite Native Persistence**
(Flash, SSD, Intel 3D XPoint)

Third-Party Persistence
(RDBMS, HDFS, NoSQL)

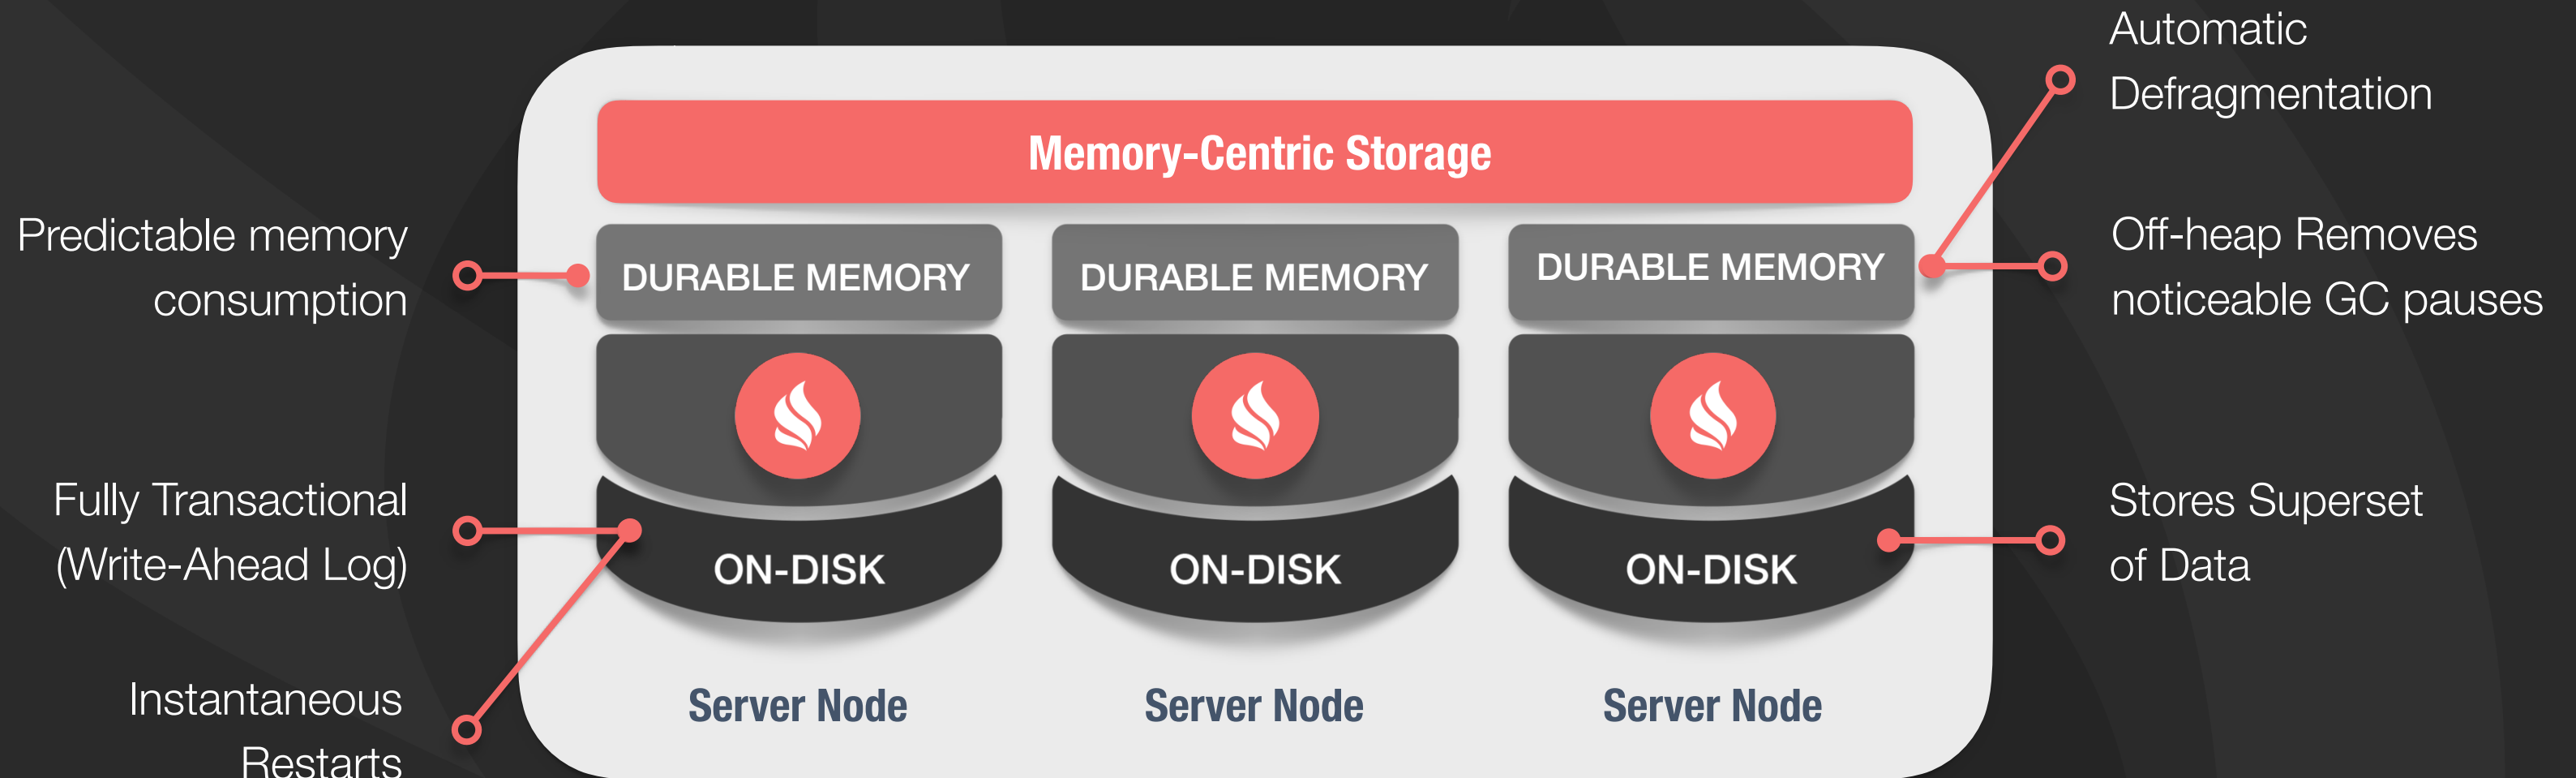
Memory-Centric Storage



Pure Ignite Deployment

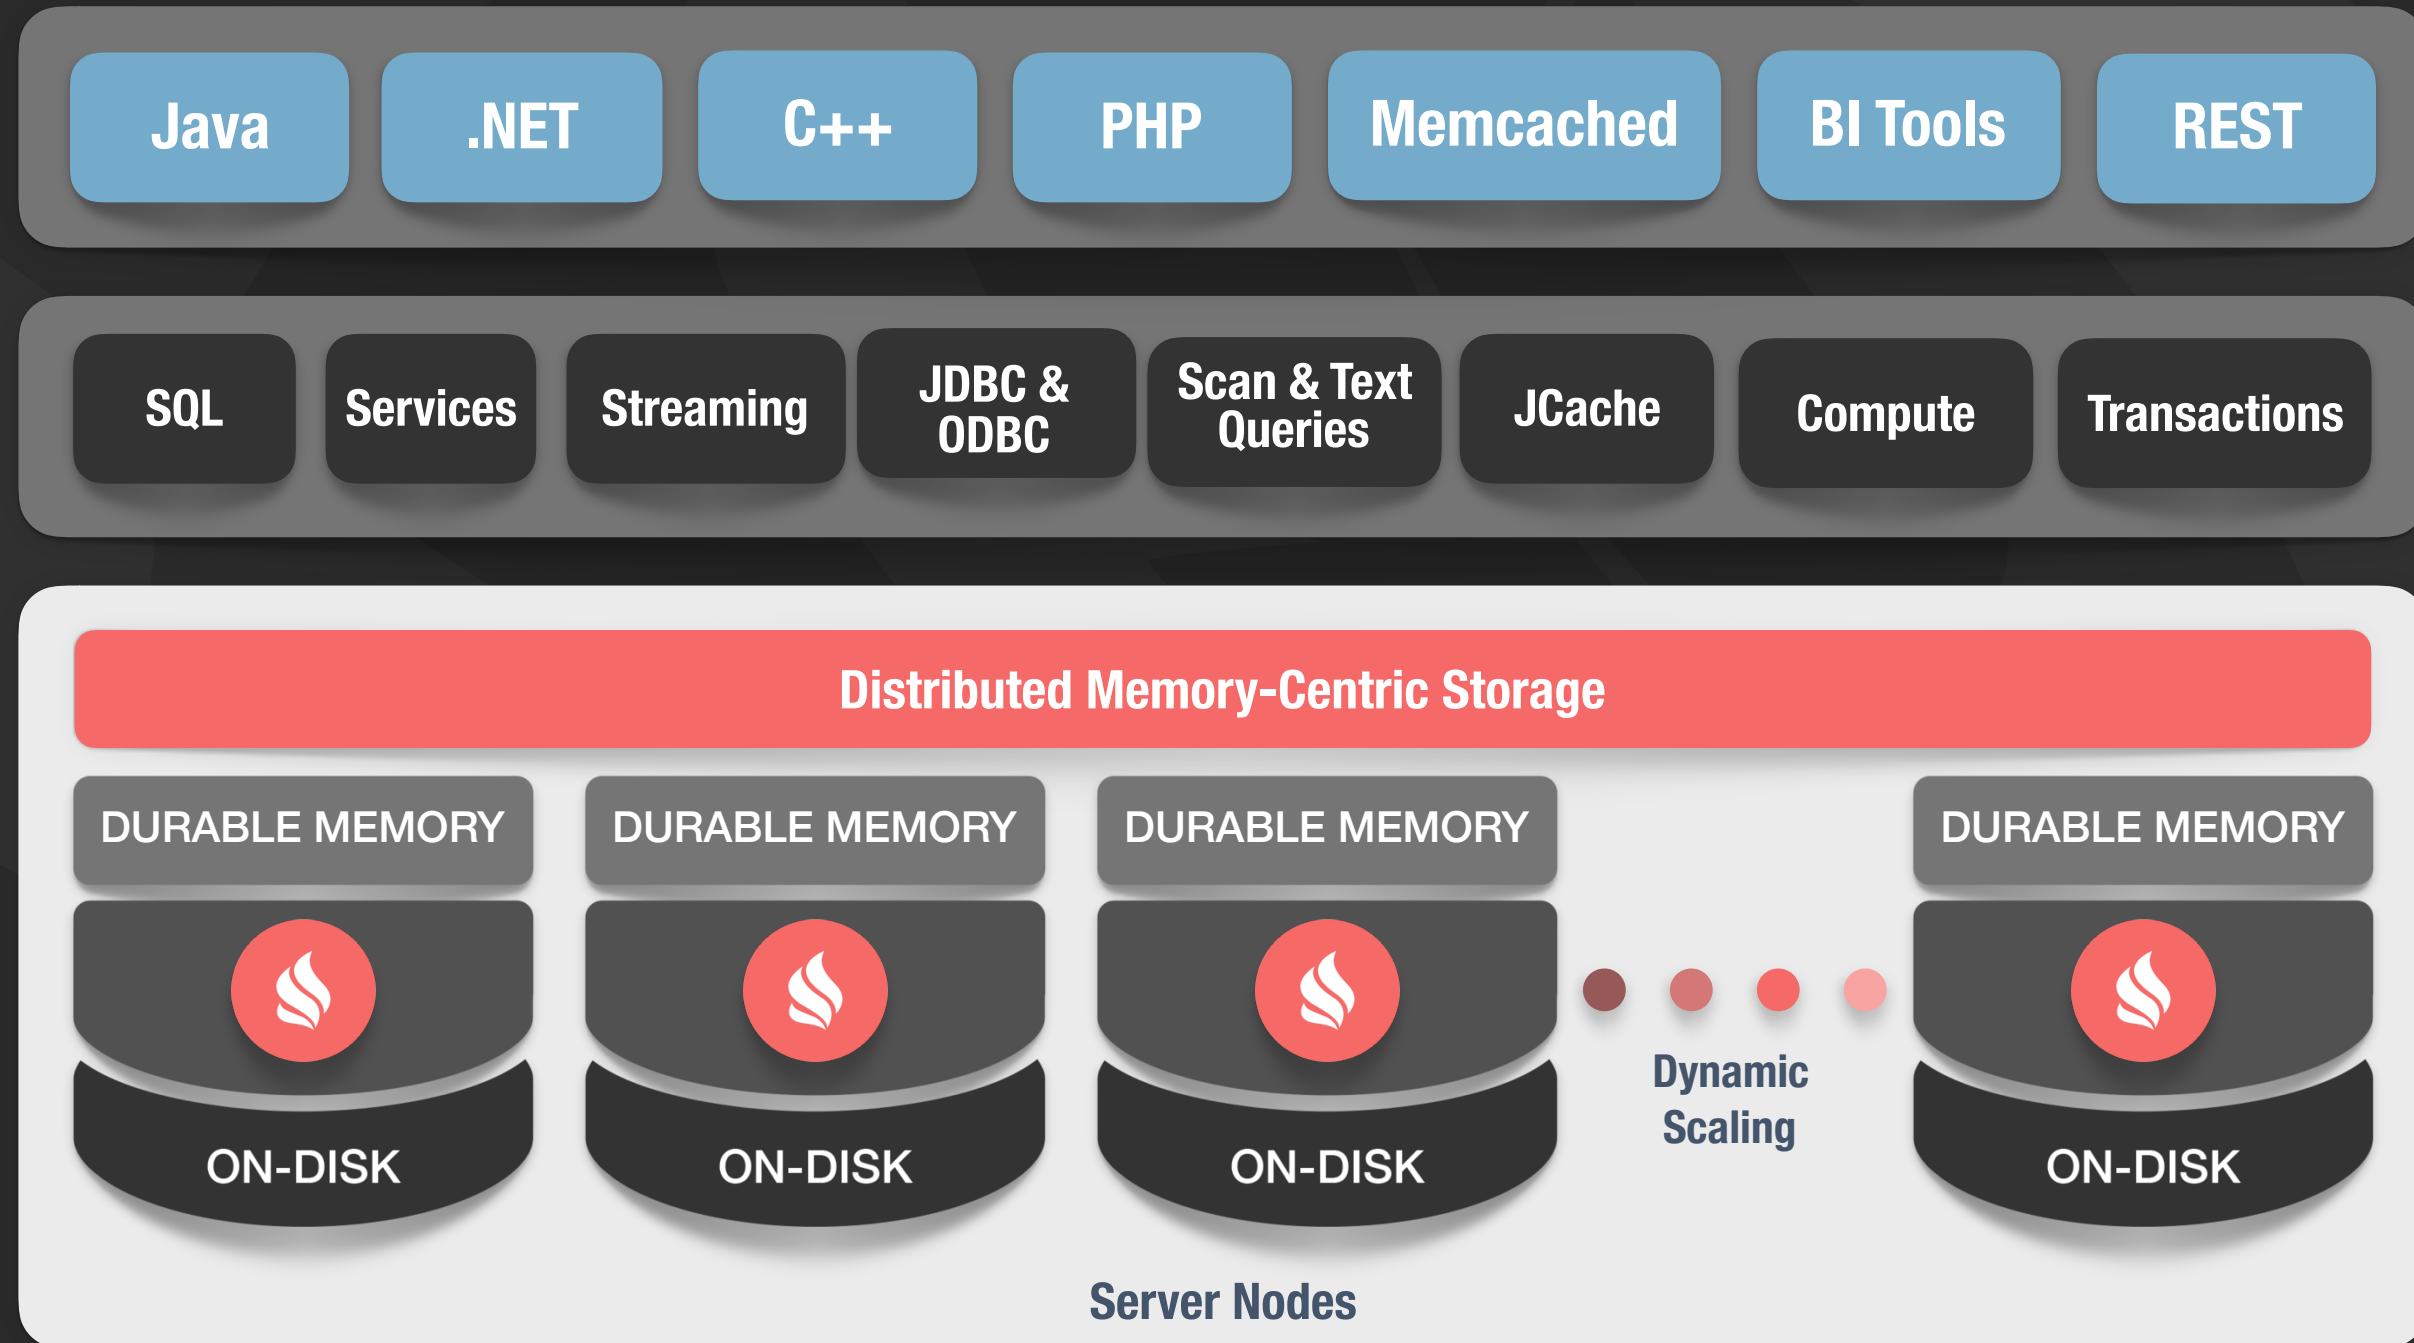


Durable Memory

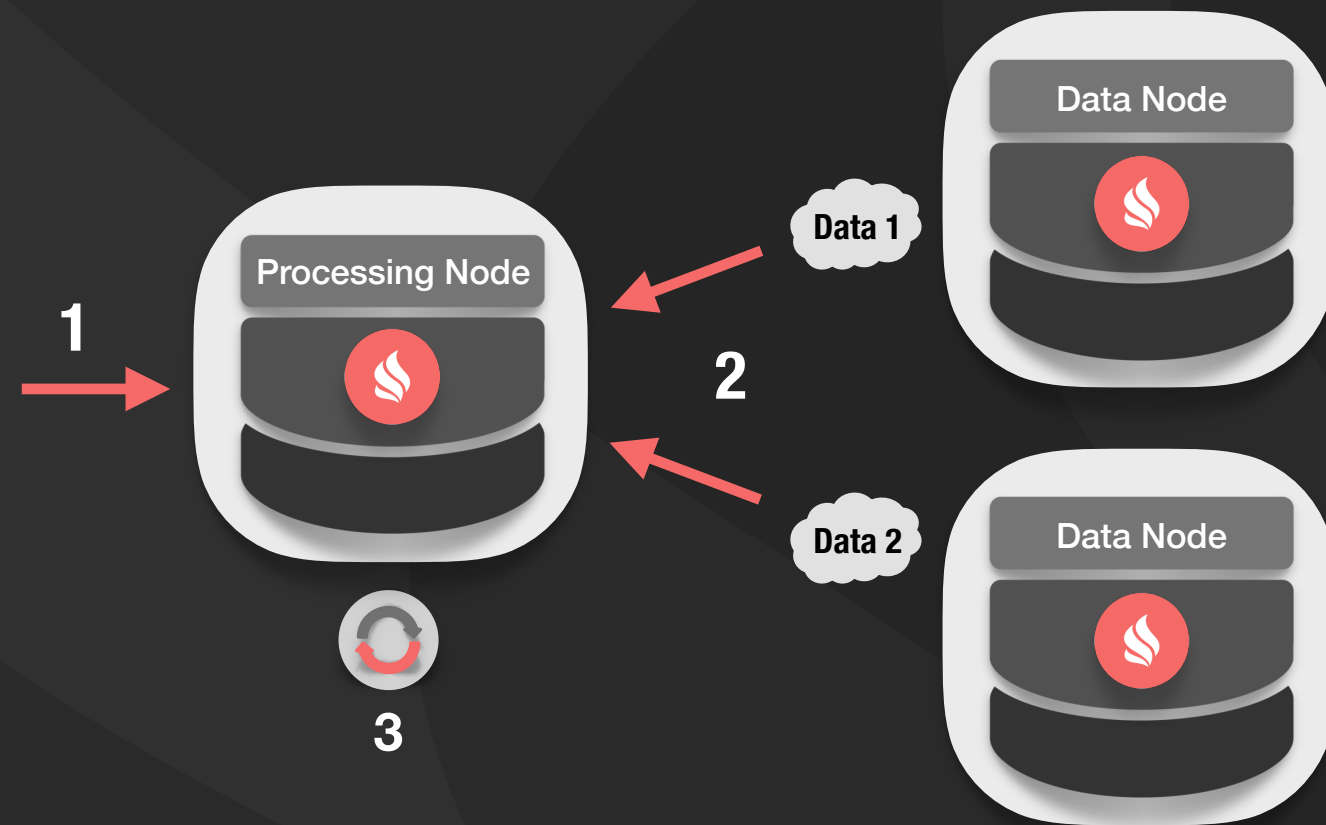


 **Ignite Server Cluster**

Apache Ignite Features




Client-Server Processing



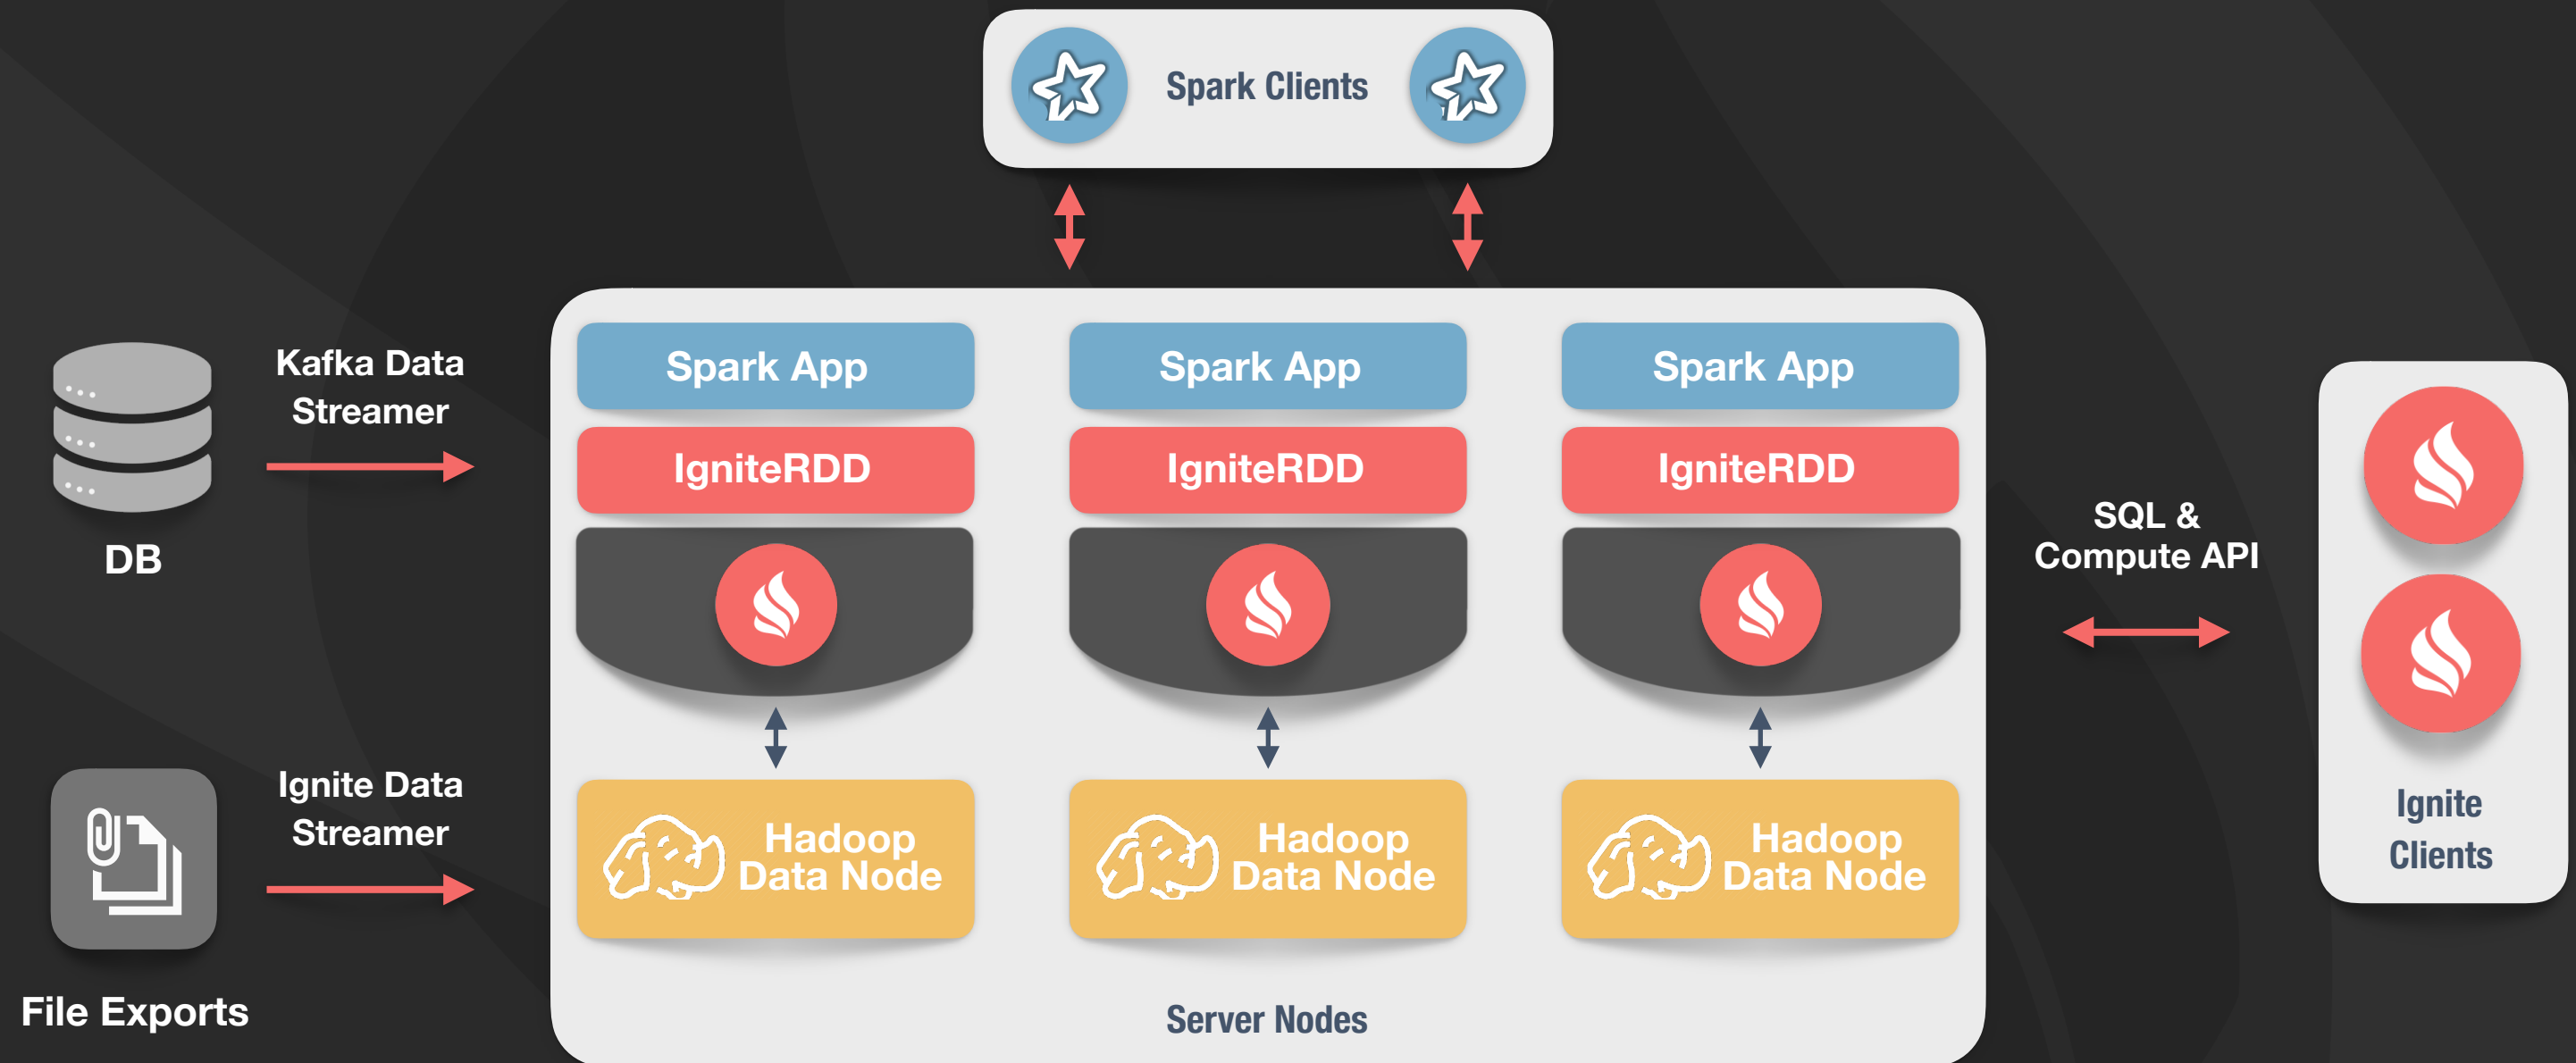
1. Initial Request
2. Fetch data from remote nodes
3. Process the entire data-set

Co-located Processing



1. Initial request
2. Co-locate processing with data
3. Reduce multiple results into one

Hadoop, Spark & Ignite Deployment



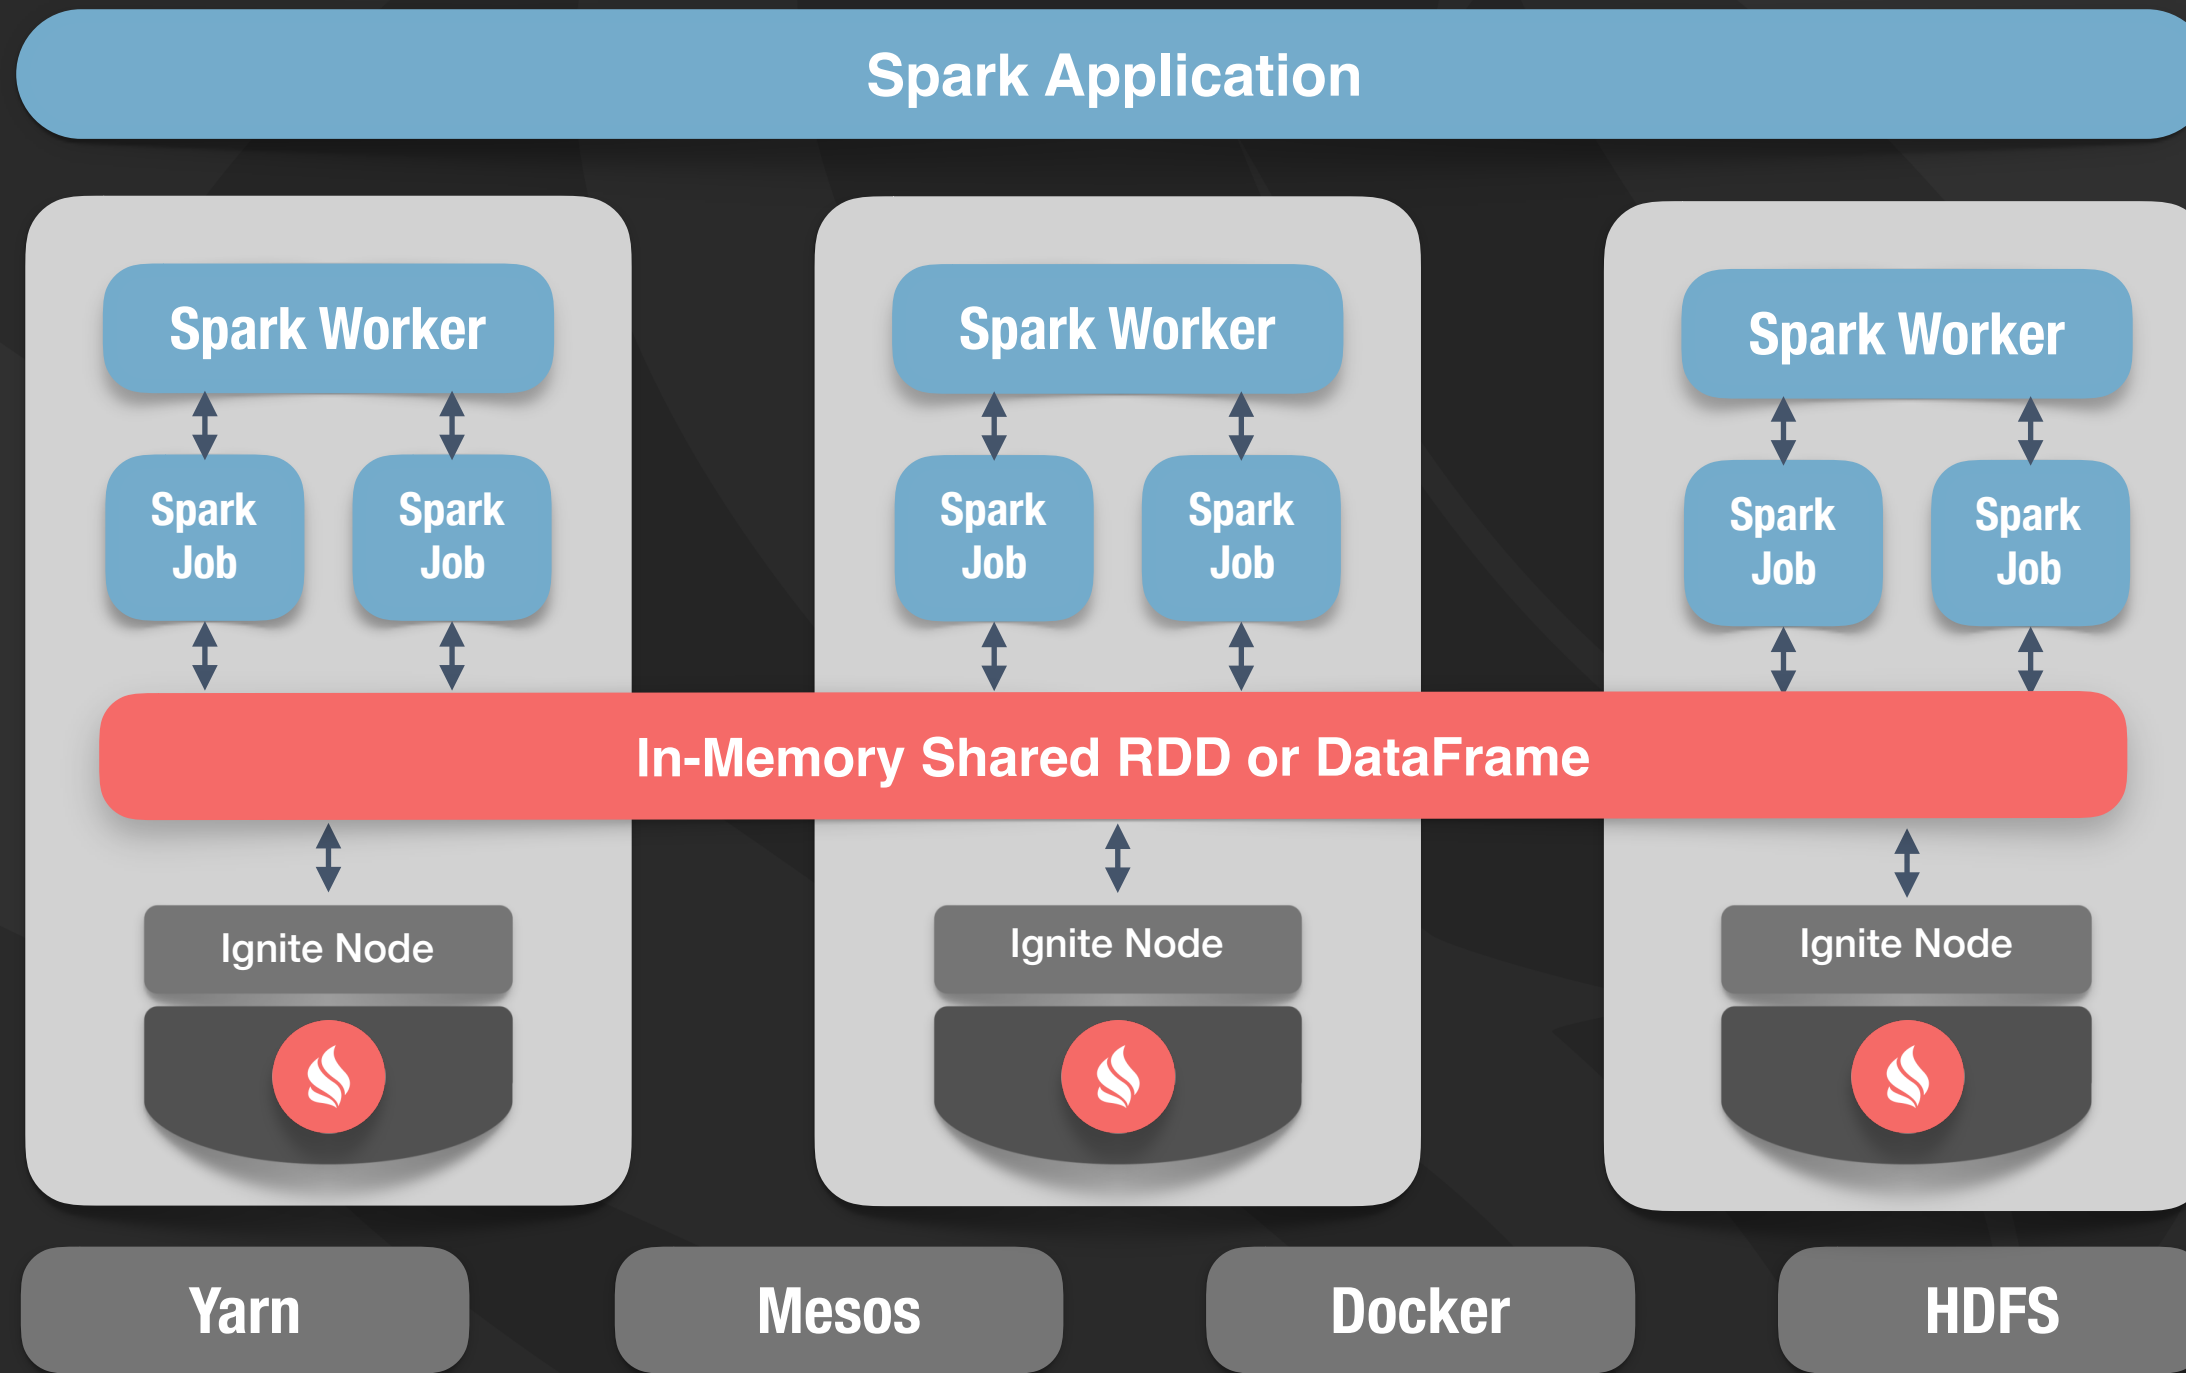
Apache Ignite Spark Integration

Share RDD
across jobs on
the host

Share RDD
Globally

In-Memory
Indexes

SQL on top of
RDDs



Working with IgniteRDD

- IgniteContext is the main entry point to Spark-Ignite integration:

```
val igniteContext = new IgniteContext[Integer, Integer]  
    (sparkContext, () => new IgniteConfiguration())
```

- Reading values from Ignite:

```
val cache = igniteContext.fromCache("myRdd")  
val result = cache.filter(_._2.contains("Ignite")).collect()
```

- Saving values to Ignite:

```
val cacheRdd = igniteContext.fromCache("myRdd")  
cacheRdd.savePairs(sparkContext.parallelize(1 to 10000, 10).map(i => (i, i)))
```

- Running SQL queries against Ignite Cache:

```
val cacheRdd = igniteContext.fromCache("myRdd")  
val result = cacheRdd.sql  
    ("select _val from Integer where val > ? and val < ?", 10, 100)
```

Working with DataFrame API

- Create an IgniteRDD

```
val companyCacheIgnite = new IgniteContext[Int, String](sc, () =>  
new IgniteConfiguration()).fromCache("CompanyCache")
```

- Create a “Company” DataFrame

```
val dfCompany = sqlContext.createDataFrame(companyCacheIgnite.map(p=>  
Company(p._1, p._2)))
```

- Register DataFrame as a table

```
dfCompany.registerTempTable("company")
```



- Data source agnostic
- Fully fledged compute engine and durable storage
- OLAP & OLTP
- Zero-deployment
- In-Memory SQL support
- Fully ACID transactions across memory and disk
- Less focused on Hadoop
- Early ML Support
- Growing Community



- Ingests data from HDFS or another distributed file system
- Inclined towards analytics (OLAP) and focused on MR-specific payloads
- Requires the creation of RDD and data and processing operations are governed by it
- Basic disk-based SQL support
- Strong ML libraries
- Big community

GridGain

- **What is GridGain?**
 - Binary build of Apache Ignite™
 - Added enterprise features for enterprise deployments
 - Earlier features and bug fixes by a few weeks
 - Fully certified & tested releases

“We develop and support the worlds leading In-Memory Computing Platform”



Any Questions?

Thank you for joining us. Follow the conversation.
<http://ignite.apache.org>

