

From pipelines to refineries: scaling big data applications

Tim Hunter

@timjhunter

Spark Summit Dublin 2017



About Me



- Tim Hunter
- Software engineer @ Databricks
- Ph.D. from UC Berkeley in Machine Learning
- Very early Spark user
- Contributor to MLlib
- Author of TensorFrames, GraphFrames, Deep Learning Pipelines

Introduction

- Spark 2.3 in the release process
- Spark 2.4 being thought about
- This is the time to discuss Spark 3
- This presentation is a personal perspective on a future Spark

Introduction

There is nothing more practical than a good theory.

James Maxwell

As Spark applications grow in complexity, what challenges lie ahead?

What are some good foundations for building big data frameworks?




Outline

- State of the union
 - What is good about Spark?
 - What are the trends?
- Classics to the rescue
 - Fighting the four horsemen of the datapocalypse
 - Laziness to the rescue
- From theory to practice
 - Making data processing great again

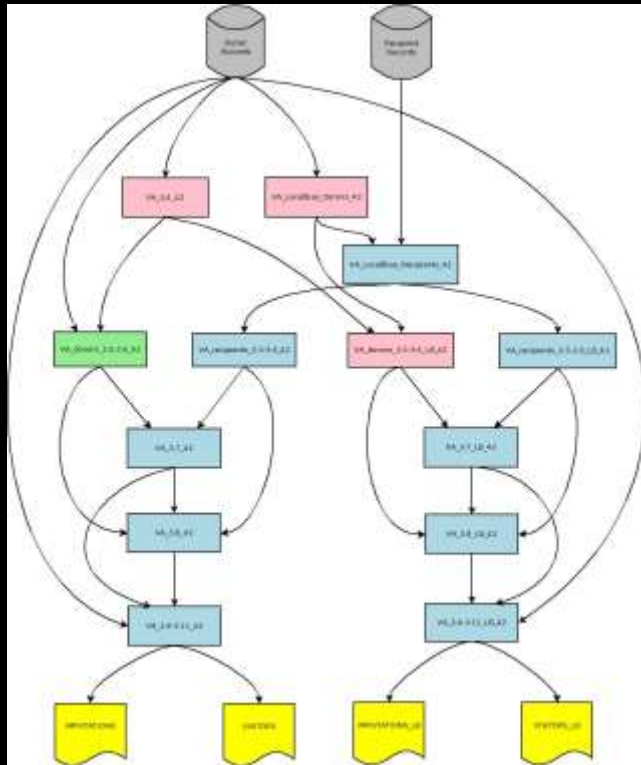
State of the union

- What we strive for



State of the union

- What we deal with:
 - Coordinating a few tasks





State of the union

- The (rapidly approaching) future
 - Hundreds of input sources
 - Thousands of concurrent requests
 - Mixing interactive, batch, streaming
- How do we enable this?



The state of the union

- The image of a pipeline gives you the illusion of simplicity
 - One input and one output
- Current big data systems: the tree paradigm
 - Combine multiple inputs into a single output
 - The SQL paradigm
 - Followed by Spark
- A forest is more than a group of trees
 - Multiple inputs, multiple outputs
 - The DAG paradigm



The ideal big data processing system:

- *Scalability*
 - in quantity (big data) and diversity (lots of sources)
- *Chaining*
 - express the dependencies between the datasets
- *Composition*
 - assemble more complex programs out of simpler ones
- *Determinism*
 - given a set of input data, the output should be unique*

How is Spark faring so far?

- You *can* do it, but it is not easy

What can go wrong with this program?

```
all_clicks = session.read.json("/tables/clicks/year=2017")
all_clicks.cache()
max_session_duration = all_clicks("session_duration").max()
top_sessions = all_clicks.filter(
    all_clicks("session_duration") >= 0.9 * max_session_duration)
top_ad_served = top_sessions("ad_id")
top_ad_served.write.parquet("/output_tables/top_ads")
```

leak

↓ a few hours...

typo

missing directory

The 4 horsemen of the datapocalypse

- Typing (schema) mismatch
- Missing source or sink
- Resource leak
- Eager evaluation

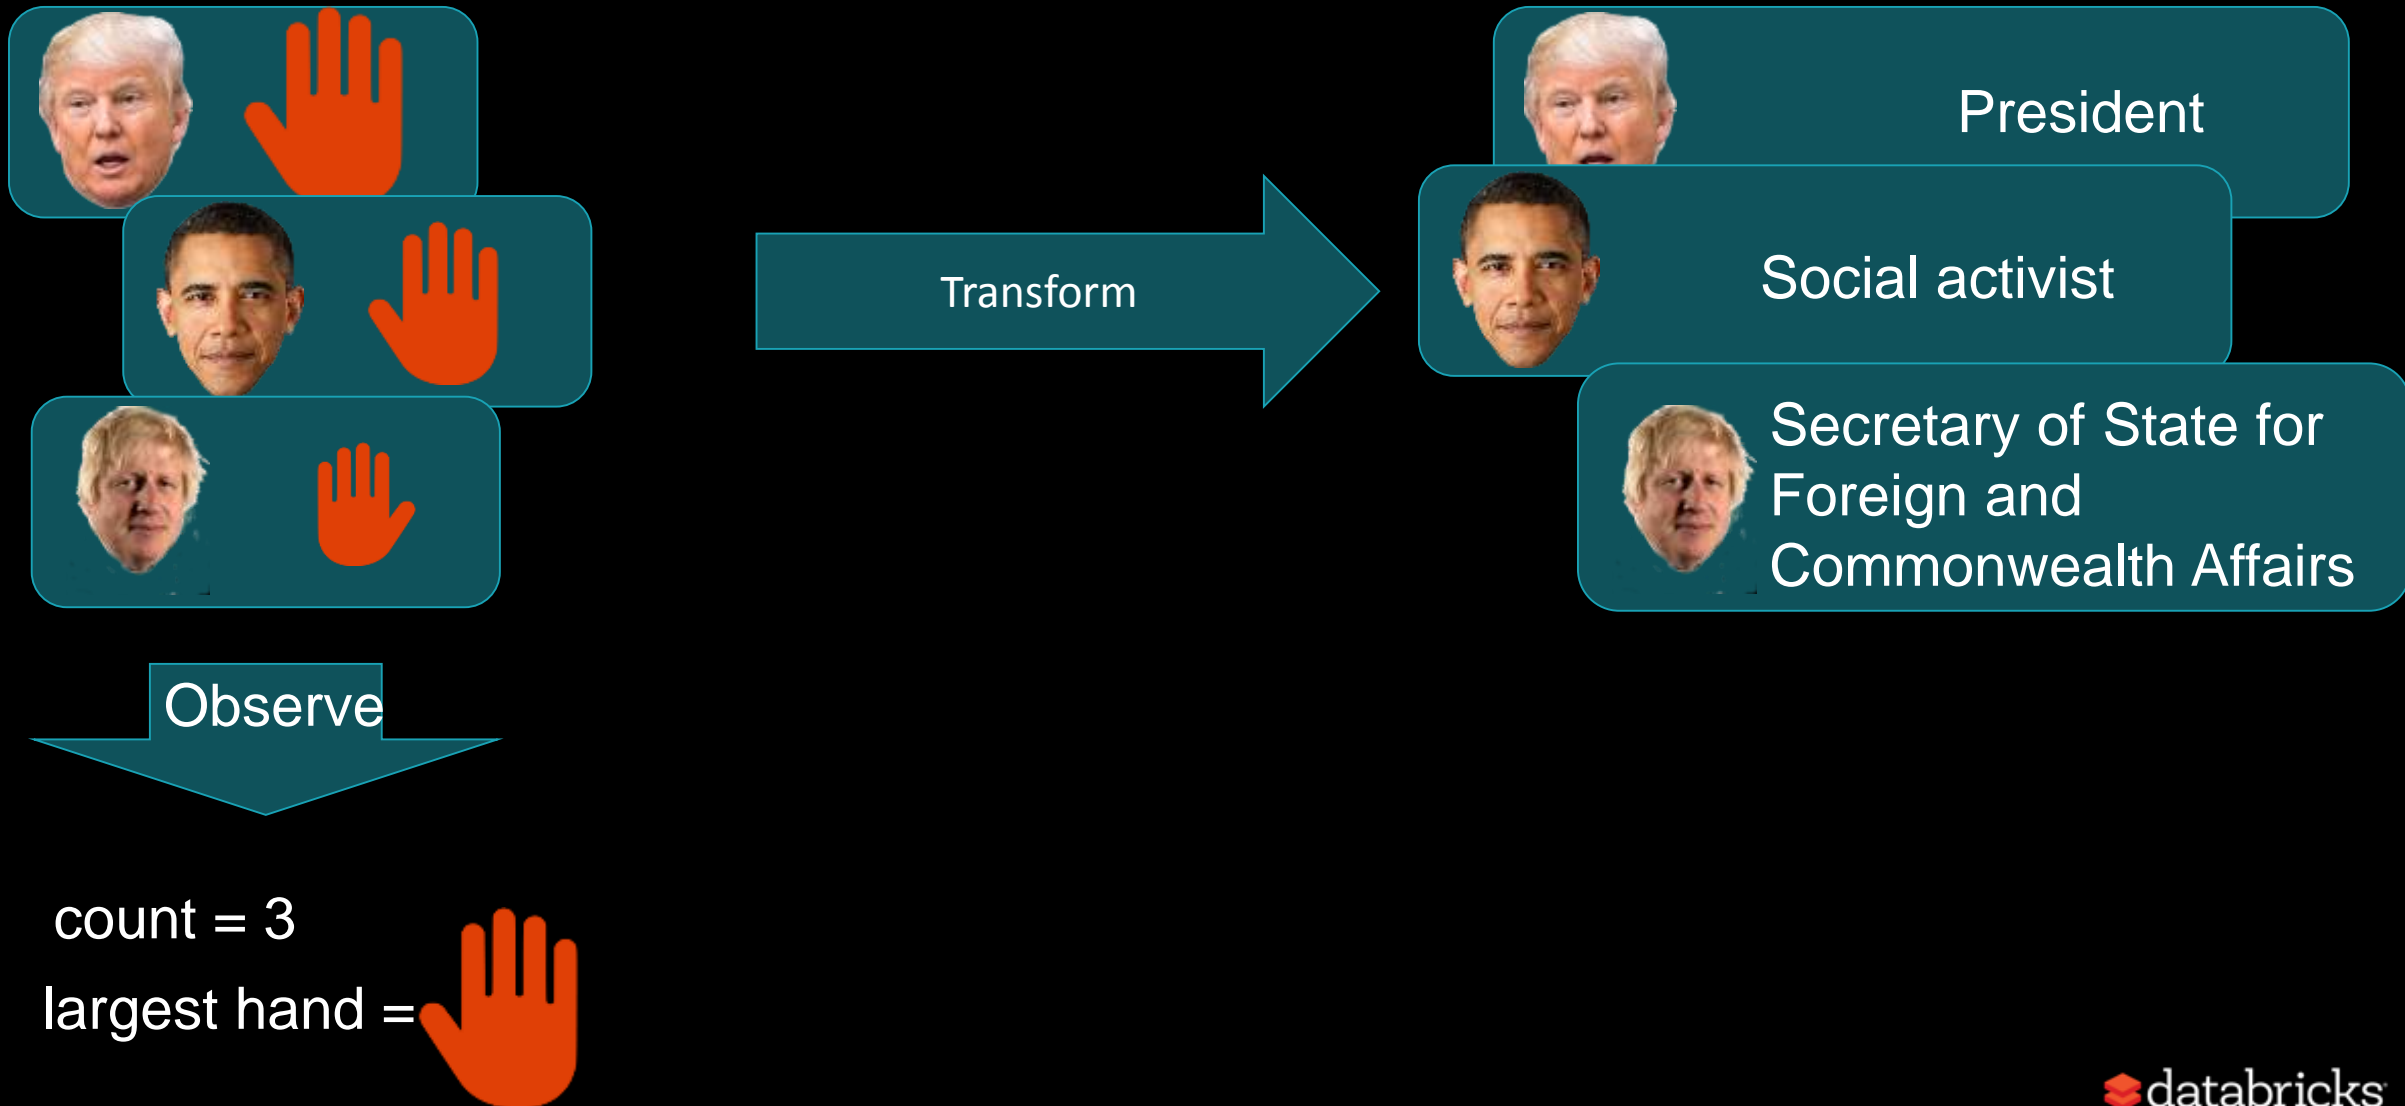


Classics to the rescue

Theoretical foundations for a data system

- A dataset is a collection of elements, all of the same type
 - Scala: `Dataset[T]`
- Principle: the content of a dataset cannot be accessed directly
 - A dataset can be *queried*
- An observable is a single element, with a type
 - intuition: dataset with a single row
 - Scala: `Observable[T]`

Theoretical foundations for a data system



Theoretical foundations for a data system

- *Principle*: the observation only depends on the content of the dataset
 - You cannot observe partitions, ordering of elements, location on disk, etc.
- Mathematical consequence: all reduction operations on datasets are monoids:
 - $f(A \cup B) = f(A) + f(B) = f(B) + f(A)$
 - $f(\text{empty}) = 0$

Theoretical foundations for a data system

- *Principle*: closed world assumption
 - All the effects are modeled within the framework
 - The inputs and the transforms are sufficient to generate the outputs
- Practical consequence: strong checks and sci-fi optimizations

Examples of operations

- They are what you expect:
 - `Dataset[Int]` : a dataset of integers
 - `Observable[Int]` : an observation on a dataset
- `max`: `Dataset[Int] => Observable[Int]`
- `collect`: `Dataset[Int] => Observable[List[Int]]`

Karps

- An implementation of these principles on top of Spark
- It outputs a graph of logical plans for Spark (or other systems)
- Makes a number of correctness checks for your program
- *Automatically converts (a subset of) Pandas programs to Spark.*

Demo 1

Enabling complex data programs

- Lazy construction of very complex programs
- Most operations in Spark can be translated to a small set of primitive actions with well-defined composition rules.
- The optimizer can then rewrite the program without changing the outcome
- Optimizations can leverage further SQL optimizations

Demo 2

Future directions

- More complete python (pandas) interface
- I/O in Python
- Finish GroupBy (cool stuff ahead)
- Tracing & Profiling
- SQL (simple and cool stuff to do in this area)

Conclusion: trends in data processing

- How to manage the complexity of data flows?
- Taking inspiration from the functional world
- Spark provides solid foundation
- Laziness, declarative APIs alleviate complexity

Trying this demo

- <https://github.com/tjhunter/karps>
- Notebooks:
 - <https://github.com/tjhunter/karps/tree/master/notebooks>



Thank You

Dealing with In and Out

- The only type of I/O: read and write datasets
- This is an observable
- Operations are deterministic + results are cached
 - -> only recompute when the data changes
- Demo

Example: Caching

