



Indicium: Interactive Querying at Scale

Arvind Heda, Kapil Malik

#EUeco9

What's in the session ...

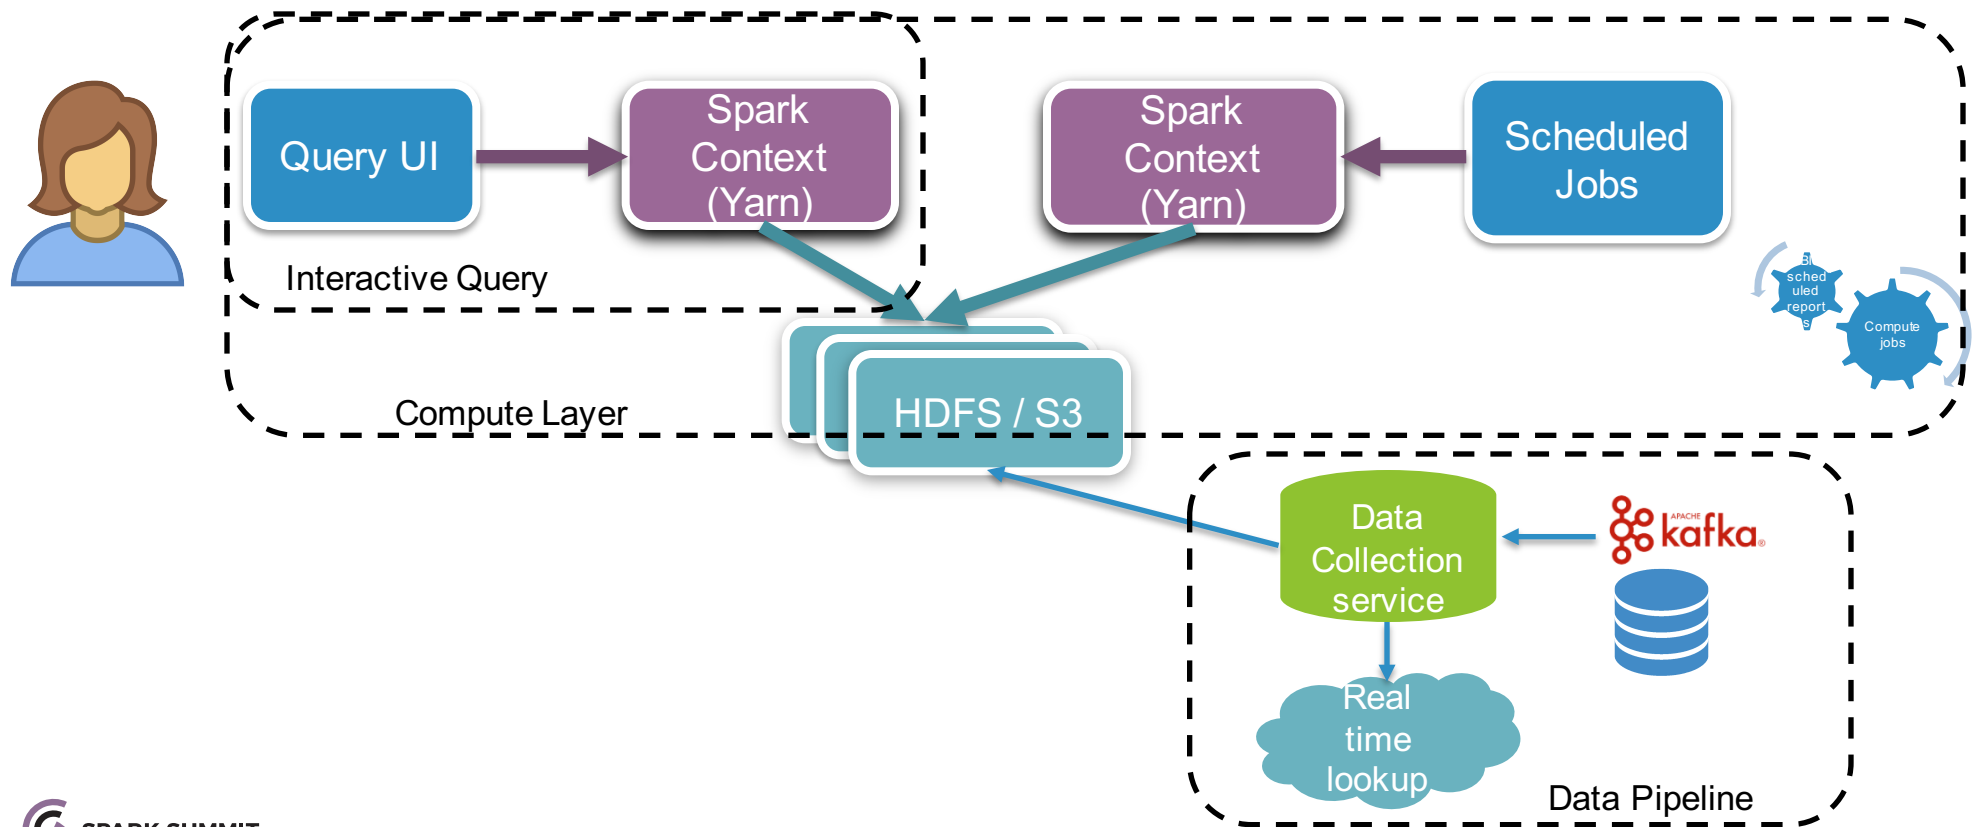
- Unified Data Platform on Spark
 - Single data source for all scheduled / ad hoc jobs and interactive lookup / queries
 - Data Pipeline
 - Compute Layer
 - Interactive Queries?
- Indicium: Part 1 (managed context pool)
- Indicium: Part 2 (smart query scheduler)

Unified Data Platform

Unified Data Platform...(for anything / everything)

- Common Data Lake for storing
 - Transactional data
 - Behavioral data
 - Computed data
- Drives all decisions / recommendations / reporting / analysis from same store.
- Single data source for all Decision Edges, Algorithms, BI tools and Ad Hoc and interactive Query / Analysis tools
- Data Platform needs to support
 - Scale – Store everything from summary to raw data.
 - Concurrency – Handle multiple requests in acceptable user response time.
 - Ad Hoc Drill down to any level – query, join, correlation on any dimension.

Unified Data Platform



Features

| Features | Details | Approach |
|-----------------------------|---|---|
| Data Persistence | Store Large Data Volume of Txn, Behavioural and Computed data; | Spark – Parquet format on S3 / HDFS |
| Data Transformations | Transformation / Aggregation – co relations and enrichments | Batch Processing - Kafka / Java / Spark Jobs |
| Algorithmic Access | Aggregated / Raw Data Access for scheduled Algorithms | Spark Processes with SQL Context based data access |
| Decision Making | Aggregated Data Access for decision in real time | In memory cache of aggregated data |
| Reporting BI / Ad Hoc Query | Aggregated / Raw Data Access for scheduled reports (BI) Aggregated / Raw Data Access for Ad Hoc Queries | BI tool with defined scheduled spark SQL queries on Data store; |
| Interactive Queries | Drill down data access on BI tools for concurrent users Ad hoc Query / Analysis on data for concurrent users | Scaling challenges for Spark SQL? |

Data Pipeline

- Kafka / Sqoop based data collection
- Live lookup store for real time decisions
- Tenant / Event and time based data partition
- Time based compaction to optimize query on sparse data
- Summary Profile data to reduce Joins
- Shared compute resources but different context for Scheduled / Ad Hoc jobs or for Algorithmic / Human touchpoints




Compute Layer

- No real 'real time' queries -- FIFO scheduling for user tasks
- Static or rigid resource allocation between scheduled and ad hoc queries / jobs
- Short lived and stateless context - no sticky ness for user defined views like temp tables.
- Interactive queries ?



What was needed for Interactive query...

- SQL like Query Tool for Ad Hoc Analysis.
- Scalability for concurrent users,
 - Fair Scheduling
 - Responsiveness
- High Availability
- Performance – specifically for scans and Joins
- Extensibility – User Views / Datasets / UDF's


Indicium ?

Latin ▾
  

Indicium Edit

English ▾
 

information

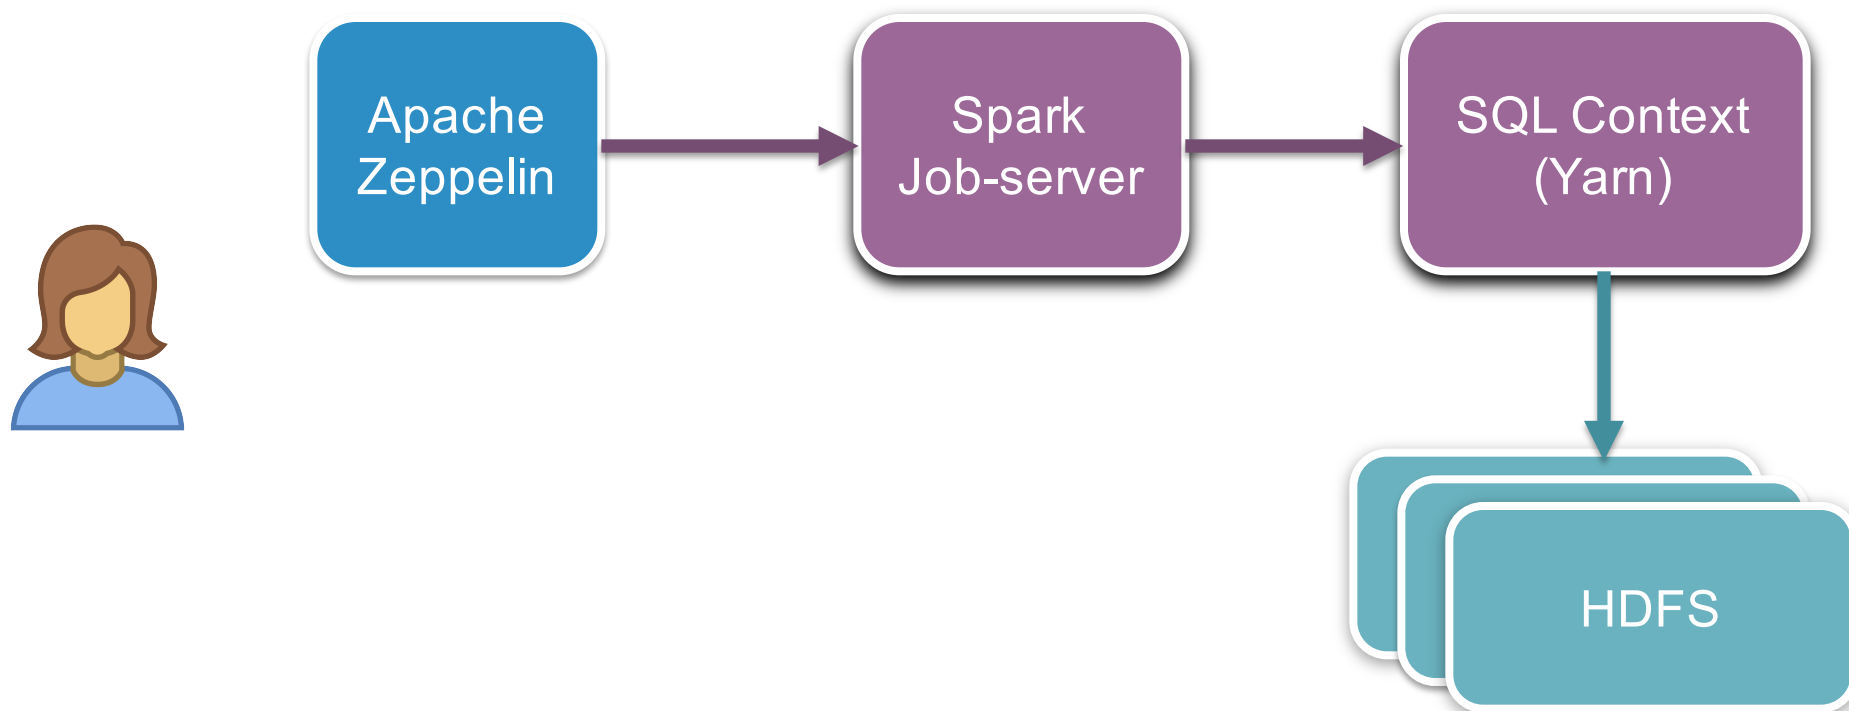
 24 more translations

[Open in Google Translate](#)[Feedback](#)

Indicium: Part 1

Managed Context Pool

Managed Context Pool



Managed Context Pool

Apache Zeppelin 0.6

- SQL like Query tool and a notebook
- Custom interpreter
 - Configuration: SJS server + context
 - Statement execution: Make asynchronous REST calls to SJS
- Concurrency - Multiple interpreters and notebooks

Spark Job-Server 0.6.x

- Custom SQL context with catalog override
- Custom application to execute queries
- High Availability: Multiple SJS servers and multiple contexts per server

Managed Context Pool

Features

- Familiar SQL interface on notebooks
- Concurrent multi-user support
- Visualization Dashboards
- Long running Spark Job – to support User Defined Views
- Access control on Spark APIs
- Custom SQL context with custom catalog
 - Intercept *lookupTable* calls to query actual data
 - Table wrappers for time windows - like *select count(*) from `lastXDays(table)`*

Managed Context Pool

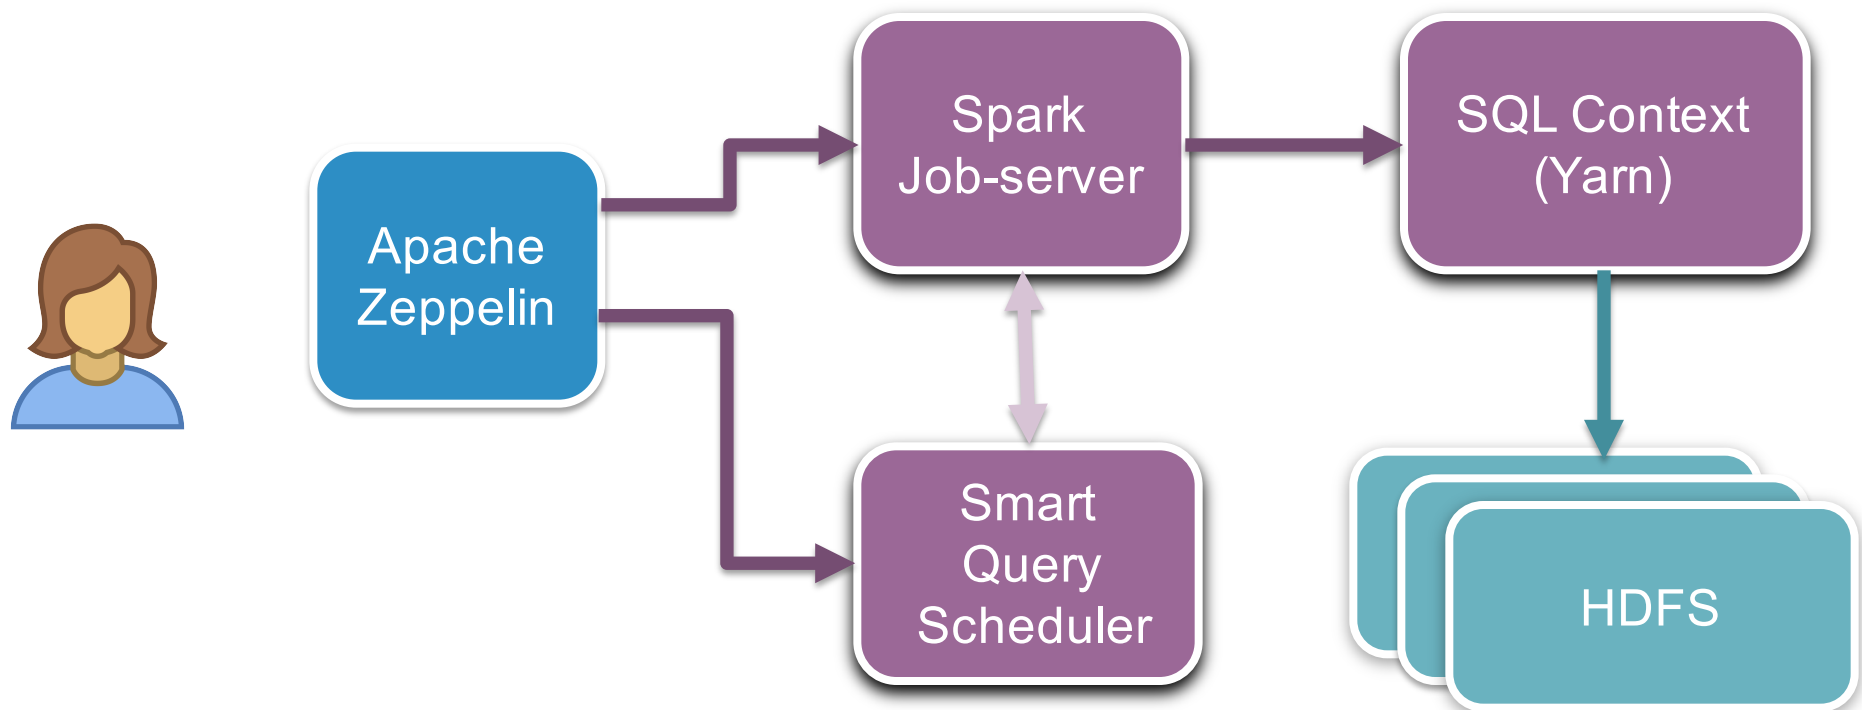
Issues

- Interpreter hard wired to a context
- FIFO scheduling: Single statement per interpreter-context pair – across notebooks / across users
- No automated failure handling
 - Detecting a dead context / SJS server
 - Recovery from the context / server failure
- No dynamic scheduling / load balancing
 - No way of identify an overloaded context
- Incompatible with Spark 2.x

Indicium: Part 2

Smart Query Scheduler

Smart Query Scheduler



Smart Query Scheduler

Zeppelin 0.7

- Supports per notebook statement execution

SJS 0.7 Custom Fork

- Support for Spark 2.x

Smart Query Scheduler:

- Scheduling: API to dynamically bind SJS server + context for every job / query

Other Optimizations:

- Monitoring: Monitor jobs running per context
- Availability: Track Health of SJS servers and contexts and ensures healthy context in pool

Smart Query Scheduler

Dynamic scheduling for every query

- Zeppelin interpreter agnostic of actual SJS / context
- Load balancing of jobs per context
- Query Classification and intelligent routing
- Dynamic scaling / de-scaling the pool size
- Shared Cache
- User Defined Views
- Workspaces or custom time window view for every interpreter

Query Classification / routing

Custom resource configurations for context dedicated for complex or asynchronous queries / jobs:

- Classify queries based on heuristics / historic data into light / heavy queries and route them to different context.
- Separate contexts for interactive vs background queries
 - An export table call does not starve an interactive SQL query

Spark Dynamic Context

Elastic scaling of contexts, co-existing on same cluster as scheduled batch jobs

- Scale up in day time, when user load is high
- Scale down in night, when overnight batch jobs are running
- Scaling also helped to create reserved bandwidth for any set of users, if needed.

Shared Cache

Alluxio to store common datasets

- Single cache for common datasets across contexts
 - Avoids replication across contexts
 - Cached data safe from executor / context crashes
- Dedicated refresh thread to release / update data consistently across contexts

Persistent User Defined Views

- Users can define a temp view for a SQL query
- Replicated across all SJS servers + contexts
- Definitions persisted in DB so that a context restart is accompanied by temp views' registration.
- Load on start to warm up load of views
- TTL support for expiry

Workspaces

- Support for multiple custom catalogs in SQL context for table resolution
- Custom time range / source / caching
 - Global
 - Per catalog
 - Per table
- Configurable via Zeppelin interpreter
- Decoupled time range from query syntax
 - Join a behavior table(refer to last 30 days) with lookup table (fetch complete data)

Automated Pool Management

- Monitoring scripts to track and restart unhealthy / unresponsive SJS servers / contexts
- APIs on SJS to stop / start / refresh context / SJS
- APIs to refresh cached tables / views;
- APIs on Router Service to reconfigure routing / pool size and resource allocation

Thank You !

Questions & Answers

kapil.ee06@gmail.com
arvind_heda@yahoo.com

References

- Apache Zeppelin: <https://zeppelin.apache.org/>
- Spark Job-server: <https://github.com/spark-jobserver/spark-jobserver>
- Alluxio: <http://www.alluxio.org/>

Scale

- Data
 - ~ 100 TB
 - ~ 1000 Event Types
- 100+ Active concurrent users
- 30+ Automated Agents
- 10000+ Scheduled / 3000+ Ad Hoc Analysis
- Avg data churn per Analysis > 200 GB