

MatFast

IN-MEMORY DISTRIBUTED MATRIX COMPUTATION
PROCESSING AND OPTIMIZATION BASED ON SPARK SQL

Mingjie Tang
Yanbo Liang

Oct, 2017



About Authors

◆ Yongyang Yu

- Machine learning, Database system, Computation Algebra
- PhD students at Purdue University

◆ Mingjie Tang

- Spark SQL, Spark ML, Database, Machine Learning
- Software Engineer at Hortonworks

◆ Yanbo Liang

- Apache Spark committer, Spark MLlib
- Staff Software Engineer at Hortonworks

◆ ... All Other Contributors

Agenda

Motivation

Overview of MatFast

Implementation and optimization

Use cases

Motivation

- ◆ Many applications rely on efficient processing of queries over big matrix data:
 - Recommender systems
 - Social network analysis
 - Predict traffic data flow
 - Anti-fraud and spam detection
 - Bioinformatics

Motivation

◆ Recommender Systems

Netflix's user-movie rating table (sample)

users \ movies					
Alice	4	?	3	5	4
Bob	?	5	4	?	?
Cindy	3	?	?	?	2

Problem: Predict the missing entries in the table

Input: User-movie rating table with missing entries

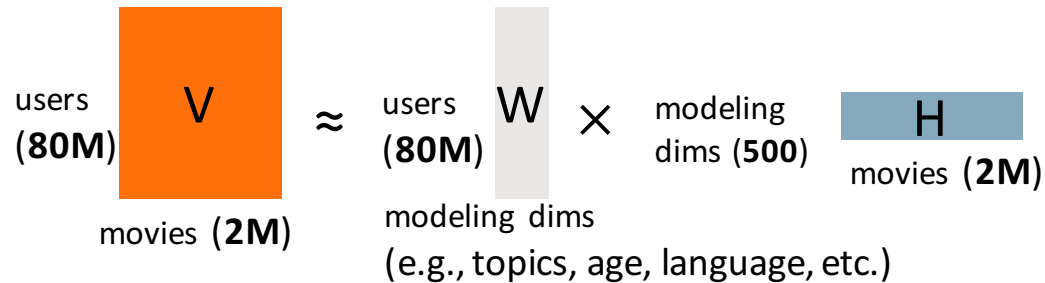
Output: Complete user-movie rating table with predictions

For Netflix, #users = 80 million, #movies = 2 million

Motivation

◆ Gaussian Non-negative Matrix Factorization (GNMF)

— Assumption: $V_{u \times m} \approx W_{u \times t} \times H_{t \times m}$



huge volume

dense/sparse storage

iterative execution

```
Initialize W and H
for i = 1 to nIter do
    H = H * (WT × V) / (WT × W × H)
    W = W * (V × HT) / (W × H × HT)
end
```

Matrix operation for GNMF Algorithm

Motivation

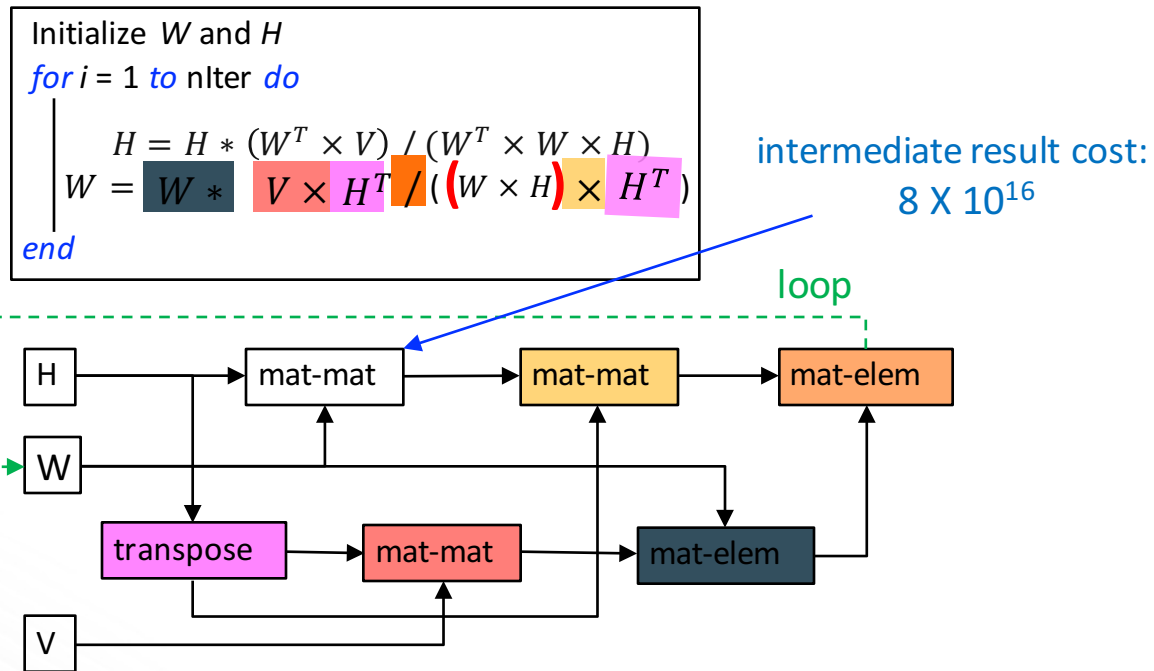
◆ User-Movie Rating Prediction with GNMF

```
val p = 200 // number of topics
val V = loadMatrix("in/V") // read matrix
val max_niter = 10 // max number of iteration
W = RandomMatrix(V.nrows, p)
H = RandomMatrix(p, V.ncols)
for (i <- 0 until max_niter) {
    H = H * (W.t %*% V) / (W.t %*% W %*% H)
    W = W * (V %*% H.t) / (W %*% H %*% H.t)
}
(H %*% W).saveToHive()
```

State of the art solution in Spark ecosystem

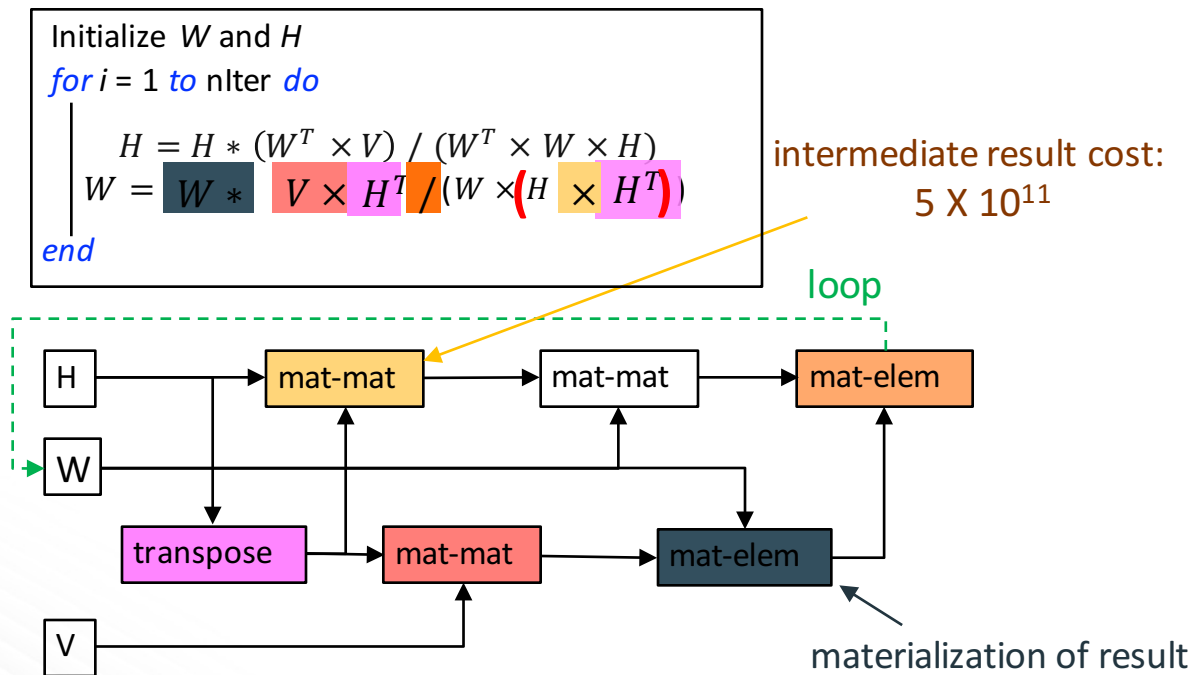
- ◆ Alternative Least Square approach in Spark (ALS)
 - Experiment on Spotify data
 - 50+ million users x 30+ million songs
 - 50 billion ratings For rank 10 with 10 iterations
 - ~1 hour running time
- ◆ How to extend ALS to other matrix computation?
 - SVD
 - PCA
 - QR

Observation



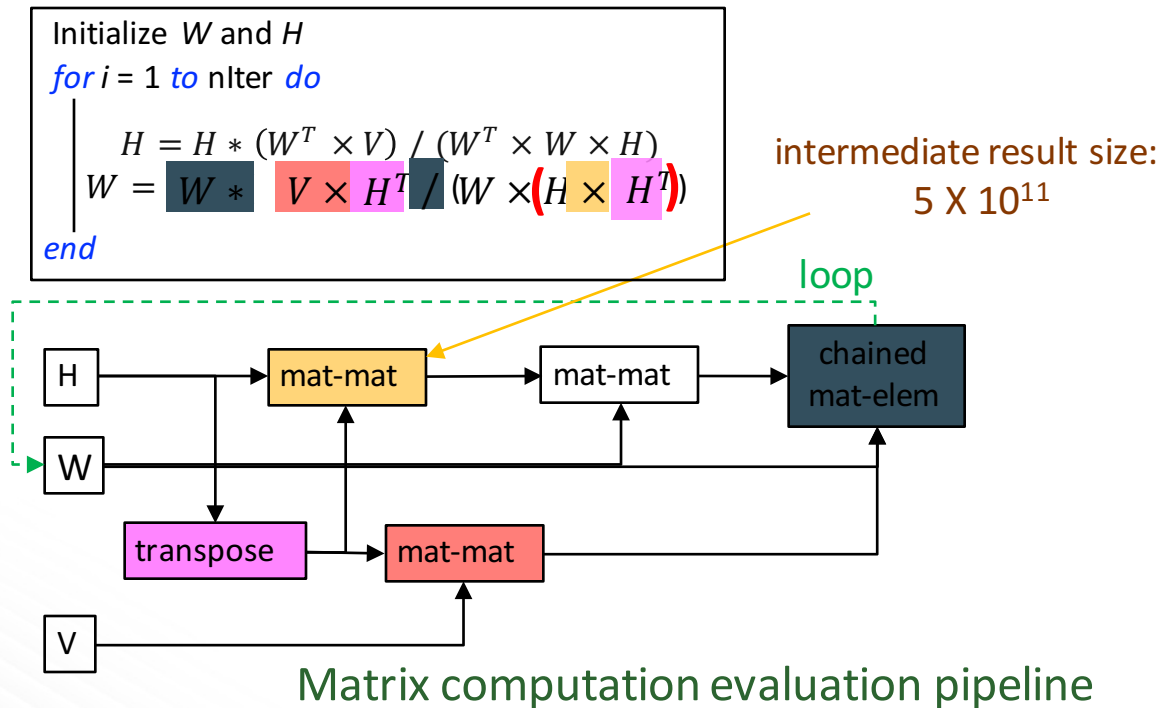
Matrix computation evaluation pipeline

Observation



Matrix computation evaluation pipeline

Observation



Overview of MatFast

Matrix operators

◆ Unary operator

- Transpose: $\mathbf{B} = \mathbf{A}^T$

◆ Binary operators

- $\mathbf{B} = \mathbf{A} + \beta; \mathbf{B} = \mathbf{A} * \beta;$
- $\mathbf{C} = \mathbf{A} \star \mathbf{B}, \star \in \{+, *, /\};$
- $\mathbf{C} = \mathbf{A} \times \mathbf{B} (\mathbf{A} \%*\% \mathbf{B})$

matrix-matrix multiplication

◆ Others

- return a matrix: `abs(A)`, `pow(A, p)`
- return a vector: `rowSum(A)`, `colSum(A)`
- return a scalar: `max(A)`, `min(A)`

Optimization targets

- MATFAST generates a computation- and communication-efficient execution plan:
 - Optimize a *single matrix operator* in an expression ✓
 - Optimize *multiple operators* in an expression
 - Exploit *data dependency* between different expressions ✓

Comparison with other systems

	Single	Distributed w. multiple nodes				
	R	ScaLAPACK	SciDB	SystemML	MLlib	DMac
huge volume.		✓	✓	✓	✓	✓
sparse comp.	✓		~	✓	~	~
multiple operators	✓	✓	✓	✓	✓	✓
partition w. dependency						✓
opt. exec. plan				✓		
interface	R script	C/Fortran	SQL-like	R-like	Java/Scala	Scala
fault tolerance			✓	✓	✓	✓
open source	✓	✓	~	✓	✓	

Compare with Spark SQL

	Matrix operators	SQL relational query
Data type	matrix	relational table
Operators	transpose, mat-mat, mat-scalar, mat-elem	join, select, group by, aggregate
Execution scheme	iterative	acyclic

System framework

Applications: Image processing, Text processing, Collaborative filtering, Spatial computation, etc.

ML algorithms: SVD, PCA, NMF, PageRank, QR, etc

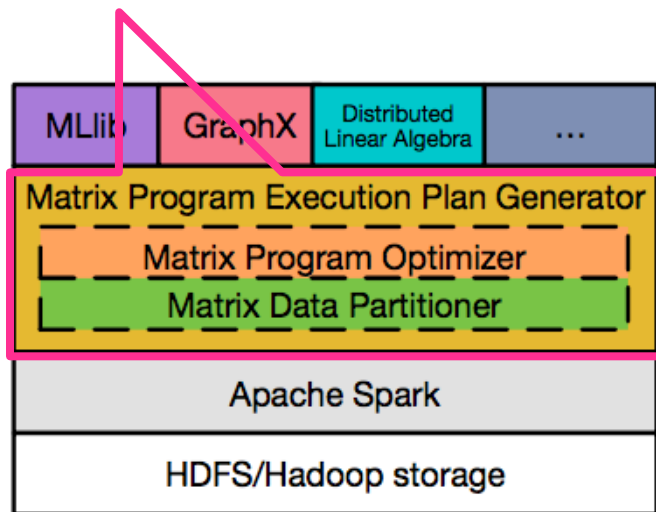
Spark SQL

MATFAST

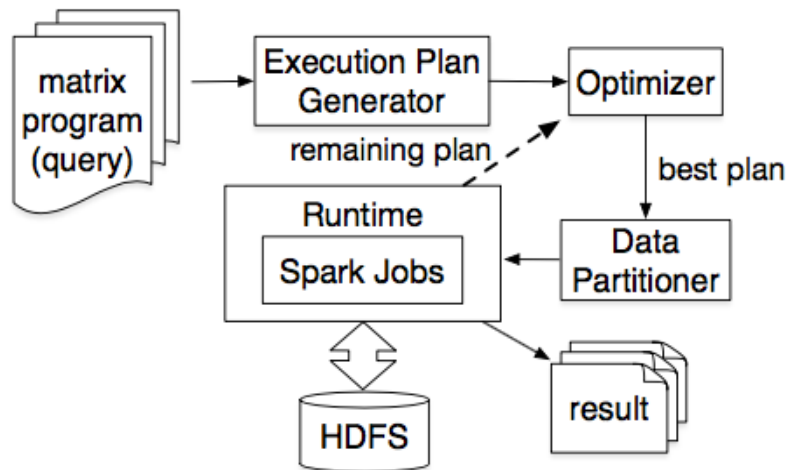
Spark RDD

System framework

MATFAST



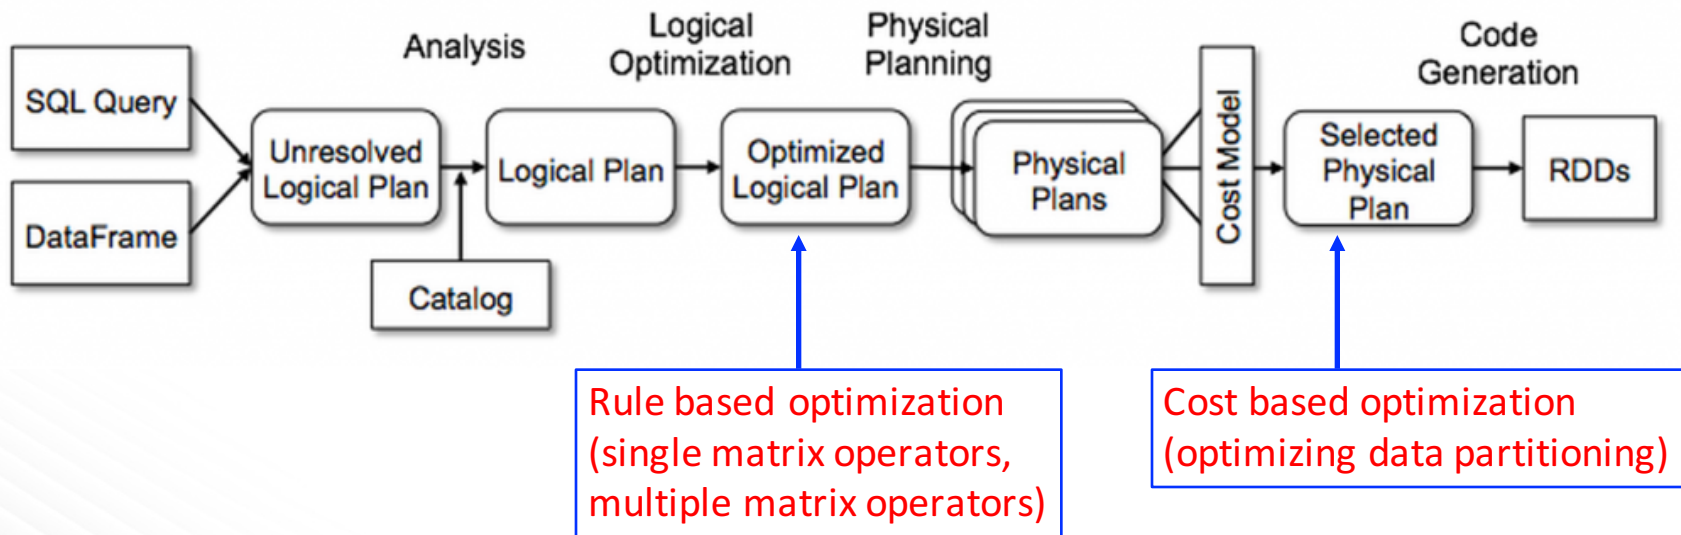
Components



Architecture

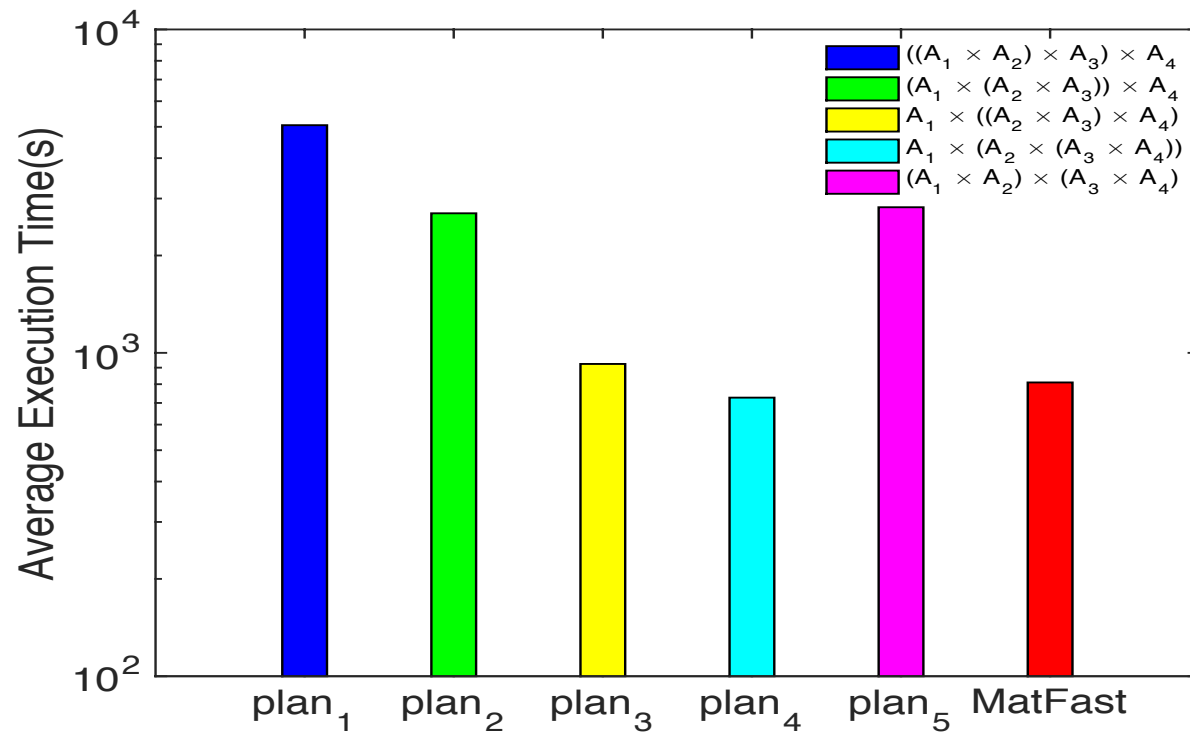
MatFast within Spark Catalyst

- Extend Spark Catalyst



Implementation and optimization

Optimization 1: a Single Operator - Cost Based Optimization

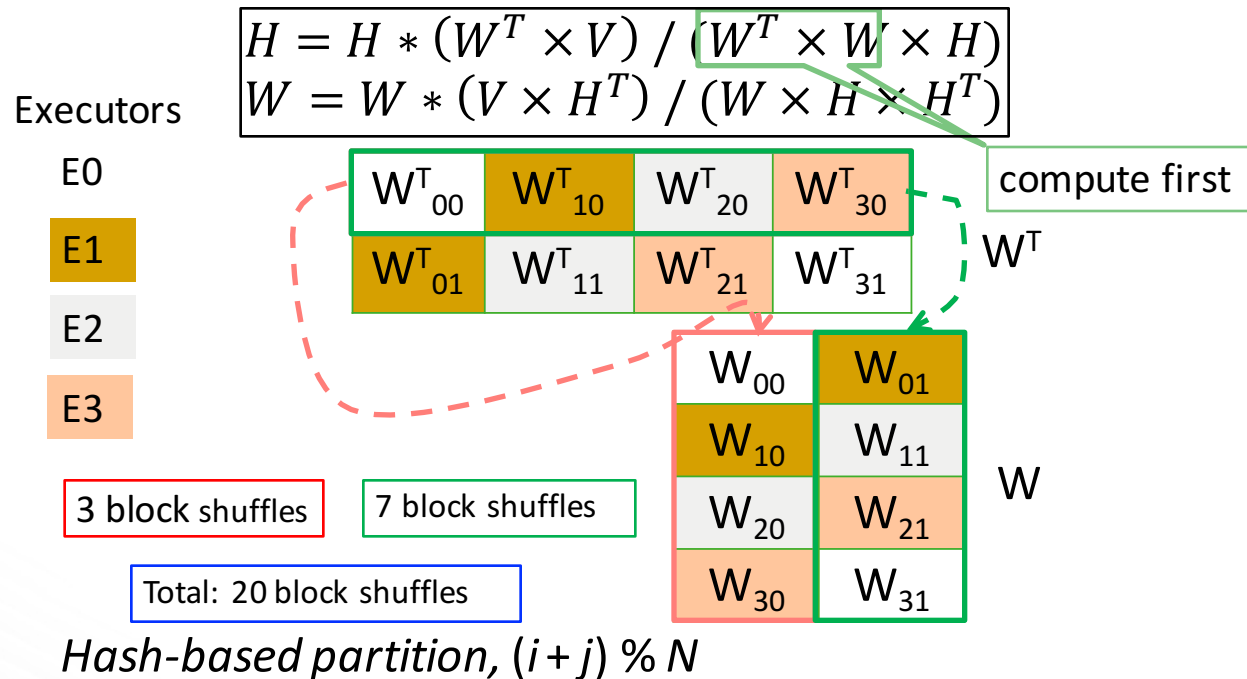


Optimization 2: optimizing data partitioning in pipeline

- ◆ Distribute matrix data over a set of workers
- ◆ How to determine the data partitioning scheme for a matrix such that minimum shuffle cost is introduced for the entire pipeline?
- ◆ Partitioning schemes
 - Row scheme (“ r ”)
 - Column scheme (“ c ”)
 - Block-Cyclic scheme (“ $b-c$ ”)
 - Broadcast scheme (“ b ”)

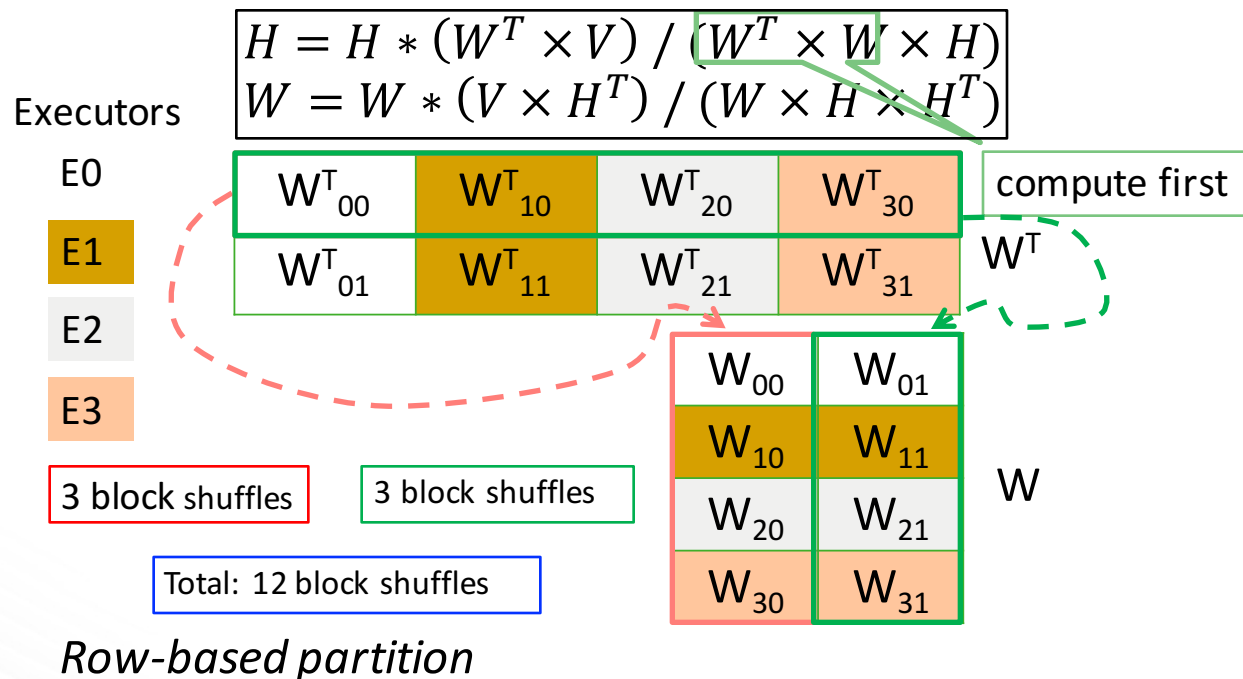
Optimization 2: optimizing data partitioning in pipeline

- How to determine the data partitioning scheme for a matrix such that minimum shuffle cost is introduced for the entire pipeline?



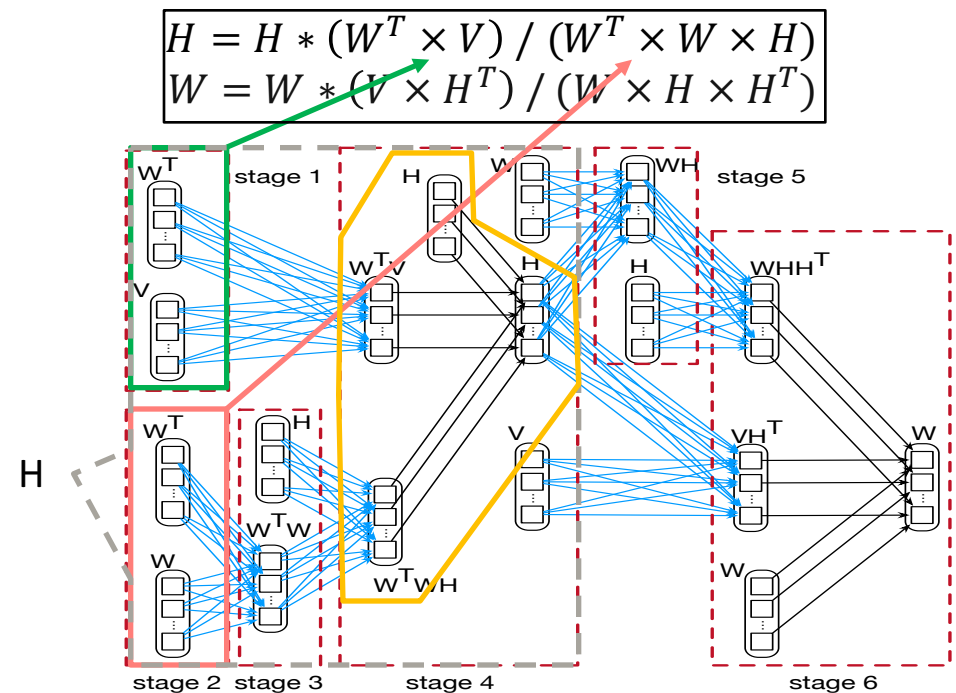
Optimization 2: optimizing data partitioning in pipeline

- How to determine the data partitioning scheme for a matrix such that minimum shuffle cost is introduced for the entire pipeline?



Optimization 2: optimizing data partitioning in pipeline

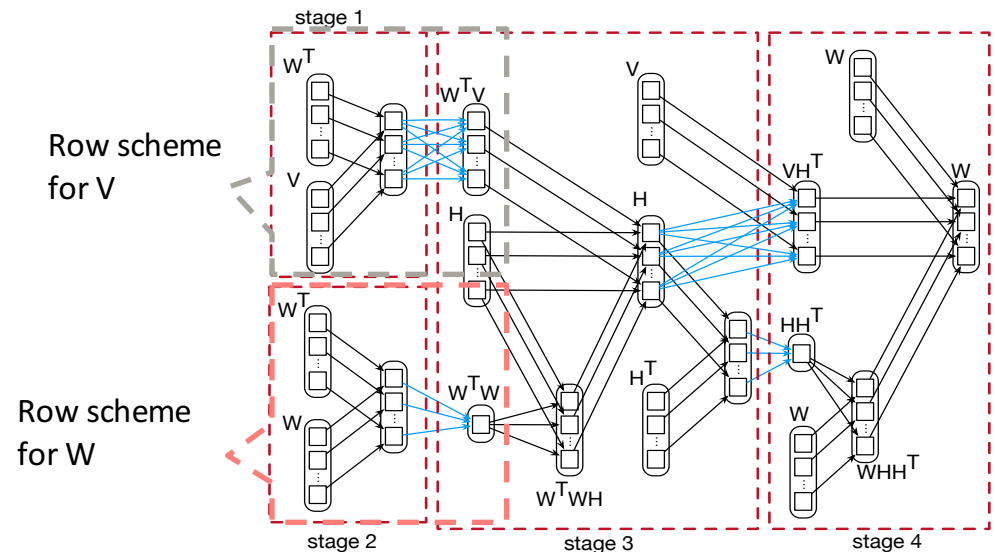
- We need an optimized plan to determine an optimized data partitioning scheme for each matrix such that minimum shuffle overhead is introduced for the entire pipeline.
- For example, with hash-based data partitioning, the computation pipeline involves multiple shuffles for aligning the data blocks.



Optimization 2: optimizing data partitioning in pipeline

- MatFast determines the partitioning scheme for an input matrix with min shuffle cost according to the cost model.
- Greedily** optimizes each operator

$$s_{i1}(i2) \leftarrow \underset{s_{i1}(i2)}{\operatorname{argmin}} C_{comm}(op, s_{i1} [s_{i2}], s_o)$$



- Physical execution plan with optimized data partitioning

Case studies

Experiments

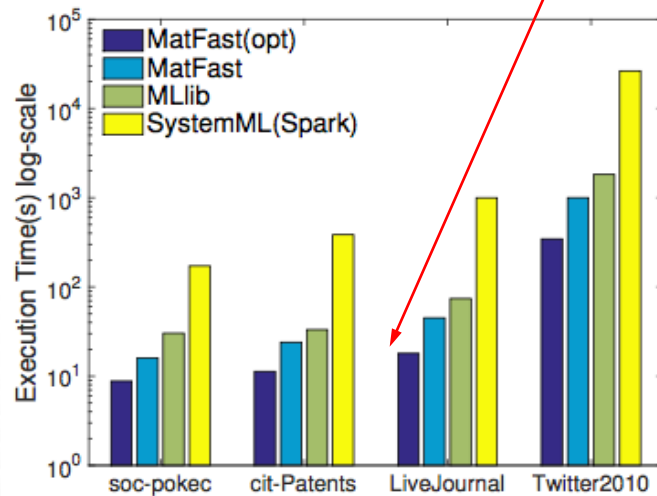
- Dataset APIs
 - [Code examples link](#)
- Compare with state-of-the-art systems
 - Spark MLlib (provided matrix operation)
 - SystemML (Spark)
 - ScaLAPACK
 - SciDB
- Netflix data
 - 100,480,507 ratings
 - 17,770 movies from 480,189 customers
- Social network data

Graph	#nodes	#edges
soc-pokec	1,632,803	30,622,564
cit-Patents	3,774,768	16,518,978
LiveJournal	4,847,571	68,993,773
Twitter2010	41,652,230	1,468,365,182


TABLE V: Statistics of the social network datasets

PageRank on different datasets

MATFAST



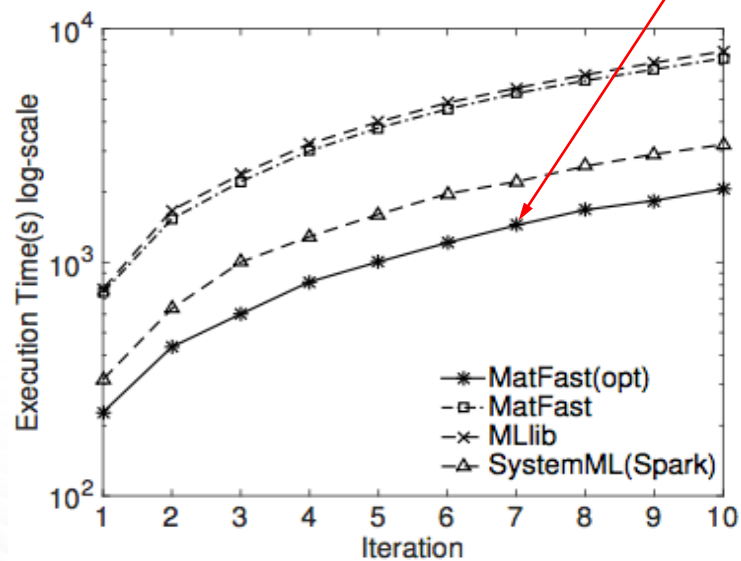
(a) Execution time



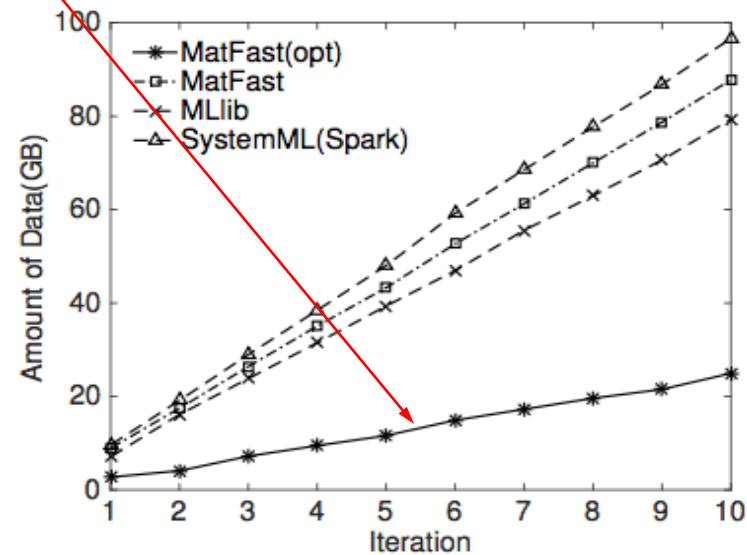
(b) Communication cost

GNMF on the Netflix dataset

MatFAST



(a) Execution time



(b) Communication cost

Future plan

- ◆ More user friend APIs
- ◆ Advanced plan optimizer
- ◆ Python and R interface
- ◆ Vertical applications

Conclusion

- ◆ Proposed and realized MATFAST, an in-memory distributed platform that optimizes query pipelines of matrix operations
- ◆ Take advantage of dynamic cost-based analysis and rule-based heuristics to generate a query execution plan
- ◆ Communication-efficient data partitioning scheme assignment

Reference

- Yongyang Yu, MingJie Tang, Walid G. Aref, Qutaibah M. Malluhi, Mostafa M. Abbas, Mourad Ouzzani:
In-Memory Distributed Matrix Computation Processing and Optimization. ICDE 2017: 1047-1058

Thanks

Q & A

mtang@hortonworks.com