



Real-time Machine Learning with Redis-ML and Apache Spark

Dvir Volk - Redis Labs

Agenda

- Intro to Redis and Redis Labs - 5 min
- Using Redis-ML for Model Serving - why and how - 10 min
- Building a recommendation system using Spark-ML and Redis-ML - 10 min
- Q&A

Redis Labs – Home of Redis



The commercial company behind Open Source Redis

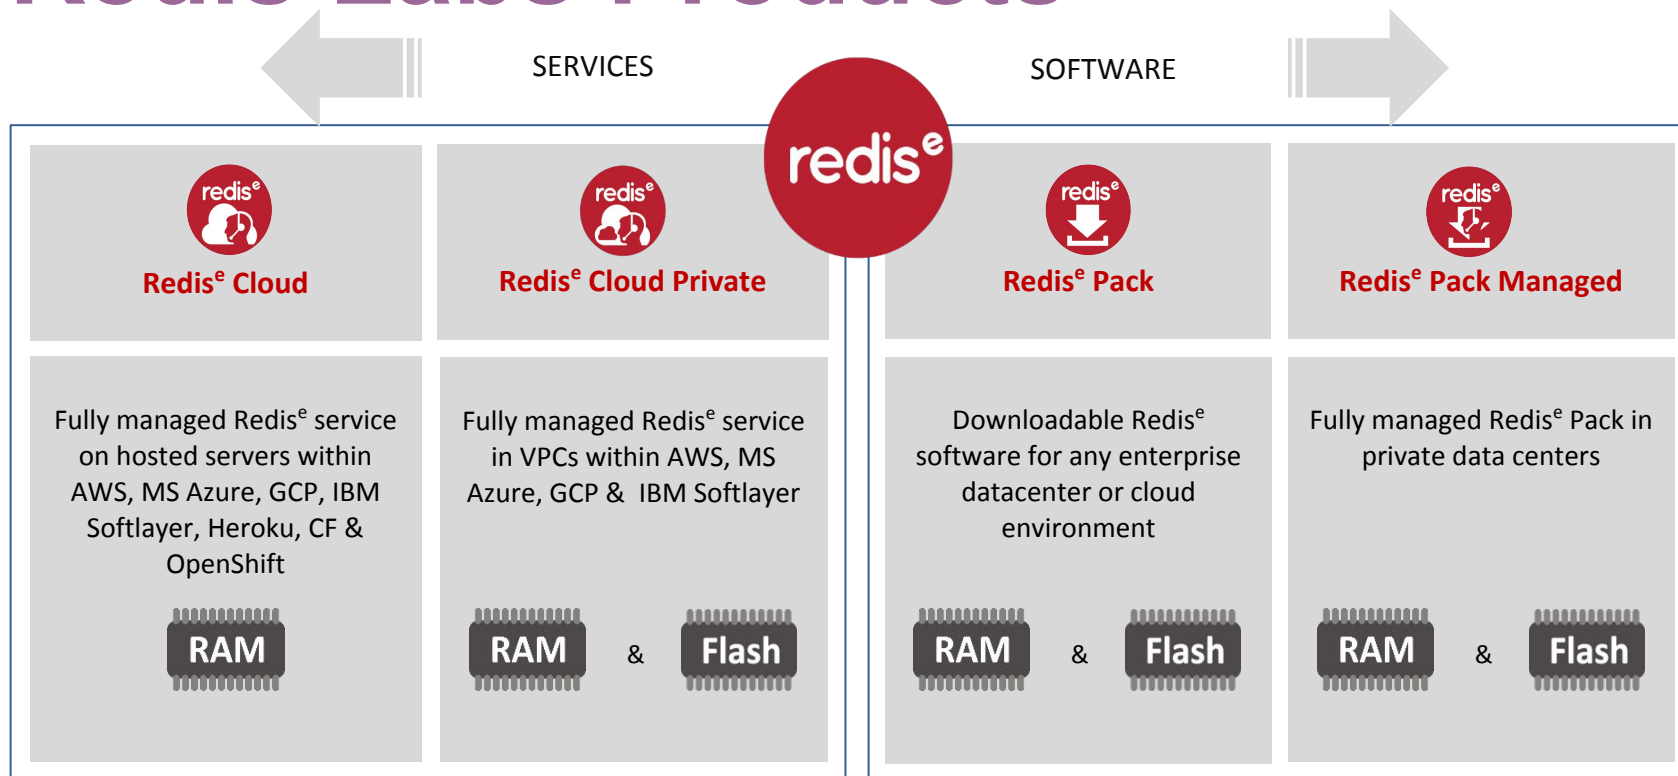


Provider of the **Redis Enterprise (Redis^e)** technology, platform and products

Founded in 2011

HQ in Mountain View CA, R&D center in Tel-Aviv IL

Redis Labs Products




A Brief Overview of Redis

- Started in 2009 by Salvatore Sanfilippo
- Most popular KV store
- In memory - disk backed
- Notable Users:
 - Twitter, Netflix, Uber, Groupon, Twitch
 - Many, many more...




Redis Main Differentiators



Performance

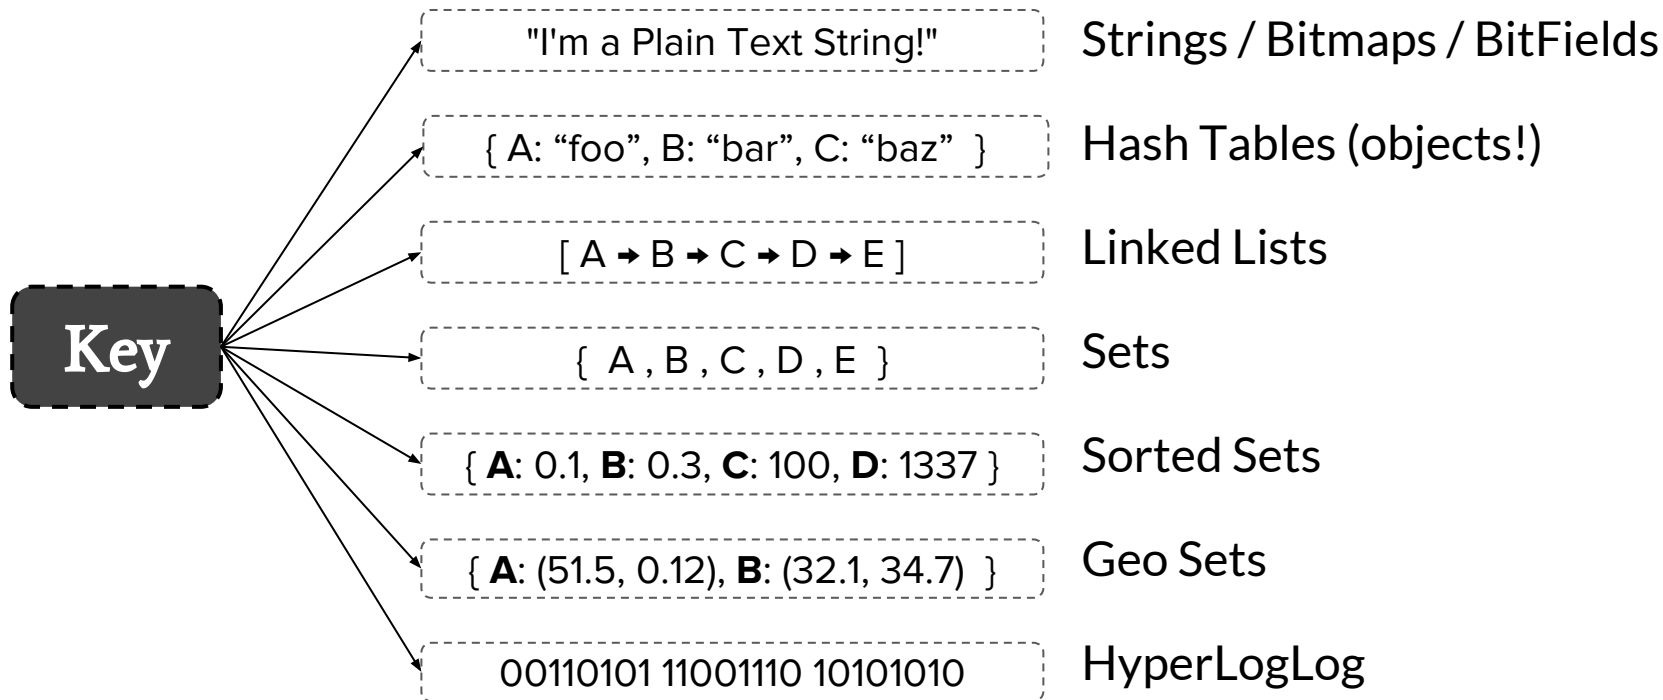


Simplicity
(through Data Structures)



Extensibility
(through Redis Modules)

A Quick Recap of Redis



Simple Redis Example (string)

```
127.0.0.1:6379> SET spark summit
```

```
OK
```

```
127.0.0.1:6379> GET spark
```

```
"summit"
```

```
127.0.0.1:6379>
```


Simple Redis Example (hash)

```
127.0.0.1:6379> HMSET spark_hash org apache version 2.1.1
```

```
OK
```

```
127.0.0.1:6379> HGET spark_hash version
```

```
"2.1.1"
```

```
127.0.0.1:6379> HGETALL spark_hash
```

```
1) "org"
```

```
2) "apache"
```

```
3) "version"
```


```
4) "2.1.1"
```

Another Redis Example (ZSET)

```
127.0.0.1:6379> ZADD my_sorted_set 1 foo
(integer) 1
127.0.0.1:6379> ZADD my_sorted_set 5 bar
(integer) 1
127.0.0.1:6379> ZADD my_sorted_set 3 baz
(integer) 1
127.0.0.1:6379> ZRANGE my_sorted_set 0 2
1) "foo"
2) "baz"
3) "bar"
```

Redis Modules: A New Hope

- Dynamic libraries loaded to redis
- Written in C/C++
- Use a C ABI/API isolating redis internals
- Use existing or add new data-structures
- Near Zero latency access to data




Modules: A New Approach

Adapt your database to your data, not the other way around

Neural Redis Simple Neural Network Native to Redis	Redis-ML Machine Learning Model Serving	Redisearch Full Text Search Engine in Redis
ReJSON JSON Engine on Redis. Pre-released	Time Series Time series values aggregation in Redis	Graph Graph database on Redis based on Cypher language
Rate Limiter Based on Generic Cell Rate Algorithm (GCRA)	Crypto Engine Wrapper Secure way to store data in Redis via encrypt/decrypt with various Themis primitives	Secondary Index/RQL Indexing + SQL-like syntax for querying indexes. Pre-released





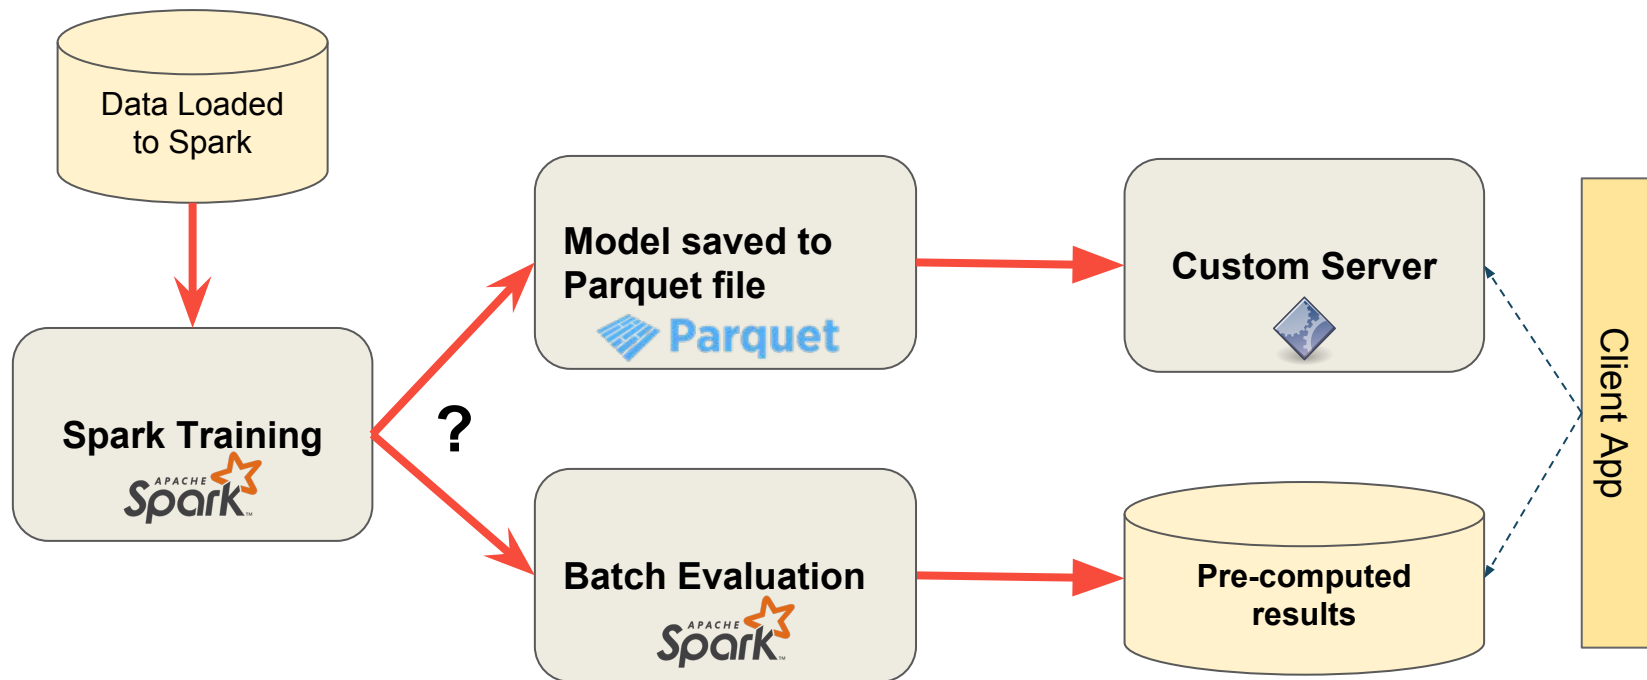
Redis ML

Machine Learning Model Server

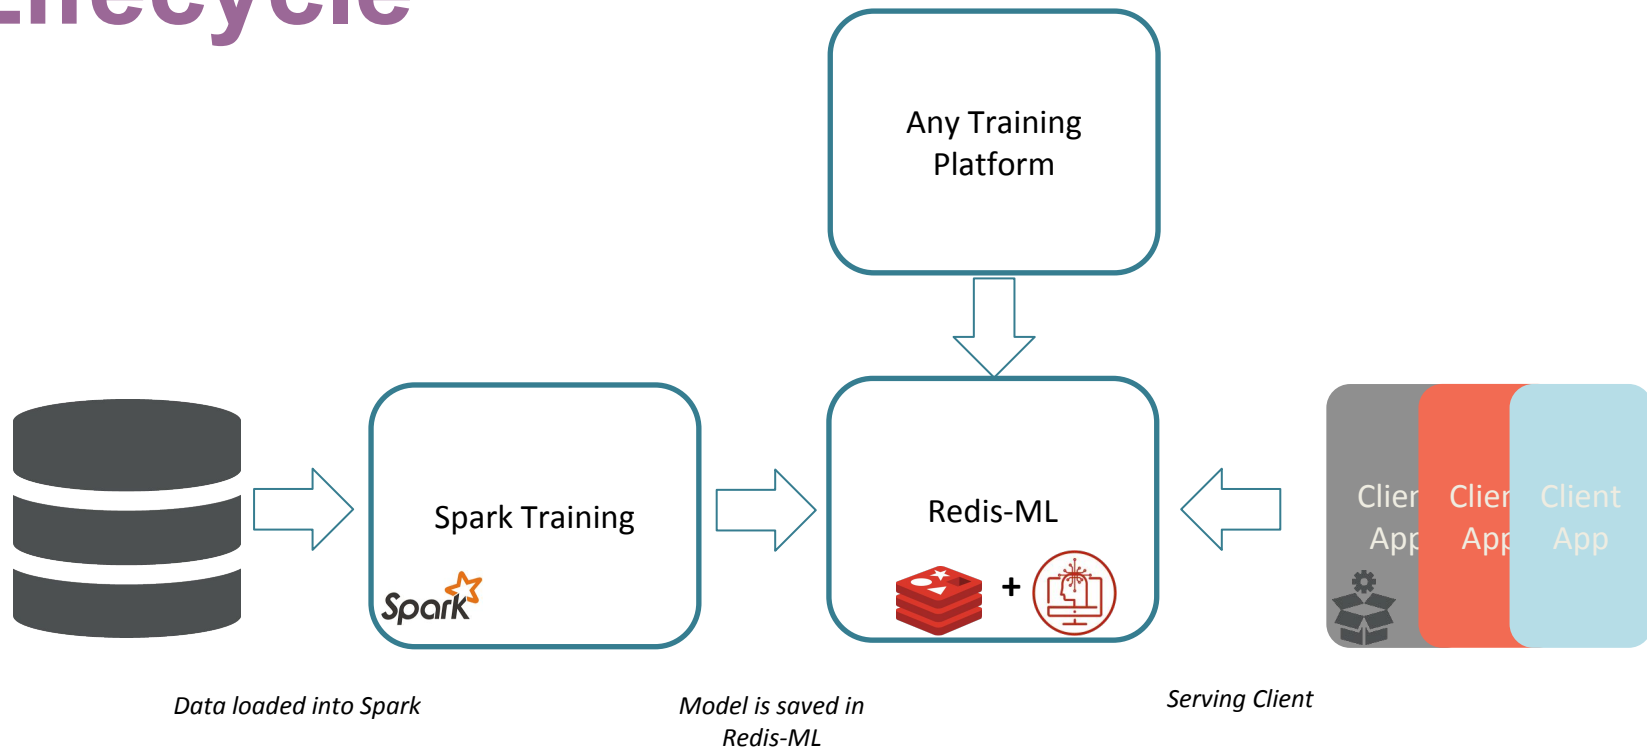
ML Models Serving Challenges

- Models are becoming bigger and more complex
- Can be challenging to deploy & serve
- Do not scale well, speed and size
- Can be very expensive

Spark-ML End-to-End Flow




A Simpler Machine Learning Lifecycle



Redis-ML – ML Serving Engine

- Store training output as “hot model”
- Perform evaluation directly in Redis
- Enjoy the performance, scalability and HA of Redis

Redis-ML



ML Models

Tree Ensembles

Linear Regression


Logistic Regression

Matrix + Vector Operations


More to come...

Random Forest Model

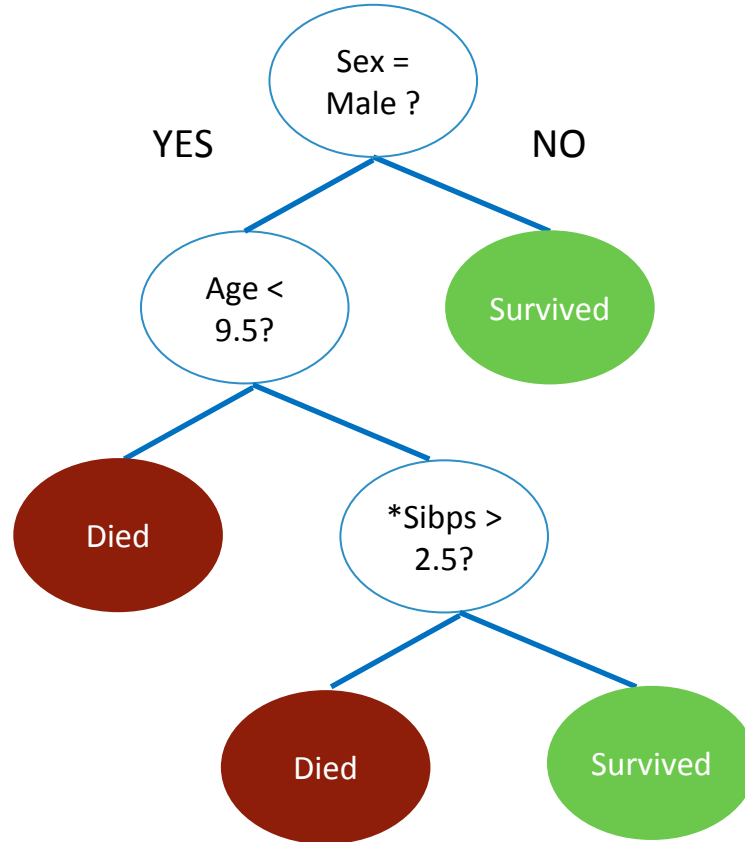
- A collection of decision trees
- Supports classification & regression
- Splitter Node can be:
 - Categorical (e.g. day == “Sunday”)
 - Numerical (e.g. age < 43)
- Decision is taken by the majority of decision trees



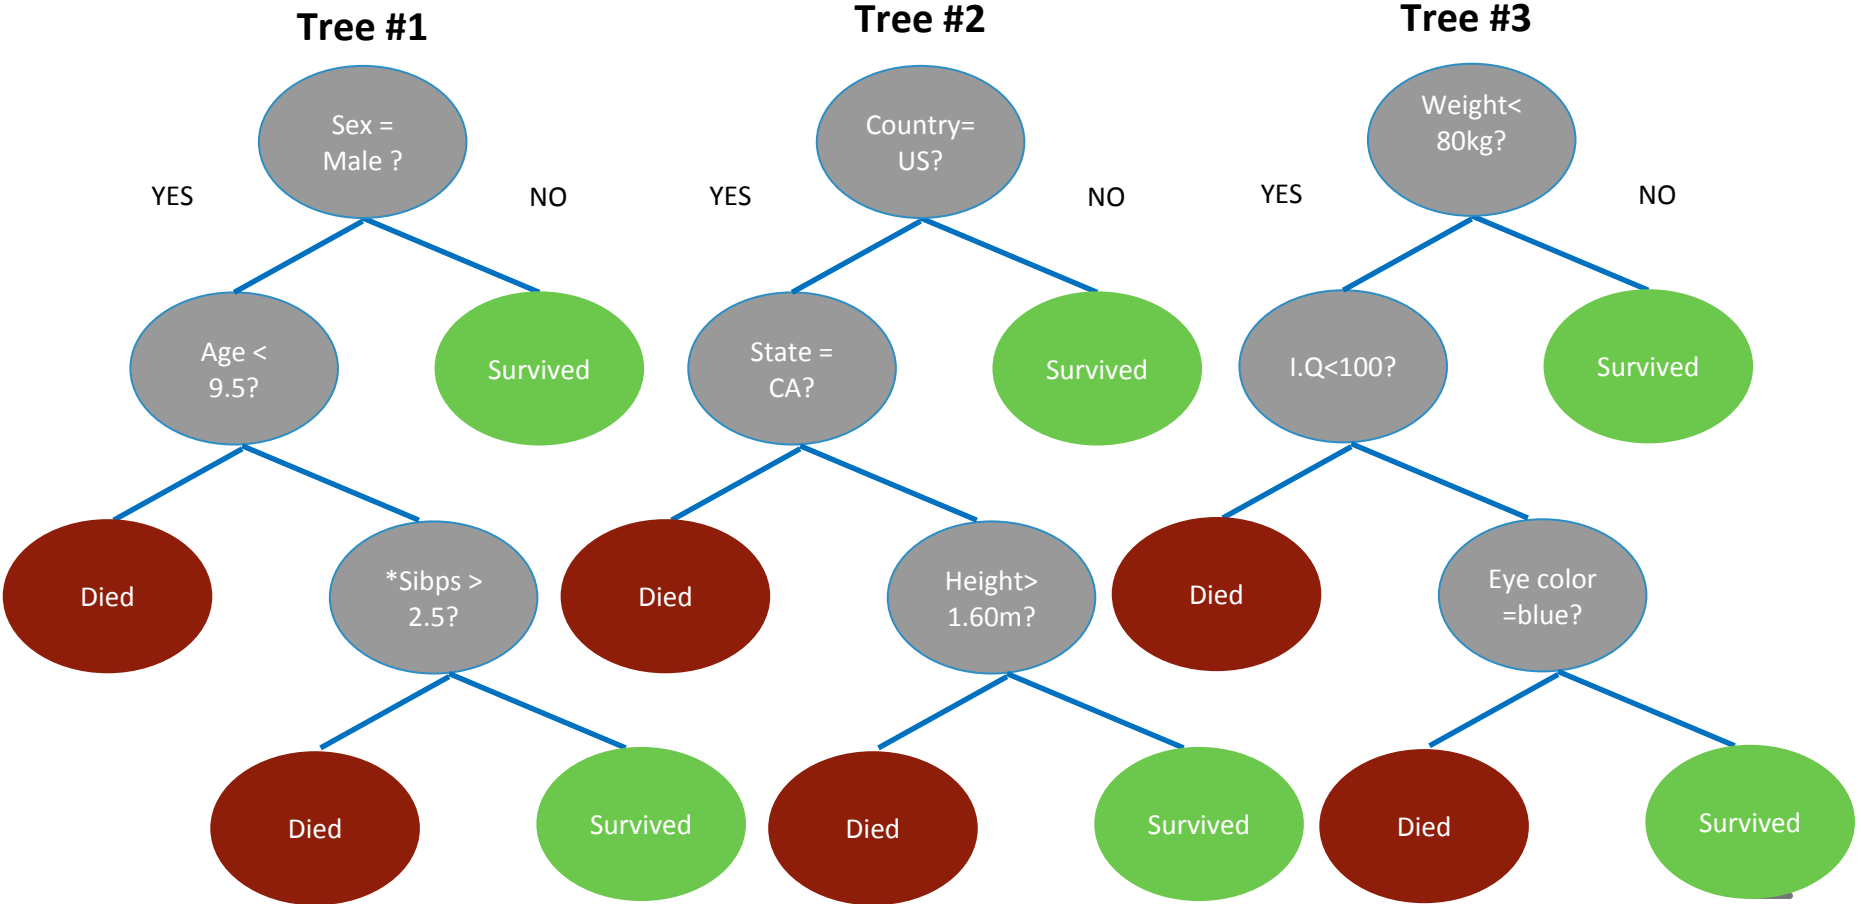
Titanic Survival Predictor on a Decision Tree



*Sibps = siblings + spouses



Titanic Survival Predictor on a Random Forest



Would John Survive The Titanic

- John's features:
{male, 34, married + 2, US, CA, 1.78m, 78kg, 110iq, blue eyes}
- Tree#1 – Survived
- Tree#2 – Failed
- Tree#3 – Survived
- Random forest decision - Survived

Forest Data Type Example

```
> MODULE LOAD "./redis-ml.so"
```

```
OK
```

```
> ML.FOREST.ADD myforest 0 . CATEGORIC sex "male" .L LEAF 1 .R LEAF 0
```

```
OK
```

```
> ML.FOREST.RUN myforest sex:male
```

```
"1"
```

```
> ML.FOREST.RUN myforest sex:no_thankx
```

```
"0"
```

Using Redis-ML With Spark

```
scala> import com.redislabs.client.redisml.MLClient
scala> import com.redislabs.provider.redis.ml.Forest

scala> val jedis = new Jedis("localhost")
scala> val rfModel = pipelineModel.stages.last.asInstanceOf[RandomForest]

// Create a new forest instance
scala> val f = new Forest(rfModel.trees)

// Load the model to redis
scala> f.loadToRedis("forest-test", "localhost")

// Classify a feature vector
scala> jedis.getClient().sendCommand(MLClient.ModuleCommand.FOREST_RUN,
"forest-test", makeInputString(0))

scala> jedis.getClient().getStatusCodeReply
res53: String = 1
```

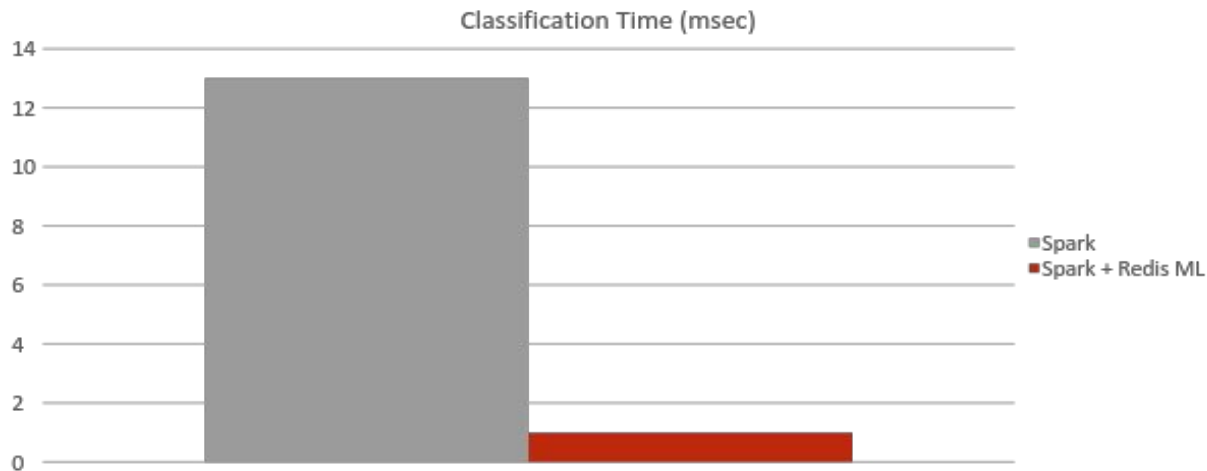

Real World Challenge

- Ad serving company
- Need to serve 20,000 ads/sec @ 50msec data-center latency
- Runs 1k campaigns → 1K random forest
- Each forest has 15K trees
- On average each tree has 7 levels (depth)
- Would require < 1000 x c4.8xlarge

Redis ML with Spark ML


40x Faster

Classification Time Over Spark




Real World Example: Movie Recommendation System

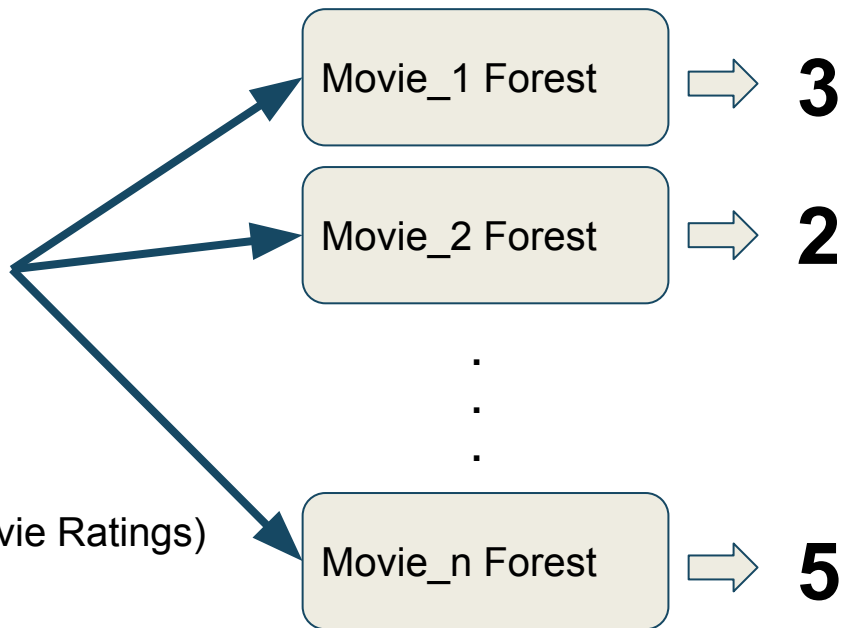
Overview



Concept: One Forest For Each Movie



User Features:
(Age, Gender, Movie Ratings)



The Tools


Transform:




Train:



Classify:



Containers:



Using the Dockers

```
$ docker pull shaynativ/redis-ml
$ docker run --net=host shaynativ/redis-ml &
$
$ docker pull shaynativ/spark-redis-ml
$ docker run --net=host shaynativ/spark-redis-ml
```

Step 1: Get The Data

- Download and extract the [MovieLens 100K Dataset](#)
- The data is organized in separate files:
 - Ratings: user id | item id | rating (1-5) | timestamp
 - Item (movie) info: movie id | genre info fields (1/0)
 - User info: user id | age | gender | occupation
- Our classifier should return the expected rating (from 1 to 5) a user would give the movie in question

Step 2: Transform

- The training data for each movie should contain 1 line per user:
 - class (rating from 1 to 5 the user gave to this movie)
 - user info (age, gender, occupation)
 - user ratings of other movies (movie_id:rating ...)
 - user genre rating averages (genre:avg_score ...)
- Run `gen_data.py` to transform the files to the desired format

Step3: Train and Load to Redis

```
// Create a new forest instance
```

```
val rf = new
```

```
RandomForestClassifier().setFeatureSubsetStrategy("auto").setLabelCol("indexedLabel").setFeaturesCol("indexedFeatures").setNumTrees(500)
```

```
....
```

```
// Train model
```

```
val model = pipeline.fit(trainingData)
```

```
....
```

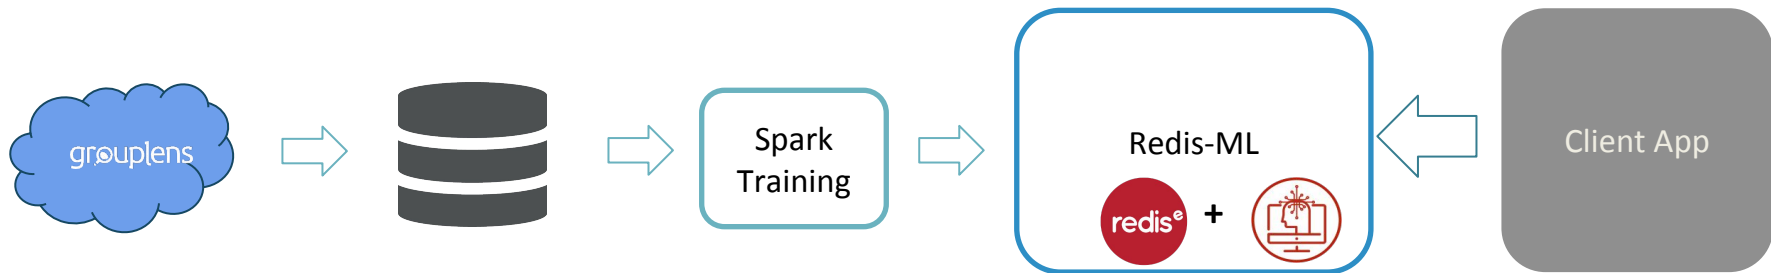
```
val rfModel = model.stages(2).asInstanceOf[RandomForestClassificationModel]
```

```
// Load the model to redis
```

```
val f = new Forest(rfModel.trees)
```

```
f.loadToRedis("movie-10", "127.0.0.1")
```

Step 3: Execute in Redis



Python Client Example

```
>> import redis
>> config = {"host":"localhost", "port":6379}
>> r = redis.StrictRedis(**config)
>> user_profile = r.get("user_shay_profile")
>> print(user_profile)
12:1.0,13:1.0,14:3.0,15:1.0,17:1.0,18:1.0,19:1.0,20:1.0,23:1.0,24:5.0,1.0,115:1.0,116:2.
0,117:2.0,119:1.0,120:4.0,121:2.0,122:2.0,
.....
1360:1.0,1361:1.0,1362:1.0,
1701:6.0,1799:435.0,1801:0.2,1802:0.11,1803:0.04,1812:0.04,1813:0.07,1814:0.24,1815:0.09
,1816:0.32,1817:0.06
>> r.execute_command("ML.FOREST.RUN", "movie-10", user_profile)
'3'
```

Redis CLI Example

```
>keys *
```

```
127.0.0.1:6379> KEYS *
```

```
1) "movie-5"
```

```
2) "movie-1"
```

```
.....
```

```
10) "movie-10"
```

```
11) "user_1_profile"
```

```
>ML.FOREST.RUN movie-10
```

```
12:1.0,13:1.0,,332:3.0,333:1.0,334:1.0,335:2.0,336:1.0,357:2.0,358:1.0,359
```

```
:
```

```
.....
```

```
412:2.0,423:1.0,454:1.0,455:1.0,456:1.0,457:3.0,458:1.0,459:1.0,470:1"
```

Performance

```
Redis time: 0.635129ms, res=3
```

```
Spark time: 46.657662ms, res=3.0
```

```
-----
```

```
Redis time: 0.644444ms, res=3
```

```
Spark time: 49.028983ms, res=3.0
```

```
-----
```

```
Classification averages:
```

```
redis: 0.9401250000000001 ms
```

```
spark: 58.01970206666667 ms
```

```
ratio: 61.71488053893542
```

```
diffs: 0.0
```

Python Script

```
r = redis.StrictRedis()

user_profile = r.get("user-1-profile")
results = {}

for i in range(1, 11):
    results[i] = r.execute_command("ML.FOREST.RUN", "movie-{}".format(i), user_profile)

sorted_results = sorted(results.items(), key=operator.itemgetter(1), reverse=True)
for k,v in sorted_results:
    print "movie-{}:{}".format(k,v)

print "\nRecommended movie: movie-{}".format(sorted_results[0][0])
```

The Results...

```
$ ./classify_user.py 1
```

```
Movies sorted by scores:
```

```
movie-4:3
```

```
movie-3:2
```

```
movie-6:2
```

```
movie-7:2
```

```
movie-8:2
```

```
movie-9:2
```

```
movie-1:1
```

```
movie-2:1
```


```
movie-5:1
```

```
movie-10:0
```

```
Recommended movie for user 1: movie-4
```


Summary

- Train with Spark, Serve with Redis
- 97% resource cost serving
- Simplify ML lifecycle
- Redis^e (Cloud or Pack):
 - Scaling, HA, Performance
 - PAYG – cost optimized
 - Ease of use



Data loaded into Spark

*Model is saved in
Redis-ML*

Serving Client

Resources

- Redis-ML: <https://github.com/RedisLabsModules/redis-ml>
- Spark-Redis-ML: <https://github.com/RedisLabs/spark-redis-ml>
- Databricks Notebook: <http://bit.ly/sparkredism1>
- Dockers: <https://hub.docker.com/r/shaynativ/redis-ml/>
<https://hub.docker.com/r/shaynativ/spark-redis-ml/>



Thank You.

dvir@redislabs.com