



RUNNING SPARK INSIDE CONTAINERS

Haohai Ma, IBM

Khalid Ahmed, IBM

Myself

- “How High”
- Software Architect
- IBM Spectrum Computing
- Toronto Canada

Agenda

- Why container?
- Migrate spark workload to container
- Spark instance on Kubernetes
 - Architecture
 - Workflow
 - Multi-tenancy
- Future work

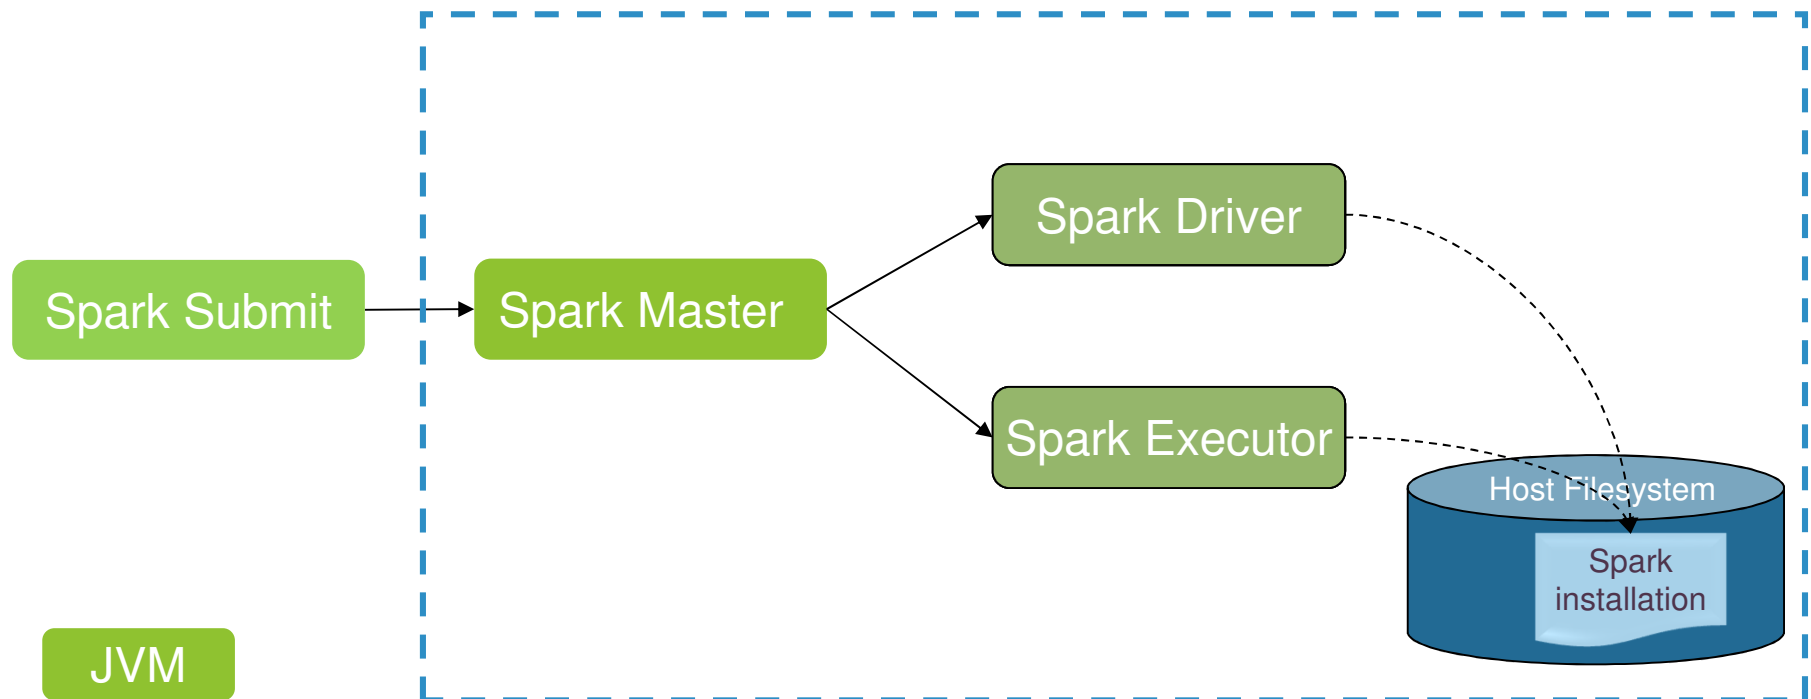
Why use containers?

- To enforce the CPU and memory bounds.
 - CPU shares are proportional to the allocated slots
 - `spark.driver.memory` & `spark.executor.memory`
- To completely isolate the file system
 - Solve the dependency conflicts
- To create and ship images
 - Develop once and run everywhere

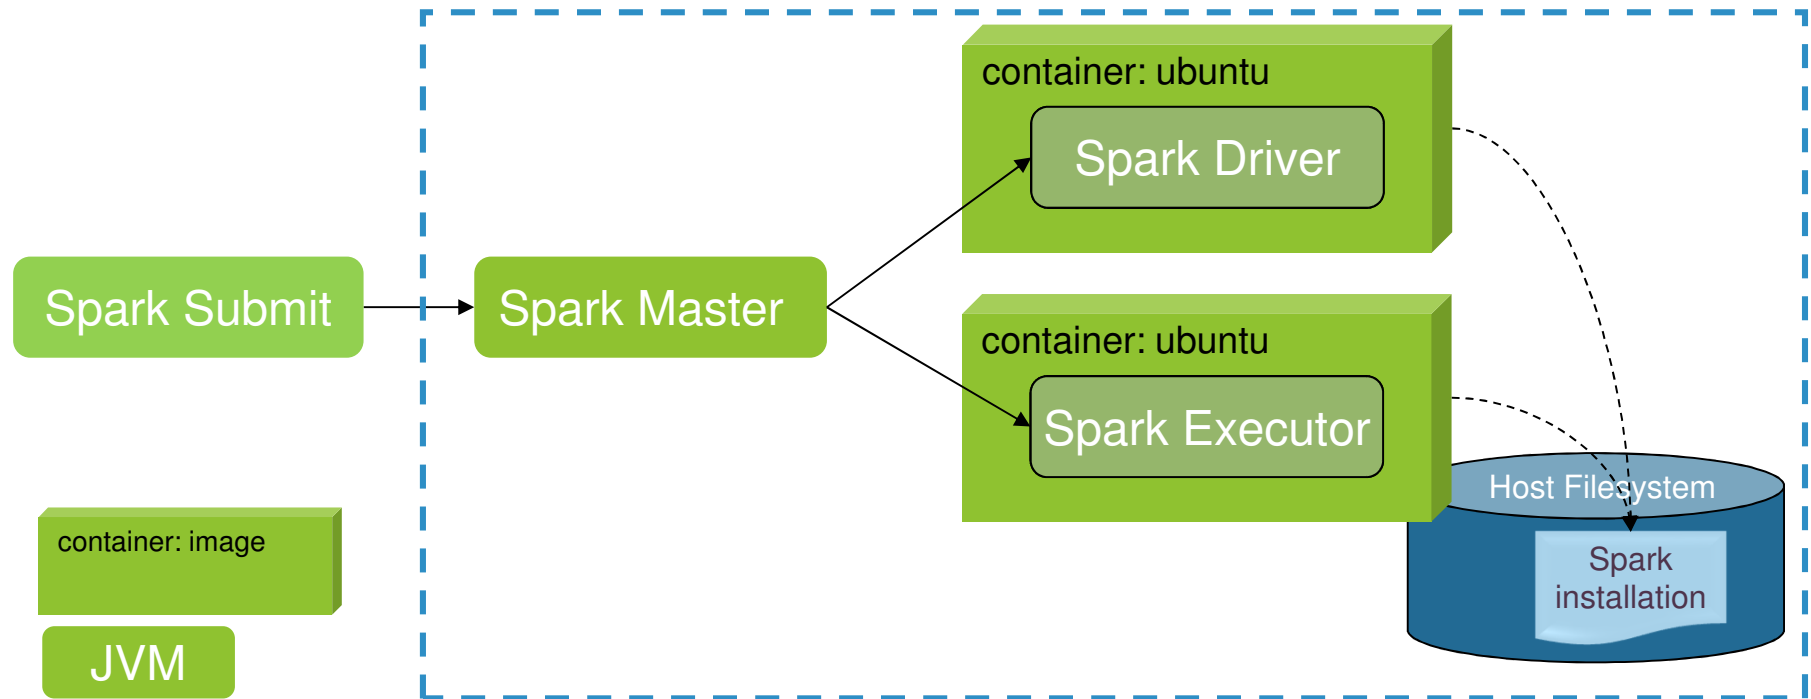
No prebuilt Spark image

- A running container needs an application image
 - Independent to Spark versions
- Seamlessly migrate Spark workloads to a container based environment
 - Assume: Spark is distributed onto the host file system

Regular Spark workload



Running in containers



Creating a container definition for an application image

Extra dependency from
host file system

Docker Container Definition

* Definition name: MyAppDef

* Docker image name: myappimage.v1

Image registry URL: localhost:5111

Data volumes are mount points for the running Docker container. File paths that are defined in your Spark configuration are automatically mounted.

Add a data volume

Data volume

* Host path: /opt/infobatch/lib

* Container path: /opt/infobatch/lib

Environment variable:

☐ Writable

OK Cancel

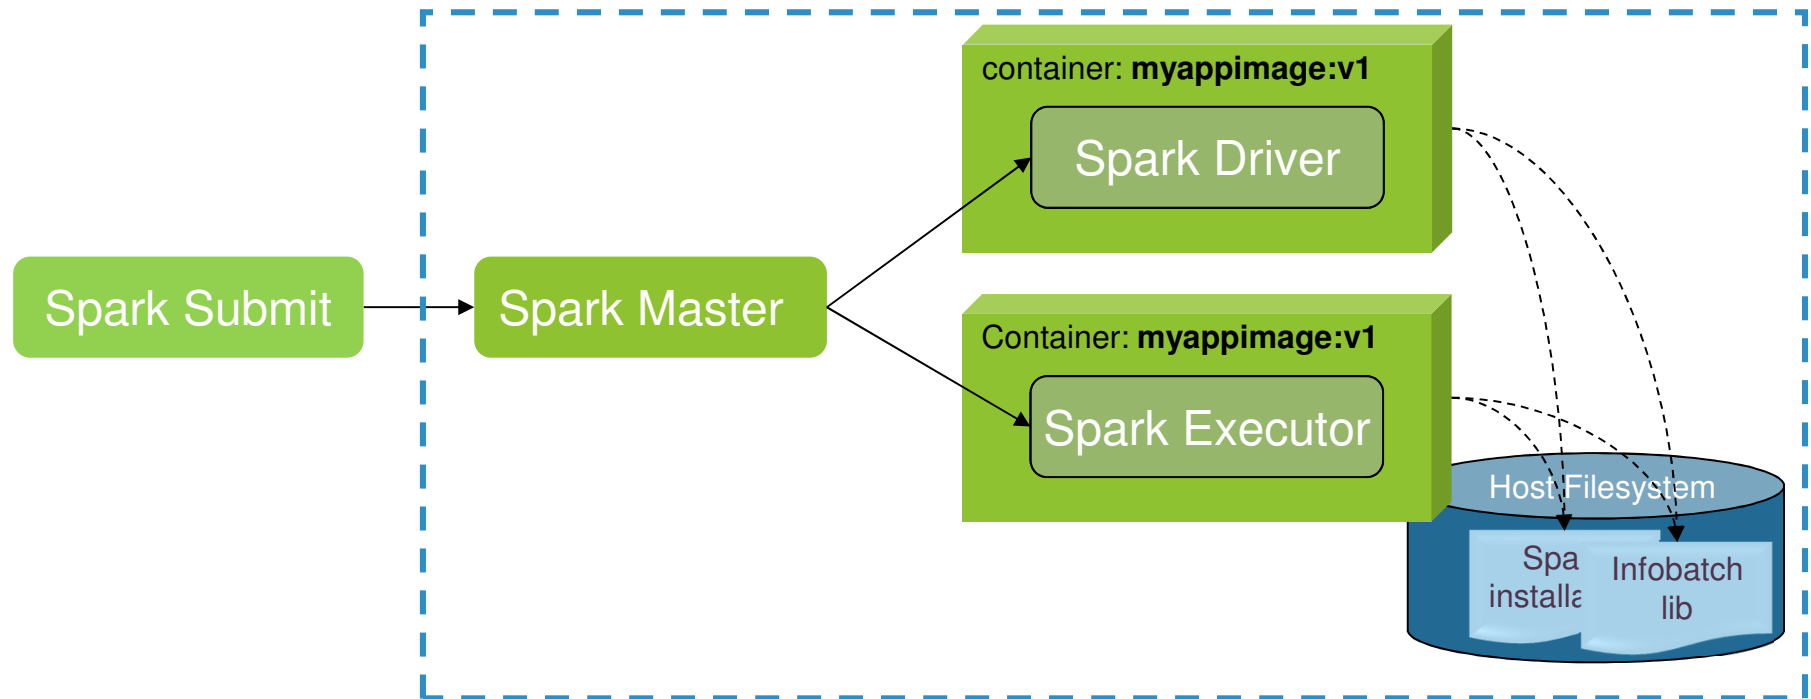
Submitting workload with the container definition

```
spark-submit --class<main-class> --master<master-url> --deploy-mode cluster \  
--conf spark.ego.driver.docker.definition= MyAppDef \  
--conf spark.ego.executor.docker.definition= MyAppDef \  
<application-jar> \  
[application-arguments]
```

Cluster Mode:

Define container specifications for
the drivers and executors

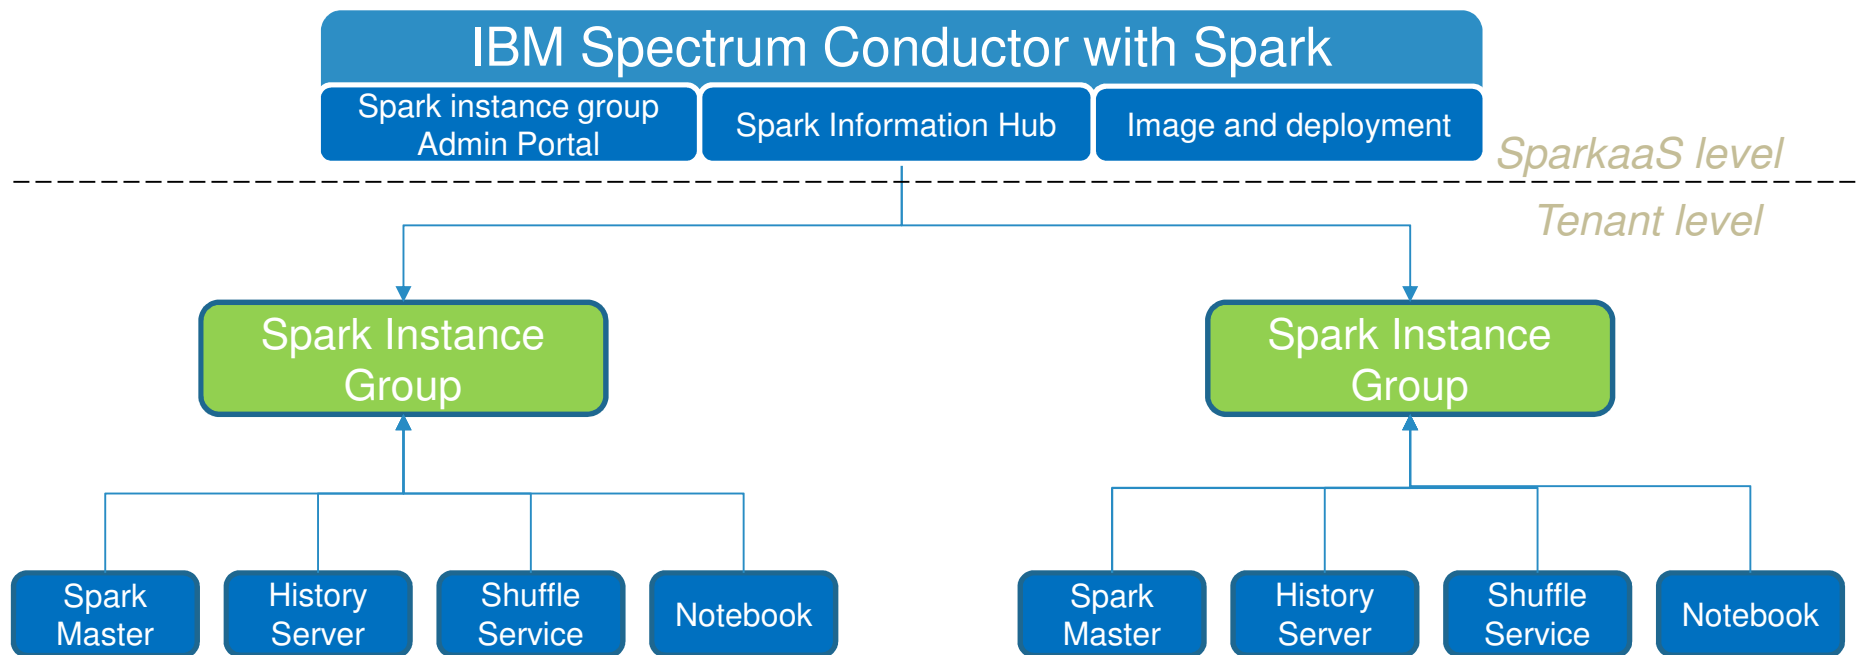
Running in containers



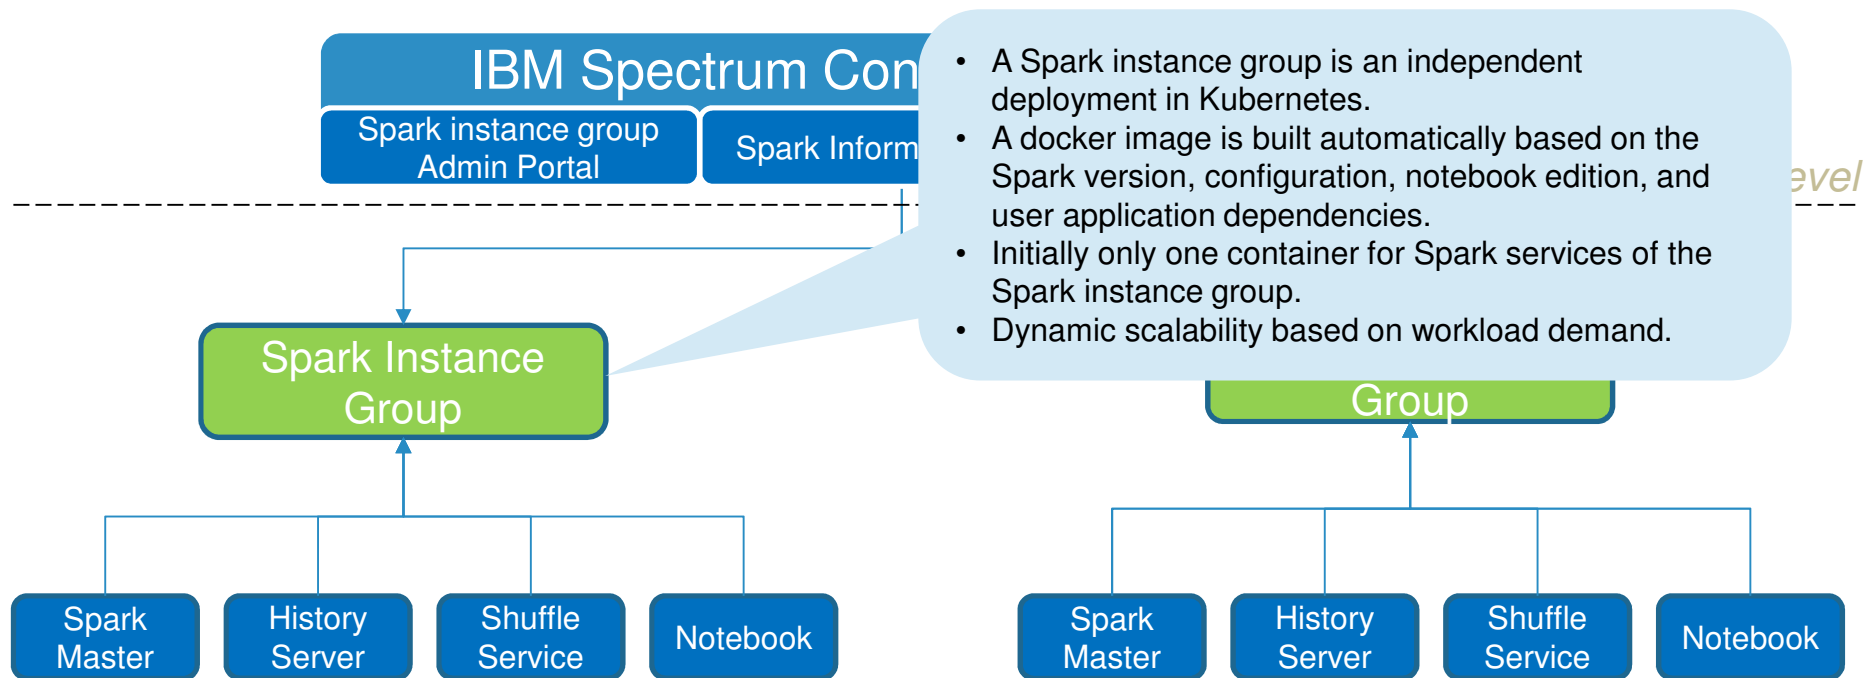
Spark Instance on Kubernetes

- Increase resource utilization
 - Share nodes between Spark and surrounding ecosystem
- Isolation between tenants and apply resource enforcement
 - Each tenant gets a dedicated Spark working instance
 - Tenant price plan can directly map to its resource quota
- Simplify deployment and roll out

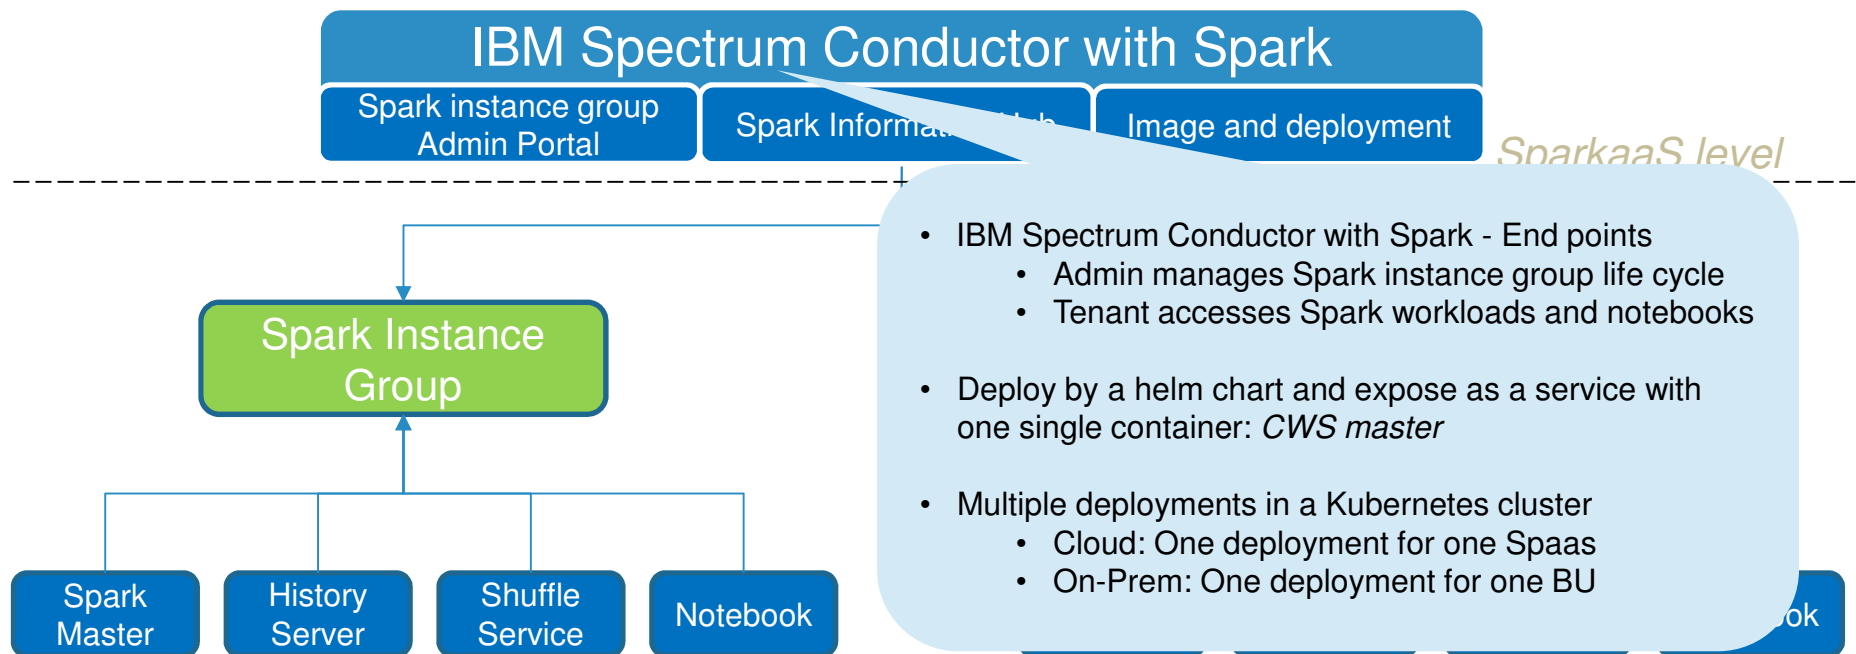
Architecture



Architecture



Architecture



Creating a master container

Kubernetes

```
helm install conductor-spark --name spaas4bu1 --namespace ns4bu1
```

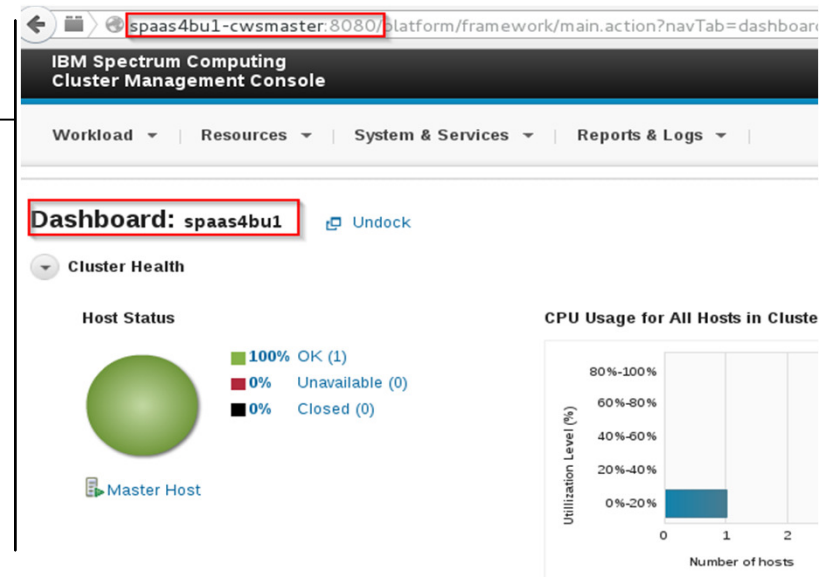
Creating a master container

Kubernetes

Namespace: *ns4bu1*

container: CWS

spaas4bu1_cwsmaster



Creating a Spark instance group

Kubernetes

Namespace: *ns4bu1*

container: CWS

spaas4bu1_cwsmaster

Registry

tenant1
Image

New Spark Instance Group

☆ Basic Settings

Containers Optional

Additional Packages Optional

Required Fields

Instance group name:

tenant1

Spark deployment directory:

/tenant

Execution user for instance group:

root

Spark configuration:

☒ Spark 2.1.0

Configuration

☐ Spark 2.0.1

Configuration

☐ Spark 1.6.1

Configuration

☐ Spark 1.5.2

Configuration

Notebooks:

☒ Jupyter 4.1.0

Configuration

Base data directory:

Deploying a Spark instance group

Kubernetes

Namespace: *ns4bu1*

container: CWS

spaas4bu1_cwsmaster

container: tenant1

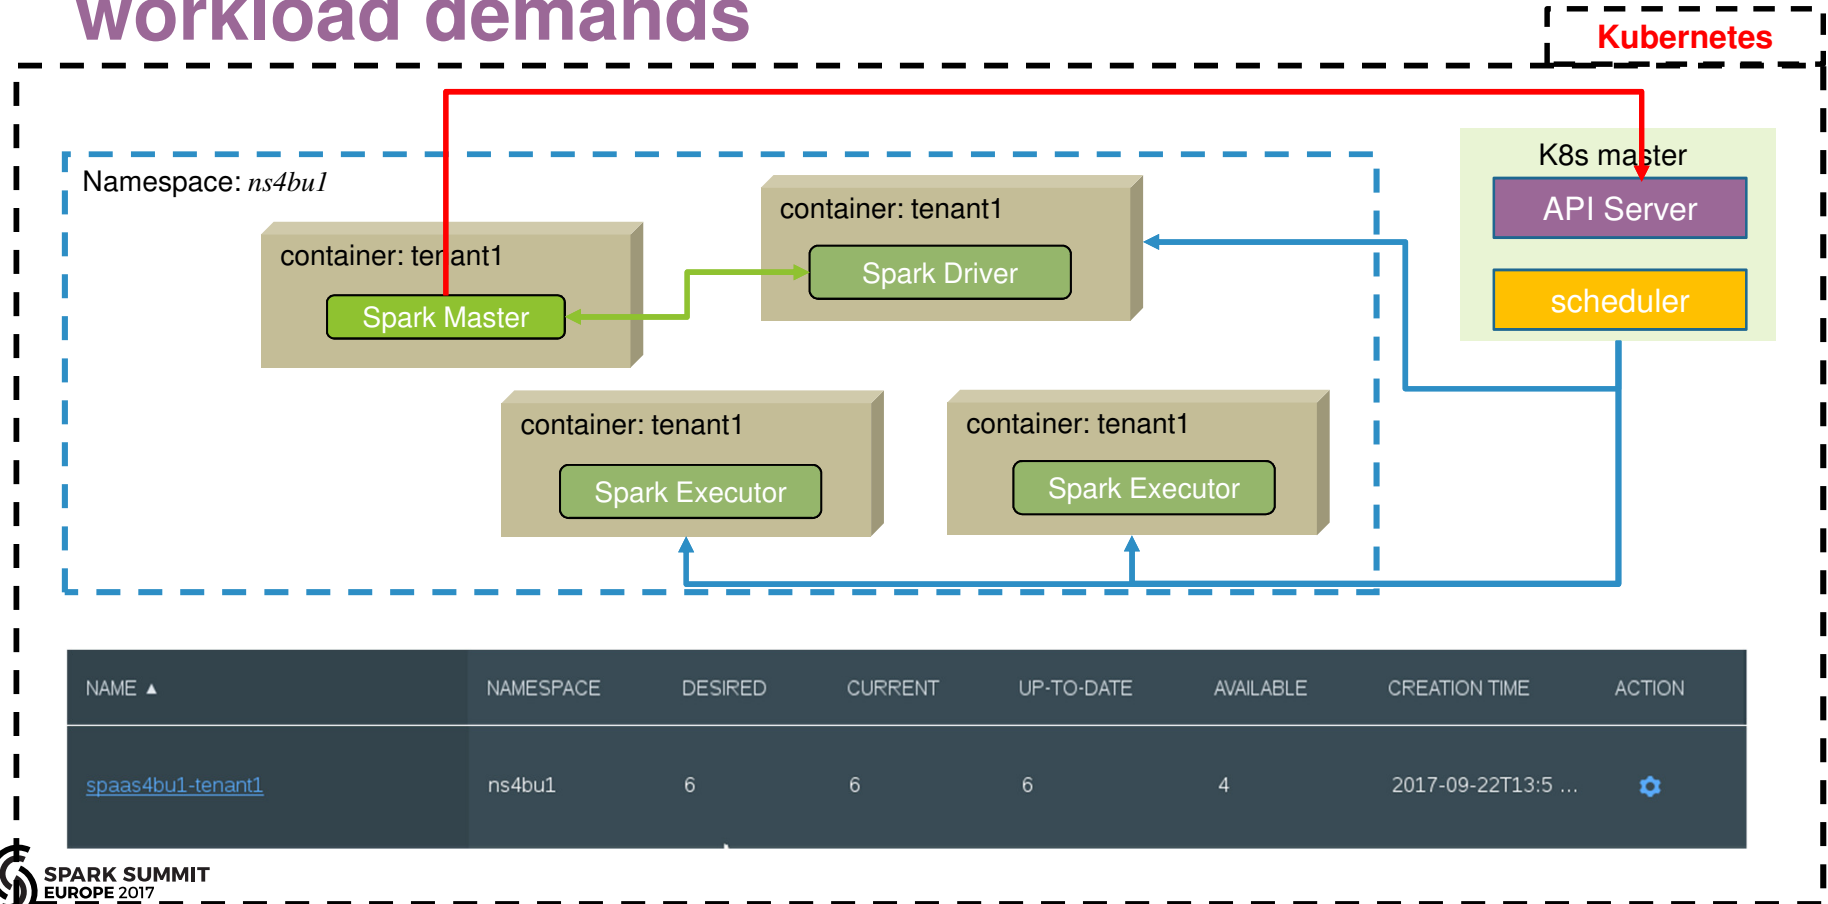
spaas4bu1_tenant1

Registry

tenant1
Image

NAME ▲	NAMESPACE	DESIRED	CURRENT	UP-TO-DATE	AVAILABLE	CREATION TIME	ACTION
spaas4bu1-cwsmaster	ns4bu1	1	1	1	1	2017-09-22T01:4 ...	⚙
spaas4bu1-tenant1	ns4bu1	1	1	1	1	2017-09-22T13:5 ...	⚙

Scaling the Spark instance group based on workload demands




Performance

- Without Dynamic Scaling

Running tasks



Completed tasks




Performance

Running tasks




Completed tasks



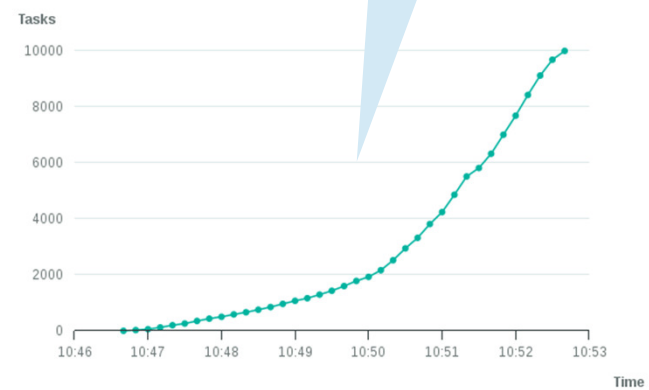
- Without Dynamic Scaling

- With Dynamic Scaling

Running tasks

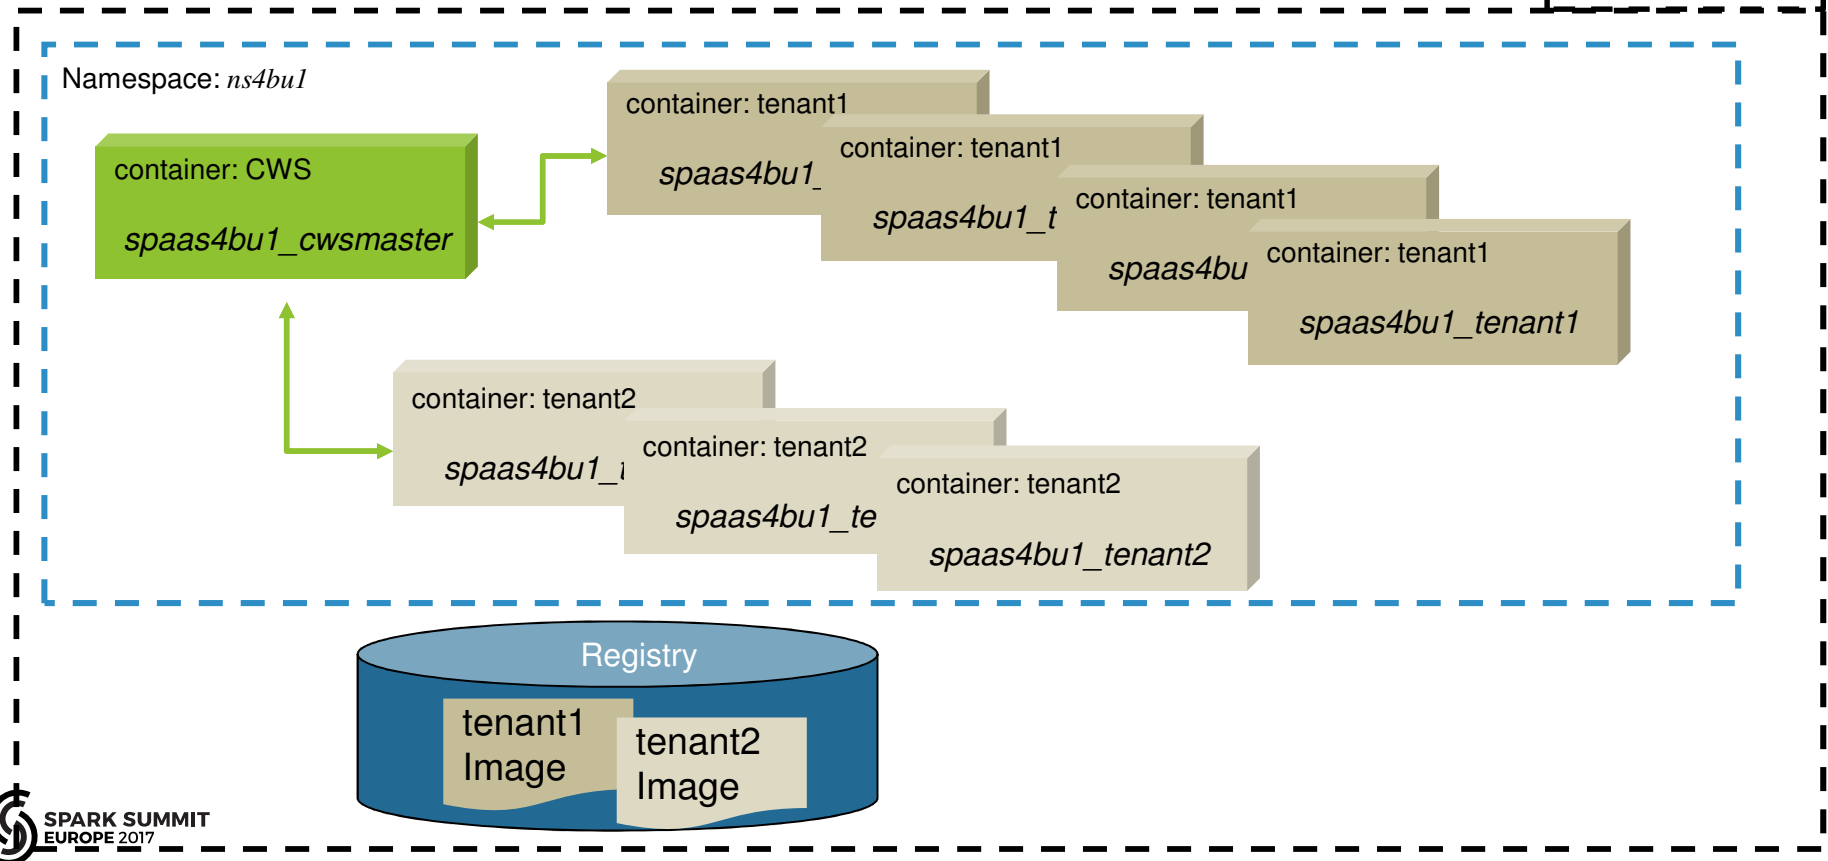


Completed tasks



Multitenancy with Spark instance groups

Kubernetes

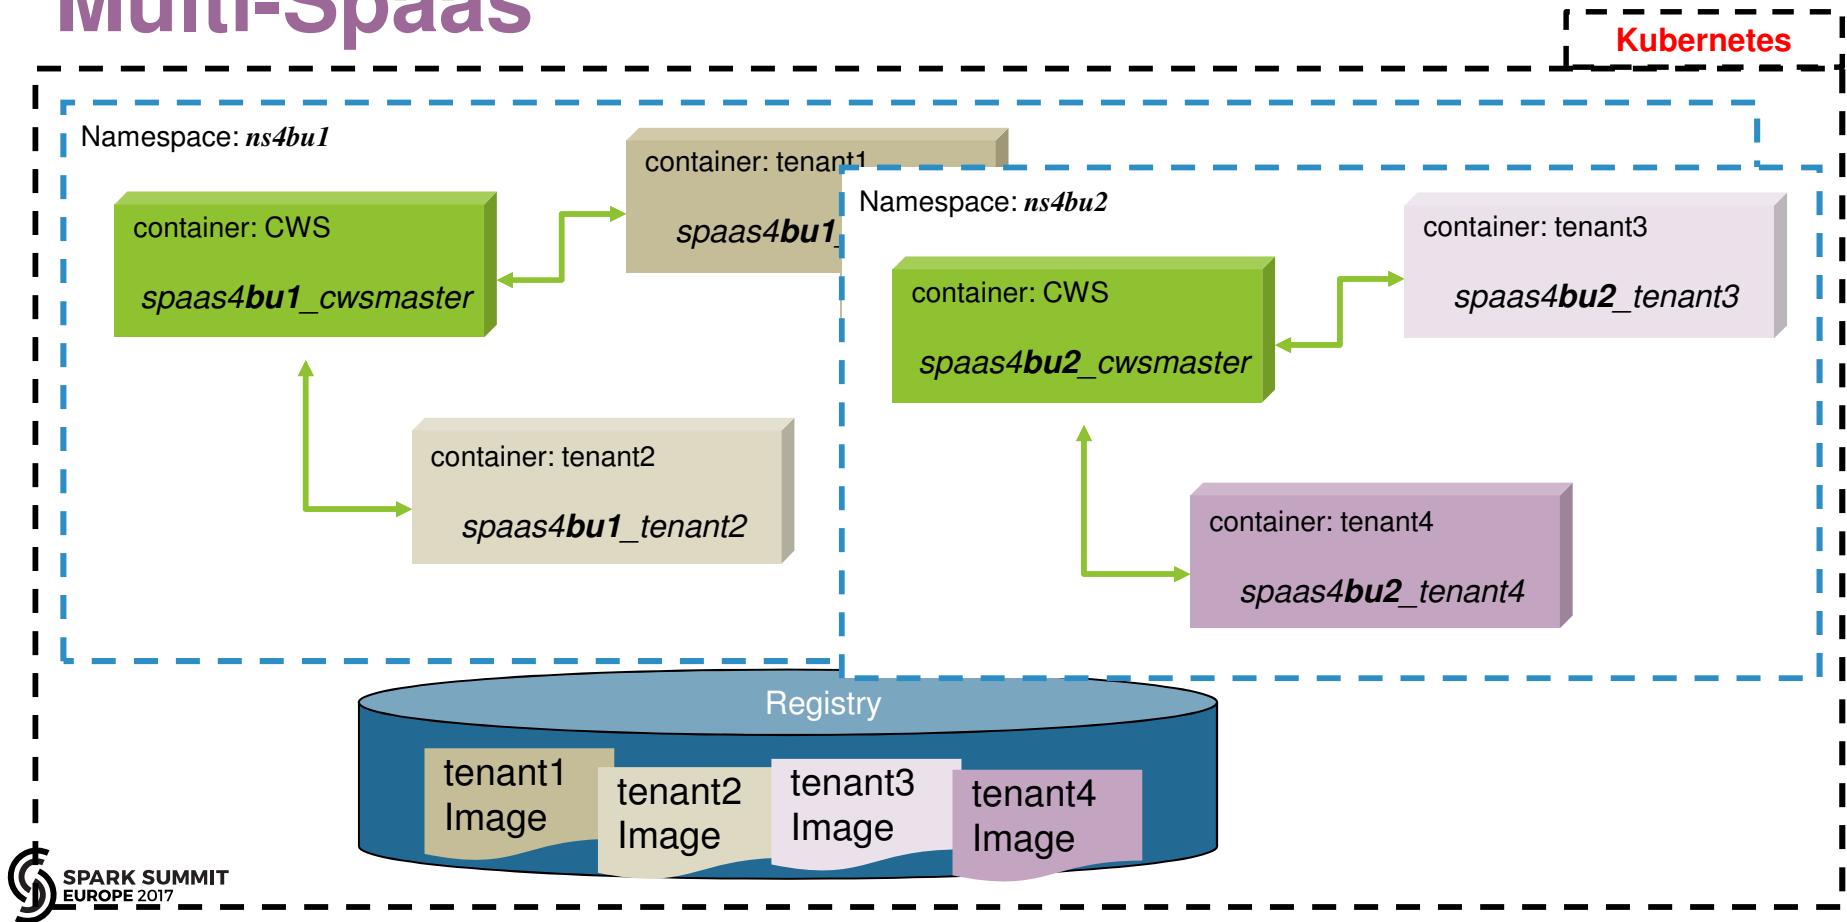


Multi-Spaas

Kubernetes

```
helm install conductor-spark --name spaas4bu1 --namespace ns4bu1  
helm install conductor-spark --name spaas4bu2 --namespace ns4bu2
```

Multi-Spaas



Survey: Spark on Kubernetes

	SPARK-18278	Standalone	IBM Spectrum Conductor with Spark on Kubernetes
Dynamic allocation on demand	Yes	Static	Yes
K8s interaction granularity	Job level	Instance level – static	Instance level – dynamic
Deployment Automation <ul style="list-style-type: none">• Simple deploy by helm charts	No	Yes	Yes
Spark instance per tenant <ul style="list-style-type: none">• Multi-job/workflow/user• Image with user applications• Security	No	limited	Yes

Future work

- Integration with Kubernetes batch workload scheduler
 - Kube-arbitrator (<https://github.com/kubernetes-incubator/kube-arbitrator>)
- Performance comparison with other Spark on Kubernetes solutions



**SPARK
SUMMIT**

Thank You

www.ibm.com/spectrum-conductor

hma@ca.ibm.com