



**SPARK  
SUMMIT**  
EUROPE 2017

# Storage Engine Considerations for your Apache Spark Applications

Mladen Kovacevic, Senior Solutions Architect  
Cloudera Inc.

**#EUdev10**

# Outline

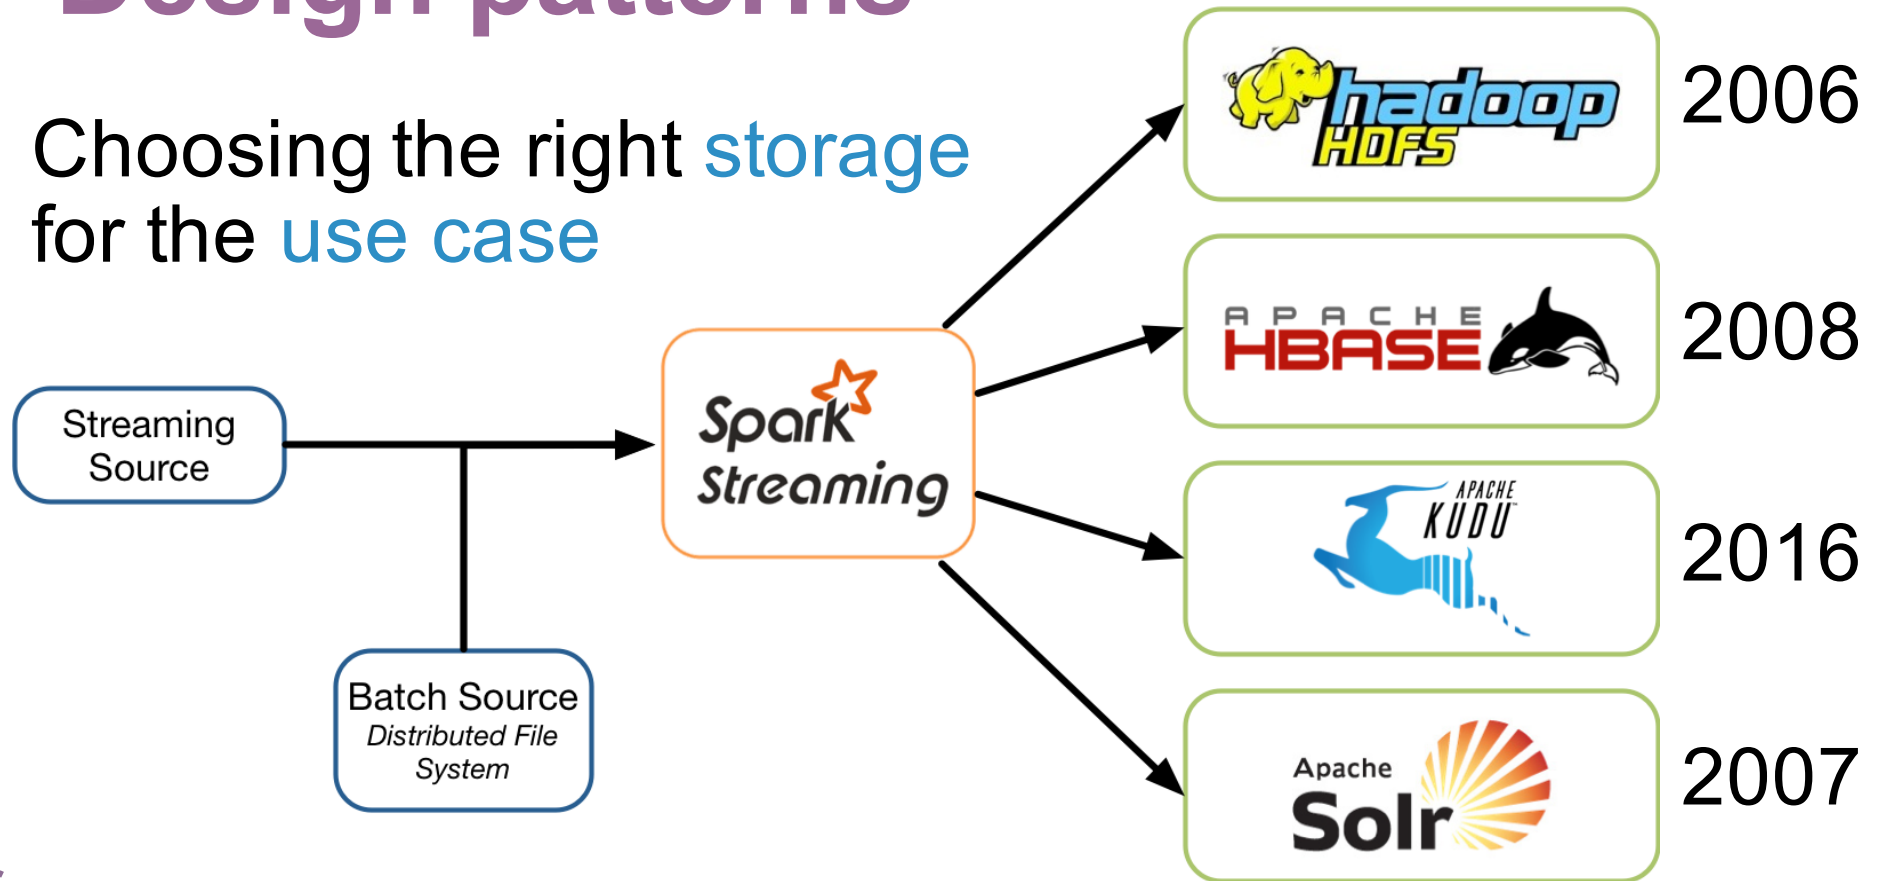
- Motivation – store your data – where *exactly*?
- Storage Capabilities:
  - HDFS
  - HBase
  - Kudu
  - Solr
- Asking the right questions
- Decide on right storage solution

# Motivation

- Spark, SparkStreaming, SparkSQL – great for processing – need a place to store content
- Integration with variety of storage systems
- *Ingest* and *consumption* requirements – use case!

# Design patterns

Choosing the right **storage**  
for the **use case**



# HDFS


- Distributed file system – cheap, scalable, storage
- Immutable – “record” changes are painful
- Columnar file formats - ideal for analytics
- SQL overlays (SparkSQL, Hive Metastore, more) to define schema

## Highlights

Very high throughput, painful random IO, **batch** oriented, coding overhead (ie. dealing with small files problems), any file

# HDFS design pattern

```
df.write.parquet("/data/person_table")
```



- Small files accumulate
- External processes, or additional application logic to manage these files
- Partition management
- Manage metadata carefully (depends on ecosystem)
- Considerations- changing dimensions (fast/slow)
- Late arriving data

# HBase

- NoSQL engine, manages files on HDFS
- Key-value, distributed storage engine
- No data types – just binary fields
- Thousands to millions of columns
- Store **entity** data (profiles of people, devices, accounts)

## Highlights

Very fast random IO, low throughput, **NRT** oriented, challenging BI, no strict data types

# HBase design pattern



- HBase Connection anywhere in Spark/SparkStreaming app
- SparkSQL/DataFrames, Bulk Load
- Primary storage for ingestion or complementary preserving state
- NoSQL store vs. structured
- Near-real-time

CDH hbase-spark: [https://github.com/cloudera/hbase/tree/cdh5-1.2.0\\_5.13.0/hbase-spark](https://github.com/cloudera/hbase/tree/cdh5-1.2.0_5.13.0/hbase-spark)

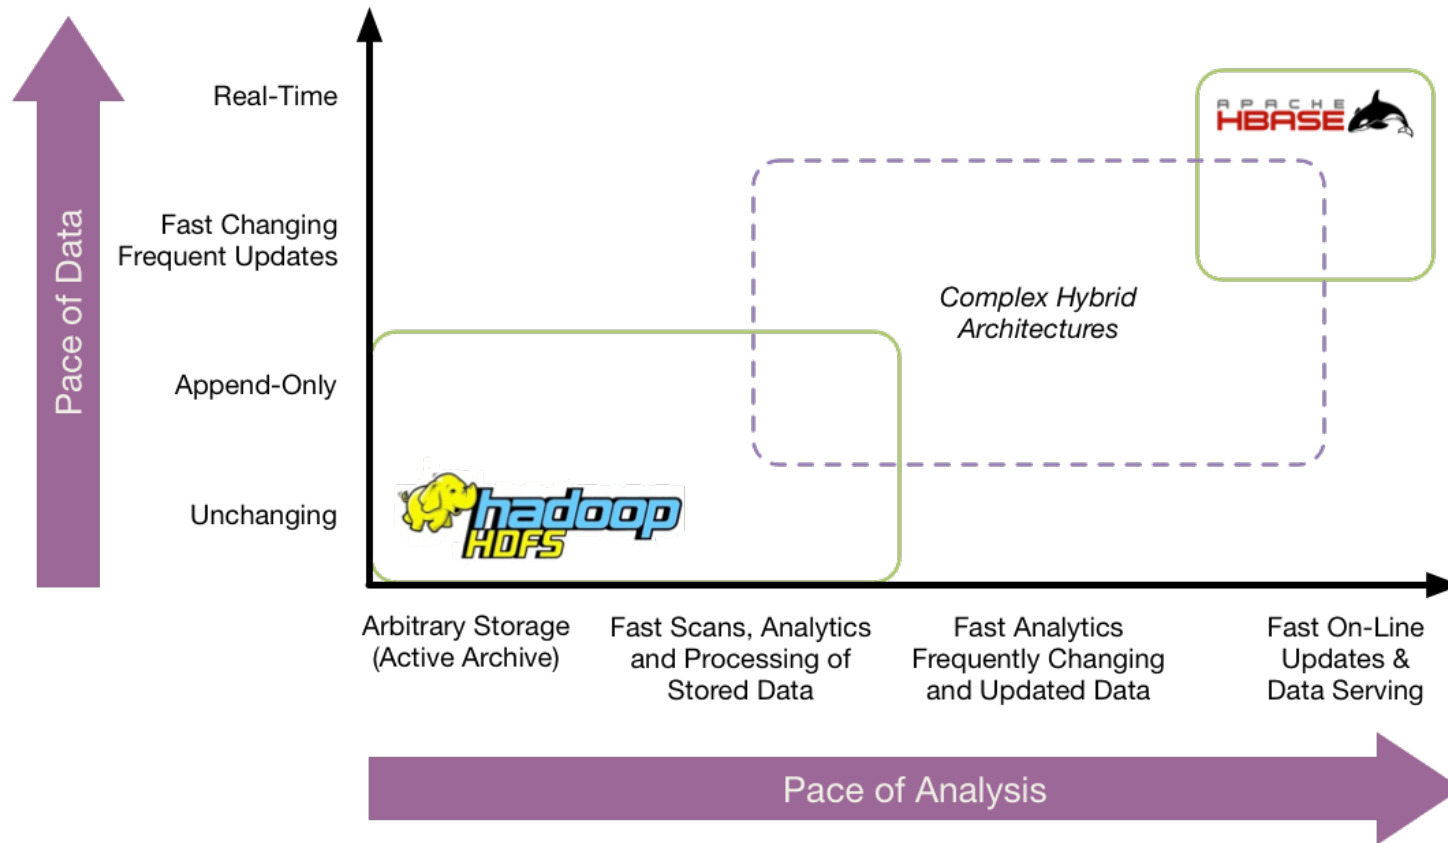
CDH HBase and Spark docs: <http://archive.cloudera.com/cdh5/cdh/5/hbase-1.2.0-cdh5.13.0/book.html#spark>

Upstream hbase-spark (watch for updates in HBase 2.x release): <https://github.com/apache/hbase/tree/master/hbase-spark/>

Upstream HBase and Spark book: <http://hbase.apache.org/book.html#spark>



# Analytic Gap



# Kudu


- Storage system for tables of structured data
- Bring-your-own-SQL (SparkSQL, Impala), NoSQL-like API, integration with Spark, MapReduce, more..
- Columnar, key partitioning by range and/or hash
- Limited number of columns (strongly typed)

## Highlights

Fast random IO, fast throughput, NRT oriented, terrific for BI, structured data

# Kudu design pattern

```
df.write.options(kuduOptions).mode("append").kudu  
OR  
kuduContext.insertRows()
```

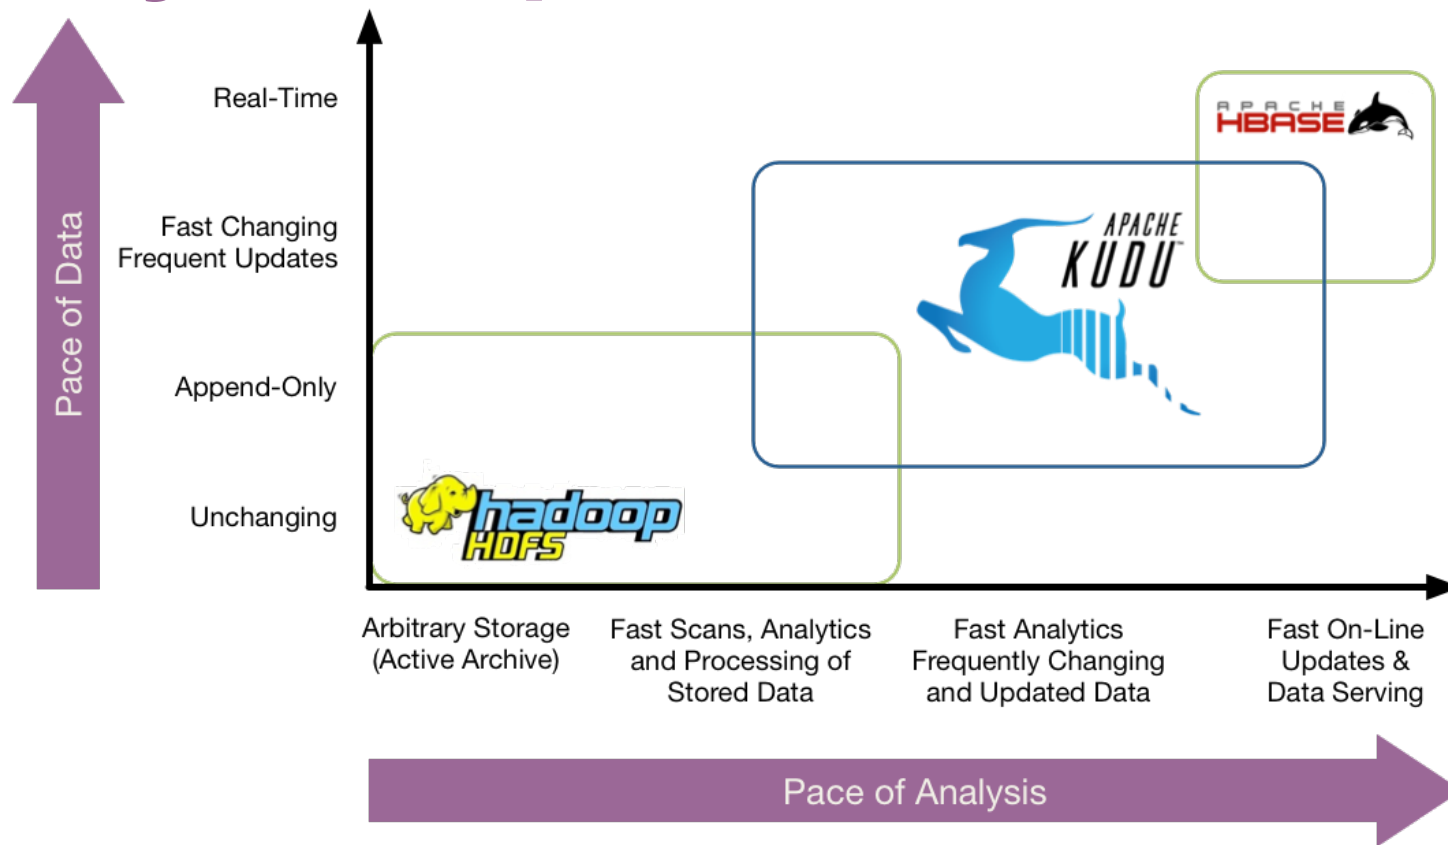


- DataFrame perfect match for Kudu (structured)
- Data available *immediately* to SQL engines (Impala, SparkSQL)
- Ideal case is append with moderate updates

Kudu Integration with Spark: [http://kudu.apache.org/docs/developing.html#kudu\\_integration\\_with\\_spark](http://kudu.apache.org/docs/developing.html#kudu_integration_with_spark)

Up and running with Apache Spark on Apache Kudu: <https://blog.cloudera.com/blog/2017/02/up-and-running-with-apache-spark-on-apache-kudu/>

# Analytic Gap Filled



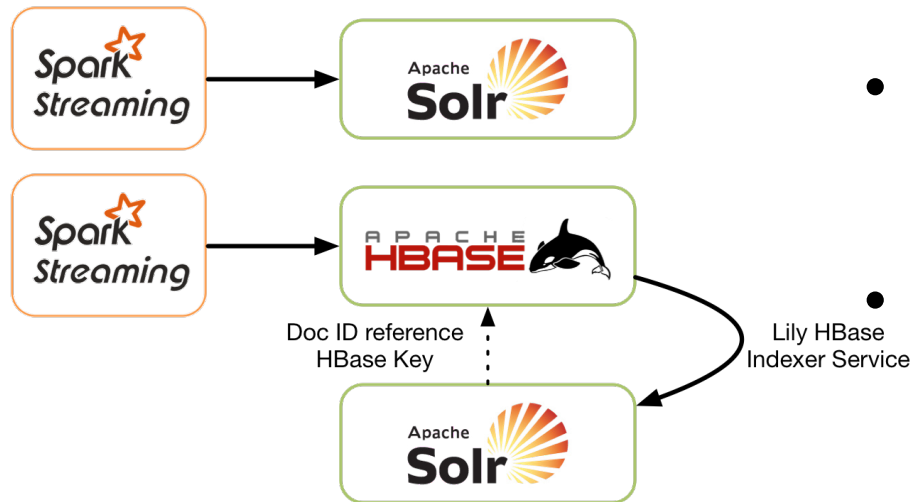
# Solr

- Distributed **index** enabling **search** capabilities (Lucene)
- Typed, REST API based, search index query processing
- Search interface, faceting, integration with HBase storing content (typically) in HDFS

## Highlights

High random IO, low throughput, multi-faceted use cases, **NRT** oriented, terrific for BI with the right tools (non-SQL), loose schema, data types

# Solr design pattern



- Prepare Solr document, add to SolrCloud directly OR
- Write to HBase, leverage Lily HBase Indexer service to update Solr
- Store complete record in HBase, while indexed fields for search in Solr
- NRT availability (short soft commits)

# Questions we ask (1)

- How many voters have cast their ballots by city thus far in the election, by the second?
  - streaming data into 'voter' table, aggregate query, immediate data availability : **Kudu**
- How many people watched last nights game compared to the night before?
  - daily batch, aggregate query : **HDFS parquet**

## Questions we ask (2)





























- What version is my device running and how many dropped packets do I have?
  - streaming entity profile data, metrics may change per release, many updates, specific device, NRT: **HBase**
- Which tweets talk to the housing market, in the 21-30 age group?
  - streaming, keyword search, facet filtering : **Solr**







# Use case questionnaire

- Consumption interface: **SQL** (JDBC/ODBC) vs. **API**
- Near-real-time requirement for **consumers**
- Ingestion **rate** (can we keep up?)
- **Entity** vs. **Events** (time-based)
- **Append-only** vs **moderate** updates vs **many** updates
- **Distinct values** in dataset

# Storage considerations (1)

Criteria				
SQL interface				
API interface				
Near-real-time ingestion				
Append-only + available for query				
Appends with moderate updates				
Mostly updates				

# Storage considerations (2)

Criteria				
Entity based data	●	●	●	●
Event based data (time-series)	●	●	●	●
High distinct values	●	●	●	●
Many and unknown attributes	●	●	●	●
Binary data (Images, PDFs, etc)	●	●	●	●
Analytics	●	●	●	●

# Wrap-up

- Review entire use-case *end-to-end* early
- Understand storage capabilities
- Ask the right questions (upstream/consumers)
- Consider security, architecture and development costs
- Decide on the right storage solution