



# Web-Scale Graph Analytics with Apache Spark

Tim Hunter, Databricks

#EUds6

# About Me




- Tim Hunter
- Software engineer @ Databricks
- Ph.D. from UC Berkeley in Machine Learning
- Very early Spark user
- Contributor to MLlib
- Co-author of TensorFrames, GraphFrames, Deep Learning Pipelines

# Outline

- Why GraphFrames
- Writing scalable graph algorithms with Spark
  - *Where is my vertex?* Indexing data
  - *Connected Components*: implementing complex algorithms with Spark and GraphFrames
  - *The Social Network*: real-world issues
- Future of GraphFrames

# Graphs are everywhere

Example: airports & flights between them



*Vertices:*

id	City	State
"JFK"	"New York"	NY

*Edges:*

src	dst	delay	tripID
"JFK"	"SEA"	45	1058923

# Apache Spark's GraphX library

- General-purpose graph processing library
- Built into Spark
- Optimized for fast distributed computing
- Library of algorithms: PageRank, Connected Components, etc.

## Issues:

- No Java, Python APIs
- Lower-level RDD-based API (vs. DataFrames)
- Cannot use recent Spark optimizations: Catalyst query optimizer, Tungsten memory management


# The GraphFrames Spark Package

## Brings DataFrames API for Spark

- Simplifies interactive queries
- Benefits from DataFrames optimizations
- Integrates with the rest of Spark ecosystem

Collaboration between Databricks, UC Berkeley & MIT

# Dataframe-based representation



vertices DataFrame

id	City	State
----	------	-------

"JFK"	"New York"	NY
-------	------------	----

"SEA"	"Seattle"	WA
-------	-----------	----

edges DataFrame

src	dst	delay	tripID
-----	-----	-------	--------

"JFK"	"SEA"	45	1058923
-------	-------	----	---------

"DFW"	"SFO"	-7	4100224
-------	-------	----	---------

# Supported graph algorithms

- Find Vertices:
  - PageRank
  - **Motif finding**
- Communities:
  - **Connected Components**
  - Strongly Connected Components
  - Label propagation (LPA)
- Paths:
  - **Breadth-first search**
  - Shortest paths
- Other:
  - **Triangle count**
  - SVD++

(Bold: native DataFrame implementation)



# Assigning integral vertex IDs

... lessons learned

# Pros of integer vertex IDs

GraphFrames take arbitrary vertex IDs.

→ convenient for users

Algorithms prefer integer vertex IDs.

→ optimize in-memory storage

→ reduce communication

Our task: Map unique vertex IDs to unique (long) integers.

# The hashing trick?

- Possible solution: hash vertex ID to long integer
- What is the chance of collision?
  - $1 - (N-1)/N * (N-2)/N * \dots$
  - seems unlikely with long range  $N=2^{64}$
  - with 1 billion nodes, the chance is ~5.4%
- Problem: collisions change graph topology.

Name	Hash
Sue Ann	84088
Joseph	-2070372689
Xiangrui	264245405
Felix	67762524

# Generating unique IDs

Spark has built-in methods to generate unique IDs.

- **RDD:** `zipWithUniqueId()`, `zipWithIndex()`
- **DataFrame:** `monotonically_increasing_id()`

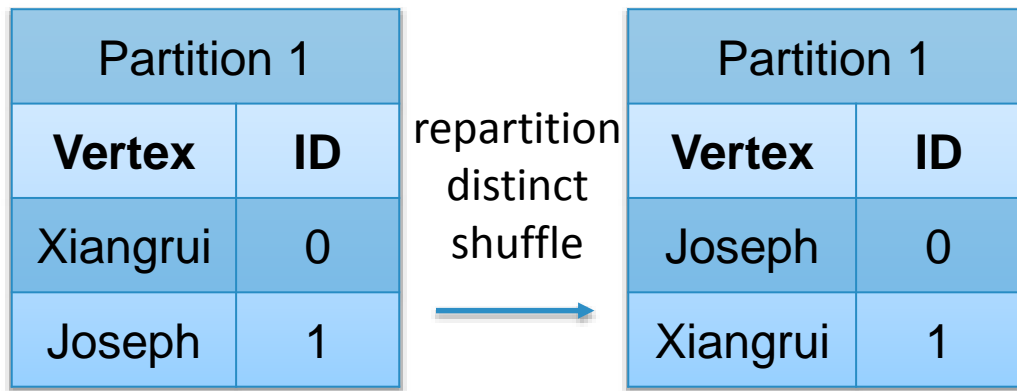
Possible solution: just use these methods

# How it works

Partition 1		Partition 2		Partition 3	
Vertex	ID	Vertex	ID	Vertex	ID
Sue Ann	0	Xiangrui	$100 + 0$	Veronica	$200 + 0$
Joseph	1	Felix	$100 + 1$	...	$200 + 1$

# ... but not always

- DataFrames/RDDs are immutable and reproducible by design.
- However, records do not always have stable orderings.
  - distinct
  - repartition
- `cache()` does not help.



# Our implementation

We implemented (v0.5.0) an expensive but correct version:

1. (hash) re-partition + distinct vertex IDs
2. sort vertex IDs within each partition
3. generate unique integer IDs

# Connected Components



# Connected Components

Assign each vertex a component ID such that vertices receive the same component ID iff they are connected.


## Applications:

- fraud detection

- [Spark Summit 2016 keynote from Capital One](#)

- clustering

- entity resolution



# Naive implementation (GraphX)

1. Assign each vertex a unique component ID.
2. Iterate until convergence:
  - For each vertex  $v$ , update:  
component ID of  $v \leftarrow$  Smallest component ID in neighborhood of  $v$

Pro: easy to implement


Con: slow convergence on large-diameter graphs

# Small-/large-star algorithm

Kiveris et al. "Connected Components in MapReduce and Beyond."


1. Assign each vertex a unique ID.
2. Iterate until convergence:
  - (small-star) for each vertex,  
    connect smaller neighbors to smallest neighbor
  - (big-star) for each vertex,  
    connect bigger neighbors to smallest neighbor (or  
    itself)

# Small-star operation




(a) The small-star operation at node 8.

# Big-star operation



(b) The large-star operation at node 8.

# Another interpretation



adjacency matrix

	1	5	7	8	9
1				X	
5				X	
7				X	
8					X
9					

# Small-star operation

	1	5	7	8	9
1				X	
5				X	
7				X	
8					X
9					

rotate & lift

	1	5	7	8	9
1		X	X	X	
5					
7					
8					X
9					

# Big-star operation

	1	5	7	8	9
1				x	
5				x	
7				x	
8					x
9					

lift

	1	5	7	8	9
1				x	x
5				x	
7				x	
8					
9					



# Convergence

	1	5	7	8	9
1	x	x	x	x	x
5					
7					
8					
9					

# Properties of the algorithm

- Small-/big-star operations do not change graph connectivity.
- Extra edges are pruned during iterations.
- Each connected component converges to a star graph.
- Converges in  $\log^2(\text{\#nodes})$  iterations

# Implementation

Iterate:

- filter
- self-join

Challenge: handle these operations at scale.

# Skewed joins

Real-world graphs contain big components.

The "Justin Bieber problem" at Twitter

→ data skew during connected components iterations

src	Component id	neighbors
0	0	2,000,000
1	0	10
2	3	5

join

src	dst
0	1
0	2
0	3
0	4
...	...
0	2,000,000
1	3
2	5

# Skewed joins

(#nbrs > 1,000,000)

src	Component id	neighbors
0	0	2,000,000

broadcast join

src	dst
0	1
0	2
0	3
0	4
...	...
0	2,000,000

---

union

1	0	10
2	3	5

hash join

1	3
2	5

# Checkpointing

We checkpoint every 2 iterations to avoid:

- query plan explosion (exponential growth)
- optimizer slowdown
- disk out of shuffle space
- unexpected node failures

# Experiments

twitter-2010 from [WebGraph datasets](#) (small diameter)

- 42 million vertices, 1.5 billion edges

16 r3.4xlarge workers on Databricks

- GraphX: 4 minutes

- GraphFrames: 6 minutes

- algorithm difference, checkpointing, checking skewness

# Experiments

uk-2007-05 from [WebGraph datasets](#)

–105 million vertices, 3.7 billion edges

16 r3.4xlarge workers on Databricks

–GraphX: 25 minutes

- slow convergence

–GraphFrames: 4.5 minutes



# Experiments

regular grid 32,000 x 32,000 (large diameter)

- 1 billion nodes, 4 billion edges

32 r3.8xlarge workers on Databricks

- GraphX: failed

- GraphFrames: 1 hour

# Future improvements

## GraphFrames

- better graph partitioning
- letting Spark SQL handle skewed joins and iterations
- graph compression

## Connected Components

- local iterations
- node pruning and better stopping criteria

# Thank you!

- <http://graphframes.github.io>
- <https://docs.databricks.com>



# 2 types of graph representations

## Algorithm-based



Standard & custom algorithms  
Optimized for batch processing

## Query-based



Motif finding  
Point queries & updates

GraphFrames: Both algorithms & queries (but not point updates)

# Graph analysis with GraphFrames

Simple queries

Motif finding

Graph algorithms


# Simple queries

SQL queries on vertices & edges

# Motif finding

Search for structural patterns within a graph.

```
val paths: DataFrame =  
  g.find("(a)-[e1]->(b);  
        (b)-[e2]->(c);  
        !(c)-[]->(a)")
```






# Motif finding

Search for structural patterns within a graph.


```
val paths: DataFrame =  
  g.find("(a) - [e1] -> (b) ;  
        (b) - [e2] -> (c) ;  
        ! (c) - [] -> (a) ")
```



# Motif finding

Search for structural patterns within a graph.


```
val paths: DataFrame =  
  g.find("(a) - [e1] -> (b) ;  
        (b) - [e2] -> (c) ;  
        ! (c) - [] -> (a) ")
```



# Motif finding

Search for structural patterns within a graph.

```
val paths: DataFrame =  
  g.find("(a) - [e1] -> (b) ;  
        (b) - [e2] -> (c) ;  
        ! (c) - [] -> (a) ")
```




# Motif finding

Search for structural patterns within a graph.

```
val paths: DataFrame =  
  g.find("(a) - [e1] -> (b) ;  
        (b) - [e2] -> (c) ;  
        ! (c) - [] -> (a) ")
```

Then filter using vertex  
& edge data.

```
paths.filter("e1.delay > 20")
```



# Graph algorithms

Find important vertices

- PageRank

Find paths between sets of vertices

- Breadth-first search (BFS)
- Shortest paths

Find groups of vertices  
(components, communities)

- Connected components
- Strongly connected components
- Label Propagation Algorithm (LPA)

Other

- Triangle counting
- SVDPlusPlus

# Saving & loading graphs

## Save & load the DataFrames.

```
vertices = sqlContext.read.parquet(...)  
edges = sqlContext.read.parquet(...)  
g = GraphFrame(vertices, edges)
```

```
g.vertices.write.parquet(...)  
g.edges.write.parquet(...)
```